

My Data Science Notes

Michael Foley

2020-02-04

Contents

Intro	7
1 Probability	9
1.1 Principles	9
1.2 Discrete Distributions	11
1.3 Continuous Distributions	16
2 Inference	27
3 Experiments	29
3.1 Example one	29
3.2 Example two	29
4 Regression	31
5 Classification	33
6 Regularization	35
7 Non-linear Models	37
7.1 Splines	37
7.2 MARS	38
7.3 GAM	40

8 Decision Trees	43
8.1 Classification Tree	44
8.2 Regression Trees	69
8.3 Bagging	84
8.4 Random Forests	84
8.5 Gradient Boosting	98
8.6 Summary	161
9 Reference	165
10 Support Vector Machines	167
10.1 Maximal Margin Classifier	167
10.2 Support Vector Classifier	168
10.3 Support Vector Machines	169
10.4 Example	170
10.5 Using Caret	180
11 Principal Components Analysis	183
12 Clustering	185
13 Text Mining	187
Appendix	189
Publishing to BookDown	191
Shiny Apps	193
14 rstudio::conf	195
14.1 Open Source Software for Data Science	196
14.2 Data, visualization, and designing with AI	196
14.3 Deploying End-to-End Data Science with Shiny, Plumber, and Pins	196
14.4 Health Connected Siloed Data Sources and Streamlined Reporting Using R	196

14.5 Building a new data science pipeline for teh FT with RStudio Connect	196
14.6 How to win an AI Hackathon without using AI	196
14.7 Production-grade Shiny Apps with golem	196
14.8 Making the Shiny Contest	196
14.9 Styling Shiny apps with Sass and Bootstrap 4	196
14.10R: Then and Now	196
14.11Journalism with RStudio, R, and teh Tidyverse	196
14.12Learning R with humorous side projects - or - Technical debt is a social problem	196
14.13FlatironKitchen: How we overhauled a Frankensteinian SQL workflow with the tidyverse	196
14.14Making better spaghetti (plots): Exploring longitudinal data with brolgar package	196
14.15Object of type ‘closure’ is not subsettable	196
14.16Branding and Packaging Reports with R Markdown	196
14.17The Glamour of Graphics	196
14.18Don’t repeat yourself, talk to yourself! Repeated reporting in the R universe	196
14.19RStudio 1.3 Sneak Preview	196
14.20One R Markdown Document, Fourteen Demos - or - Tidyverse 2019-20	196
14.21Best Practices for programing with ggplot2	196
14.22Spruce up your ggplot2 visualization with formatted text	196
14.23The little package that could: taking visualization to the next level with the scales package	196
14.24Advances in tidyeval	196
14.25	196

Intro

These notes are pulled from various classes, tutorials, books, etc. and are intended for my own consumption. If you are finding this on the internet, I hope it is useful to you, but you should know that I am just a student and there's a good chance whatever you're reading here is mistaken. In fact, that should probably be your null hypothesis... or your prior. Whatever.

Chapter 1

Probability

1.1 Principles

Here are three rules that come up all the time.

- $Pr(A \cup B) = Pr(A) + Pr(B) - Pr(AB)$. This rule generalizes to $Pr(A \cup B \cup C) = Pr(A) + Pr(B) + Pr(C) - Pr(AB) - Pr(AC) - Pr(BC) + Pr(ABC)$.
- $Pr(A|B) = \frac{P(AB)}{P(B)}$
- If A and B are independent, $Pr(A \cap B) = Pr(A)Pr(B)$, and $Pr(A|B) = Pr(A)$.

Uniform distributions on finite sample spaces often reduce to counting the elements of A and the sample space S , a process called combinatorics. Here are three important combinatorial rules.

Multiplication Rule. $|S| = |S_1||S_k|$.

How many outcomes are possible from a sequence of 4 coin flips and 2 rolls of a die? $|S| = |S_1| \cdot |S_2| \dots |S_6| = 2 \cdot 2 \cdot 2 \cdot 2 \cdot 6 \cdot 6 = 288$.

How many subsets are possible from a set of $n=10$ elements? In each subset, each element is either included or not, so there are $2^n = 1024$ subsets.

How many subsets are possible from a set of $n=10$ elements taken k at a time with replacement? Each experiment has n possible outcomes and is repeated k times, so there are n^k subsets.

Permutations. The number of *ordered* arrangements (permutations) of a set of $|S| = n$ items taken k at a time *without* replacement has $n(n-1) \dots (n-k+1)$

subsets because each draw is one of k experiments with decreasing number of possible outcomes.

$${}_nP_k = \frac{n!}{(n-k)!}$$

Notice that if $k = 0$ then there is 1 permutation; if $k = 1$ then there are n permutations; if $k = n$ then there are $n!$ permutations.

How many ways can you distribute 4 jackets among 4 people? ${}_nP_k = \frac{4!}{(4-4)!} = 4! = 24$

How many ways can you distribute 4 jackets among 2 people? ${}_nP_k = \frac{4!}{(4-2)!} = 12$

Subsets. The number of *unordered* arrangements (combinations) of a set of $|S| = n$ items taken k at a time *without* replacement has

$${}_nC_k = \binom{n}{k} = \frac{n!}{k!(n-k)!}$$

combinations and is called the binomial coefficient. The binomial coefficient is the number of different subsets. Notice that if $k=0$ then there is 1 subset; if $k=1$ then there are n subsets; if $k=n$ then there is 1 subset. The connection with the permutation rule is that there are $n!/(n-k)!$ permutations and each permutation has $k!$ permutations.

How many subsets of 7 people can be taken from a set of 12 persons? ${}_{12}C_7 = \binom{12}{7} = \frac{12!}{7!(12-7)!} = 792$

If you are dealt five cards, what is the probability of getting a “full-house” hand containing three kings and two aces (KKKAA)?

$$P(F) = \frac{\binom{4}{3}\binom{4}{2}}{\binom{52}{5}}$$

Distinguishable permutations. The number of *unordered* arrangements (distinguishable permutations) of a set of $|S| = n$ items in which n_1 are of one type, n_2 are of another type, etc., is

$$\binom{n}{n_1, n_2, \dots, n_k} = \frac{n!}{n_1!n_2!\dots n_k!}$$

How many ordered arrangements are there of the letters in the word PHILIPPINES? There are $n=11$ objects. $|P| = n_1 = 3$; $|H| = n_2 = 1$; $|I| = n_3 = 3$; $|L| = n_4 = 1$; $|N| = n_5 = 1$; $|E| = n_6 = 1$; $|S| = n_7 = 1$.

$$\binom{11}{n_1, n_2, \dots, n_k} = \frac{11!}{3!1!3!1!1!1!1!} = 1,108,800$$

How many ways can a research pool of 15 subjects be divided into three equally sized test groups?

$$\binom{n}{n_1, n_2, \dots, n_k} = \frac{15!}{5!5!5!} = 756,756$$

1.2 Discrete Distributions

1.2.1 Binomial

If X is the count of successful events in n identical and independent Bernoulli trials of success probability p , then X is a random variable with a binomial distribution $X \sim b(n, p)$ with mean $\mu = np$ and variance $\sigma^2 = np(1 - p)$. The probability of $X = x$ successes in n trials is

$$P(X = x) = \frac{n!}{x!(n - x)!} p^x (1 - p)^{n - x}.$$

What is the probability 2 out of 10 coin flips are heads if the probability of heads is 0.3?

Function `dbinom()` calculates the binomial probability.

```
dbinom(x = 2, size = 10, prob = 0.3)
```


```
## [1] 0.2334744
```

A simulation of $n = 10,000$ random samples of size 10 gives a similar result. `rbinom()` generates a random sample of numbers from the binomial distribution.

```
library(tidyverse)

data.frame(cnt = rbinom(n = 10000, size = 10, prob = 0.3)) %>%
  count(cnt) %>%
  ungroup() %>%
  mutate(pct = n / sum(n),
         X_eq_x = cnt == 2) %>%
  ggplot(aes(x = as.factor(cnt), y = n, fill = X_eq_x, label = pct)) +
  geom_col(alpha = 0.8) +
  scale_fill_manual(values = c(my_colors$grey, my_colors$red)) +
  geom_label(aes(label = round(pct, 2)), size = 3, alpha = .6) +
  theme_minimal() +
  theme(legend.position = "none") +
```

```
labs(title = "Binomial Distribution",
     subtitle = paste0("P(X=2) successes in 10 trials when p = 0.3 is ", round(dbinom(2, 10, 0.3), 4)),
     x = "Successes",
     y = "Count",
     caption = "Simulation from n = 10,000 binomial random samples.")
```



What is the probability of ≤ 2 heads in 10 coin flips where probability of heads is 0.3?

The cumulative probability is the sum of the first three bars in the simulation above. Function `pbinom()` calculates the *cumulative* binomial probability.

```
pbinom(q = 2, size = 10, prob = 0.3, lower.tail = TRUE)
```

```
## [1] 0.3827828
```

What is the expected number of heads in 25 coin flips if the probability of heads is 0.3?

The expected value, $\mu = np$, is 7.5. Here's an empirical test from 10,000 samples.

```
mean(rbinom(n = 10000, size = 25, prob = .3))
```

```
## [1] 7.5102
```

The variance, $\sigma^2 = np(1 - p)$, is 5.25. Here's an empirical test.

```
var(rbinom(n = 10000, size = 25, prob = .3))
```

```
## [1] 5.20409
```

Suppose X and Y are independent random variables distributed $X \sim b(10, .6)$ and $Y \sim b(10, .7)$. What is the probability that either variable is ≤ 4 ?

Let $P(A) = P(X \leq 4)$ and $P(B) = P(Y \leq 4)$. Then $P(A|B) = P(A) + P(B) - P(AB)$, and because the events are independent, $P(AB) = P(A)P(B)$.

```
p_a <- pbinom(q = 4, size = 10, prob = 0.6, lower.tail = TRUE)
p_b <- pbinom(q = 4, size = 10, prob = 0.7, lower.tail = TRUE)
p_a + p_b - (p_a * p_b)
```

```
## [1] 0.2057164
```

Here's an empirical test.

```
df <- data.frame(
  x = rbinom(10000, 10, 0.6),
  y = rbinom(10000, 10, 0.7)
)
mean(if_else(df$x <= 4 | df$y <= 4, 1, 0))
```

```
## [1] 0.2036
```

1.2.2 Negative-Binomial

If X is the count of trials required to reach a target number r of successful events in identical and independent Bernoulli trials of success probability p , then X is a random variable with a negative-binomial distribution $X \sim nb(r, p)$ with mean $\mu = r/p$ and variance $\sigma^2 = r(1-p)/p^2$. The probability of $X = x$ trials prior to r successes is

$$P(X = x) = \binom{x-1}{r-1} p^r (1-p)^{x-r}.$$

An oil company has a $p = 0.20$ chance of striking oil when drilling a well. What is the probability the company drills $x = 7$ wells to strike oil $r = 3$ times?

$$P(X = 7) = \binom{7-1}{3-1} (0.2)^3 (1-0.2)^{(7-3)} = 0.049.$$

Function `dnbinom()` calculates the negative-binomial probability. Parameter `x` equals the number of failures, $x - r$.

```
dnbinom(x = 4, size = 3, prob = 0.2)
```


```
## [1] 0.049152
```

Here is a simulation of $n = 10,000$ random samples. `rnbinom()` generates a random sample of numbers from the negative-binomial distribution.

```
data.frame(cnt = rnbinom(n = 10000, size = 3, prob = 0.2)) %>%
  count(cnt) %>%
  ungroup() %>%
  mutate(pct = n / sum(n),
         X_eq_x = cnt == 7-3,
         cnt = cnt + 3) %>%
  filter(cnt < 15) %>%
  ggplot(aes(x = as.factor(cnt), y = n, fill = X_eq_x, label = pct)) +
  geom_col(alpha = 0.8) +
  scale_fill_manual(values = c(my_colors$grey, my_colors$red)) +
  geom_label(aes(label = round(pct, 2)), size = 3, alpha = .6, check_overlap = TRUE) +
  theme_minimal() +
  theme(legend.position = "none") +
  labs(title = "Negative-Binomial Distribution",
       subtitle = paste0("P(X=7) trials to reach 3 successes when p = 0.2 is ", round(
         x = "Trials",
         y = "Count",
         caption = "Simulation from n = 10,000 negative-binomial random samples.")
```

Negative-Binomial Distribution

$P(X=7)$ trials to reach 3 successes when $p = 0.2$ is 0.0492.



Simulation from $n = 10,000$ negative-binomial random samples.

1.2.3 Geometric

If X is the count of independent Bernoulli trials of success probability p required to achieve the first successful trial, then X is a random variable with a geometric distribution $X \sim G(p)$ with mean $\mu = \frac{n}{p}$ and variance $\sigma^2 = \frac{(1-p)}{p^2}$. The probability of $X = n$ trials is

$$f(X = n) = p(1 - p)^{n-1}.$$

The probability of $X \leq n$ trials is

$$F(X = n) = 1 - (1 - p)^n.$$


Example. A sports marketer randomly selects persons on the street until he encounters someone who attended a game last season. What is the probability the marketer encounters $x = 3$ people who did not attend a game before the first success if $p = 0.20$ of the population attended a game?

Function `pgeom()` calculates the geometric distribution probability.

```
dgeom(x = 3, prob = 0.20)
```

```
## [1] 0.1024
```

```
data.frame(cnt = rgeom(n = 10000, prob = 0.20)) %>%
  count(cnt) %>%
  top_n(n = 15, wt = n) %>%
  ungroup() %>%
  mutate(pct = round(n / sum(n), 2),
         X_eq_x = cnt == 3) %>%
  ggplot(aes(x = as.factor(cnt), y = n, fill = X_eq_x, label = pct)) +
  geom_col(alpha = 0.8) +
  scale_fill_manual(values = c(my_colors$grey, my_colors$red)) +
  geom_text(size = 3) +
  theme_minimal() +
  theme(legend.position = "none") +
  labs(title = "Distribution of trials prior to first success",
       subtitle = paste("P(X = 3) | X ~ G(.2) = ", round(dgeom(3, .2), 4)),
       x = "Unsuccessful trials",
       y = "Count",
       caption = "simulation of n = 10,000 samples from geometric dist.")
```



1.3 Continuous Distributions

1.3.1 Normal

Random variable X is distributed $X \sim N(\mu, \sigma^2)$ if

$$f(X) = \frac{1}{\sigma\sqrt{2\pi}} e^{-.5\left(\frac{x-\mu}{\sigma}\right)^2}$$

Example

IQ scores are distributed $X \sim N(100, 16^2)$. What is the probability a randomly selected person's IQ is < 90 ?

```
my_mean = 100
my_sd = 16
my_x = 90
# exact
pnorm(q = my_x, mean = my_mean, sd = my_sd, lower.tail = TRUE)
```

```
## [1] 0.2659855
```




```
# simulated
mean(rnorm(n = 10000, mean = my_mean, sd = my_sd) <= my_x)
```

```
## [1] 0.2653
```

```
library(dplyr)
library(ggplot2)

data.frame(x = 0:1500 / 10,
           prob = pnorm(q = 0:1500 / 10,
                        mean = my_mean,
                        sd = my_sd,
                        lower.tail = TRUE)) %>%
  mutate(cdf = ifelse(x > 0 & x <= my_x, prob, 0)) %>%
  ggplot() +
    geom_line(aes(x = x, y = prob)) +
    geom_area(aes(x = x, y = cdf), alpha = 0.3) +
    labs(title = bquote('X~N(' ~mu==.(my_mean)~', '~sigma^{2}==.(my_sd)^{2}~')'),
         subtitle = bquote('P(X<=' ~.(my_x)~') when mean is' ~.(my_mean)~' and variance is' ~.(my_sd)^{2}~'),
         x = "x",
         y = "Probability")
```



1.3.2 Example

IQ scores are distributed $X \sim N(100, 16^2)$. What is the probability a randomly selected person's IQ is >140 ?

```
my_mean = 100
my_sd = 16
my_x = 140
# exact
pnorm(q = my_x, mean = my_mean, sd = my_sd, lower.tail = FALSE)
```

```
## [1] 0.006209665
```

```
# simulated
mean(rnorm(n = 10000, mean = my_mean, sd = my_sd) > my_x)
```


```
## [1] 0.0061
```

```
library(dplyr)
library(ggplot2)

data.frame(x = 0:1500 / 10,
           prob = pnorm(q = 0:1500 / 10,
                        mean = my_mean,
                        sd = my_sd,
                        lower.tail = TRUE)) %>%
  mutate(cdf = ifelse(x > my_x & x < 1000, prob, 0)) %>%
  ggplot() +
  geom_line(aes(x = x, y = prob)) +
  geom_area(aes(x = x, y = cdf), alpha = 0.3) +
  labs(title = bquote('X~N(' ~ mu == .(my_mean) ~ ', ' ~ sigma^{2} == .(my_sd)^{2} ~ ')'),
       subtitle = bquote('P(X<= ' ~ .(my_x) ~ ') when mean is ' ~ .(my_mean) ~ ' and variance is
       x = "x",
       y = "Probability")
```

$$X \sim N(\mu = 100, \sigma^2 = 16^2)$$

$P(X \leq 140)$ when mean is 100 and variance is 16^2 .



1.3.3 Example

IQ scores are distributed $X \sim N(100, 16^2)$. What is the probability a randomly selected person's IQ is between 92 and 114?

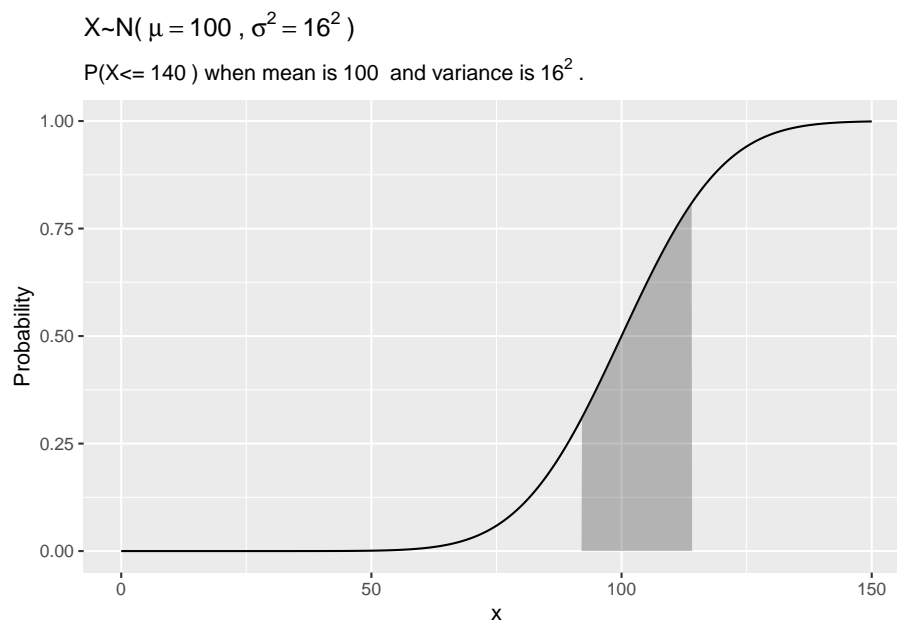
```
my_mean = 100
my_sd = 16
my_x_l = 92
my_x_h = 114
# exact
pnorm(q = my_x_h, mean = my_mean, sd = my_sd, lower.tail = TRUE) -
  pnorm(q = my_x_l, mean = my_mean, sd = my_sd, lower.tail = TRUE)
```

```
## [1] 0.5006755
```

```
library(dplyr)
library(ggplot2)

data.frame(x = 0:1500 / 10,
           prob = pnorm(q = 0:1500 / 10,
                        mean = my_mean,
                        sd = my_sd,
                        lower.tail = TRUE)) %>%
```

```
mutate(cdf = ifelse(x > my_x_l & x <= my_x_h, prob, 0)) %>%
ggplot() +
  geom_line(aes(x = x, y = prob)) +
  geom_area(aes(x = x, y = cdf), alpha = 0.3) +
  labs(title = bquote('X~N(' ~ mu == .(my_mean) ~ ', ' ~ sigma^{2} == .(my_sd)^{2} ~ ')'),
        subtitle = bquote('P(X <= ' ~ .(my_x) ~ ') when mean is ' ~ .(my_mean) ~ ' and variance is
x = "x",
y = "Probability")
```



1.3.4 Example

Class scores are distributed $X \sim N(70, 10^2)$. If the instructor wants to give A's to ≥ 85 th percentile and B's to 75th-85th percentile, what are the cutoffs?

```
my_mean = 70
my_sd = 10
my_pct_l = .75
my_pct_h = .85

qnorm(p = my_pct_l, mean = my_mean, sd = my_sd, lower.tail = TRUE)
```

```
## [1] 76.7449
```

```
qnorm(p = my_pct_h, mean = my_mean, sd = my_sd, lower.tail = TRUE)
```


```
## [1] 80.36433
```

```
library(dplyr)
library(ggplot2)

data.frame(x = 0:1000 / 10,
           prob = pnorm(q = 0:1000 / 10,
                        mean = my_mean,
                        sd = my_sd,
                        lower.tail = TRUE)) %>%
  mutate(cdf = ifelse(prob > my_pct_l & prob <= my_pct_h, prob, 0)) %>%
  ggplot() +
  geom_line(aes(x = x, y = prob)) +
  geom_area(aes(x = x, y = cdf), alpha = 0.3) +
  labs(title = bquote('X~N(' ~ mu == .(my_mean) ~ ', ' ~ sigma^{2} == .(my_sd)^{2} ~ ')'),
       subtitle = bquote('P(X<=x) = [' ~ .(my_pct_l) ~ ', ' ~ .(my_pct_h) ~ ']' when mean is ' ~ .(my_mean) ~ '
       x = "x",
       y = "Probability")
```

$X \sim N(\mu = 70, \sigma^2 = 10^2)$

$P(X \leq x) = [0.75, 0.85]$ when mean is 70 and variance is 10^2 .



1.3.5 Normal Approximation to Binomial

The CLT implies that certain distributions can be approximated by the normal distribution.

The binomial distribution $X \sim B(n, p)$ is approximately normal with mean $\mu = np$ and variance $\sigma^2 = np(1-p)$. The approximation is useful when the expected number of successes and failures is at least 5: $np \geq 5$ and $n(1-p) \geq 5$.

1.3.6 Example

A measure requires $p \geq 50\%$ popular to pass. A sample of $n=1,000$ yields $x=460$ approvals. What is the probability that the overall population approves, $P(X) > 0.5$?

```
my_x = 460
my_p = 0.50
my_n = 1000

my_mean = my_p * my_n
my_sd = round(sqrt(my_n * my_p * (1 - my_p)), 1)

# Exact binomial
pbinom(q = my_x, size = my_n, prob = my_p, lower.tail = TRUE)
```

```
## [1] 0.006222073
```

```
# Normal approximation
pnorm(q = my_x, mean = my_p * my_n, sd = sqrt(my_n * my_p * (1 - my_p)), lower.tail = FALSE)
```

```
## [1] 0.005706018
```


```
library(dplyr)
library(ggplot2)
library(tidyr)

data.frame(x = 400:600,
           Normal = pnorm(q = 400:600,
                          mean = my_p * my_n,
                          sd = sqrt(my_n * my_p * (1 - my_p)),
                          lower.tail = TRUE),
           Binomial = pbinom(q = 400:600,
                             size = my_n,
                             prob = my_p,
```

```

      lower.tail = TRUE)) %>%
gather(key = "Distribution", value = "cdf", c(-x)) %>%
ggplot(aes(x = x, y = cdf, color = Distribution)) +
geom_line() +
labs(title = bquote('X~B(n=~.(my_n)~, p=~.(my_p)~)'), 'X~N(~mu=~.(my_mean)~, ~sigma^{2}=~.(my_sd)^2~)',
      subtitle = "Normal approximation to the binomial",
      x = "x",
      y = "Probability")

```



The Poisson distribution $x \sim P(\lambda)$ is approximately normal with mean $\mu = \lambda$ and variance $\sigma^2 = \lambda$, for large values of λ .

1.3.7 Example

*The annual number of earthquakes registering at least 2.5 on the Richter Scale and having an epicenter within 40 miles of downtown Memphis follows a Poisson distribution with mean $\lambda = 6.5$. What is the probability that at least $x \geq 9$ such earthquakes will strike next year?**

```

my_x = 9
my_lambda = 6.5
my_sd = round(sqrt(my_lambda), 2)

```

```
# Exact Poisson
ppois(q = my_x - 1, lambda = my_lambda, lower.tail = FALSE)
```


```
## [1] 0.208427
```

```
# Normal approximation
pnorm(q = my_x - 0.5, mean = my_lambda, sd = my_sd, lower.tail = FALSE)
```

```
## [1] 0.216428
```

```
library(dplyr)
library(ggplot2)
library(tidyr)

data.frame(x = 0:200 / 10,
           Normal = pnorm(q = 0:200 / 10,
                          mean = my_lambda,
                          sd = my_sd,
                          lower.tail = TRUE),
           Poisson = ppois(q = 0:200 / 10,
                           lambda = my_lambda,
                           lower.tail = TRUE)) %>%
gather(key = "Distribution", value = "cdf", c(-x)) %>%
ggplot(aes(x = x, y = cdf, color = Distribution)) +
geom_line() +
labs(title = bquote('X~P('~lambda~'='~.(my_lambda)~)'), 'X~N('~mu=~.(my_lambda)~',
      subtitle = "Normal approximation to the Poisson",
      x = "x",
      y = "Probability")
```

1.3.8 From Sample to Population

Suppose a person's blood pressure typically measures 160 ± 20 mm. If one takes $n=5$ blood pressure readings, what is the probability the average will be ≤ 150 ?

```
my_mu = 160
my_sigma = 20
my_n = 5
my_x = 150

my_se = round(my_sigma / sqrt(my_n), 1)

pnorm(q = my_x, mean = my_mu, sd = my_sigma / sqrt(my_n), lower.tail = TRUE)
```

```
## [1] 0.1317762
```

```
library(dplyr)
library(ggplot2)

data.frame(x = 1000:2000 / 10,
           prob = pnorm(q = 1000:2000 / 10,
                        mean = my_mu,
                        sd = my_sigma / sqrt(my_n),
```


```

      lower.tail = TRUE)) %>%
  mutate(cdf = ifelse(x > 0 & x <= my_x, prob, 0)) %>%
  ggplot() +
    geom_line(aes(x = x, y = prob)) +
    geom_area(aes(x = x, y = cdf), alpha = 0.3) +
    labs(title = bquote('X~N(' ~ mu == .(my_mu) ~ ', ' ~ sigma^{2} == .(my_se)^{2} ~ ')'),
         subtitle = bquote('P(X <= ' ~ .(my_x) ~ ') when mean is ' ~ .(my_mu) ~ ' and variance is ' ~ .(my_se)^2 ~ ')'),
         x = "x",
         y = "Probability")

```

$X \sim N(\mu = 160, \sigma^2 = 8.9^2)$

$P(X \leq 150)$ when mean is 160 and variance is $\sigma/\sqrt{n} \ 8.9^2$.



```

knitr::include_app("https://mpfoley73.shinyapps.io/shiny_dist/",
  height = "600px")

```

Chapter 2

Inference

Chapter 3

Experiments

Some *significant* applications are demonstrated in this chapter.

3.1 Example one

3.2 Example two

Chapter 4

Regression

Chapter 5

Classification

Chapter 6

Regularization

Chapter 7

Non-linear Models

Linear methods can model nonlinear relationships by including polynomial terms, interaction effects, and variable transformations. However, it is often difficult to identify how to formulate the model. Nonlinear models may be preferable because you do not need to know the exact form of the nonlinearity prior to model training.

7.1 Splines

A regression spline fits a piecewise polynomial to the range of X partitioned by *knots* (K knots produce $K + 1$ piecewise polynomials) **James et al** (James et al., 2013). The polynomials can be of any degree d , but are usually in the range $[0, 3]$, most commonly 3 (a cubic spline). To avoid discontinuities in the fit, a degree- d spline is constrained to have continuity in derivatives up to degree $d-1$ at each knot.

A cubic spline fit to a data set with K knots, performs least squares regression with an intercept and $3 + K$ predictors, of the form

$$y_i = \beta_0 + \beta_1 X + \beta_2 X^2 + \beta_3 X^3 + \beta_4 h(X, \xi_1) + \beta_5 h(X, \xi_2) + \cdots + \beta_{K+3} h(X, \xi_K)$$

where ξ_1, \dots, ξ_K are the knots are truncated power basis functions $h(X, \xi) = (X - \xi)^3$ if $X > \xi$, else 0.

Splines can have high variance at the outer range of the predictors. A **natural spline** is a regression spline additionally constrained to be linear at the boundaries.

How many knots should there be, and Where should the knots be placed? It is common to place knots in a uniform fashion, with equal numbers of points

between each knot. The number of knots is typically chosen by trial and error using cross-validation to minimize the RSS. The number of knots is usually expressed in terms of degrees of freedom. A cubic spline will have $K + 3 + 1$ degrees of freedom. A natural spline has $K + 3 + 1 - 5$ degrees of freedom due to the constraints at the endpoints.

A further constraint can be added to reduce overfitting by enforcing smoothness in the spline. Instead of minimizing the loss function $\sum (y - g(x))^2$ where $g(x)$ is a natural spline, minimize a loss function with an additional penalty for variability:

$$L = \sum (y_i - g(x_i))^2 + \lambda \int g''(t)^2 dt.$$

The function $g(x)$ that minimizes the loss function is a *natural cubic spline* with knots at each x_1, \dots, x_n . This is called a **smoothing spline**. The larger λ is, the greater the penalty on variation in the spline. In a smoothing spline, you do not optimize the number or location of the knots – there is a knot at each training observation. Instead, you optimize λ . One way to optimize λ is cross-validation to minimize RSS. Leave-one-out cross-validation (LOOCV) can be computed efficiently for smoothing splines.

7.2 MARS

Multivariate adaptive regression splines (MARS) is a non-parametric algorithm that creates a piecewise linear model to capture nonlinearities and interactions effects. The resulting model is a weighted sum of *basis* functions $B_i(X)$:

$$\hat{y} = \sum_{i=1}^k w_i B_i(x)$$

The basis functions are either a constant (for the intercept), a *hinge* function of the form $\max(0, x - x_0)$ or $\max(0, x_0 - x)$ (a more concise representation is $[\pm(x - x_0)]_+$), or products of two or more hinge functions (for interactions). MARS automatically selects which predictors to use and what predictor values to serve as the *knots* of the hinge functions.

MARS builds a model in two phases: the forward pass and the backward pass, similar to growing and pruning of tree models. MARS starts with a model consisting of just the intercept term equaling the mean of the response values. It then assesses every predictor to find a basis function pair consisting of opposing sides of a mirrored hinge function which produces the maximum improvement in the model error. MARS repeats the process until either it reaches a predefined limit of terms or the error improvement reaches a predefined limit.

MARS generalizes the model by removing terms according to the generalized cross validation (GCV) criterion. GCV is a form of regularization: it trades off goodness-of-fit against model complexity.

The `earth::earth()` function (documentation) performs the MARS algorithm (*the term “MARS” is trademarked, so open-source implementations use “Earth” instead*). The caret implementation tunes two parameters: `nprune` and `degree`. `nprune` is the maximum number of terms in the pruned model. `degree` is the maximum degree of interaction (default is 1 (no interactions)). However, there are other hyperparameters in the model that may improve performance, including `minspan` which regulates the number of knots in the predictors.

Here is an example using the Ames housing data set (following this tutorial).


```
library(tidyverse)
library(earth)
library(caret)

# set up
ames <- AmesHousing::make_ames()
set.seed(12345)
idx <- createDataPartition(ames$Sale_Price, p = 0.80, list = FALSE)
ames_train <- ames[idx, ] %>% as.data.frame()
ames_test  <- ames[-idx, ]

m <- train(
  x = subset(ames_train, select = -Sale_Price),
  y = ames_train$Sale_Price,
  method = "earth",
  metric = "RMSE",
  minspan = -15,
  trControl = trainControl(method = "cv", number = 10),
  tuneGrid = expand.grid(
    degree = 1:3,
    nprune = seq(2, 100, length.out = 10) %>% floor()
  )
)
```

The model plot shows the best tuning parameter combination.

```
plot(m, main = "MARS Parameter Tuning")
```



```
m$bestTune
```

```
##      nprune degree
## 25      45      3
```

How does this model perform against the holdout data?

```
caret::postResample(
  pred = log(predict(m, newdata = ames_test)),
  obs = log(ames_test$Sale_Price)
)
```

```
##      RMSE  Rsquared      MAE
## 0.16515620 0.85470300 0.09319503
```

7.3 GAM

Generalized additive models (GAM) allow for non-linear relationships between each feature and the response by replacing each linear component $\beta_j x_{ij}$ with a nonlinear function $f_j(x_{ij})$. The GAM model is of the form

$$y_i = \beta_0 + \sum f_j(x_{ij}) + \epsilon_i.$$

It is called an additive model because we calculate a separate f_j for each X_j , and then add together all of their contributions.

The advantage of GAMs is that they automatically model non-linear relationships so you do not need to manually try out many different transformations on each variable individually. And because the model is additive, you can still examine the effect of each X_j on Y individually while holding all of the other variables fixed. The main limitation of GAMs is that the model is restricted to be additive, so important interactions can be missed unless you explicitly add them.

Chapter 8

Decision Trees

Decision trees, also known as classification and regression tree (CART) models, are tree-based methods for supervised machine learning. Simple *classification trees* and *regression trees* are easy to use and interpret, but are not competitive with the best machine learning methods. However, they form the foundation for **bagged trees**, **random forests**, and **boosted trees** models, which although less interpretable, are very accurate.

CART models segment the predictor space into K non-overlapping terminal nodes (leaves), A_1, A_2, \dots, A_K . Each node is described by a set of rules which can be used to predict new responses. The predicted value \hat{y} for each node is the mode (classification), or mean (regression).

CART models define the nodes through a *top-down greedy* process called *recursive binary splitting*. The process is *top-down* because it begins at the top of the tree with all observations in a single region and successively splits the predictor space. It is *greedy* because at each splitting step, the best split is made at that particular step without consideration to subsequent splits.

The best split is the predictor variable and cutpoint that minimizes a cost function. For a regression tree, the most common cost function is the sum of squared residuals,

$$RSS = \sum_{k=1}^K \sum_{i \in A_k} (y_i - \hat{y}_{A_k})^2.$$

For a classification tree, the most common cost functions are the Gini index,

$$G = \sum_{c=1}^C \hat{p}_{kc}(1 - \hat{p}_{kc}),$$

or the entropy

$$D = - \sum_{c=1}^C \hat{p}_{kc} \log \hat{p}_{kc}$$

where \hat{p}_{kc} is the proportion of training observations in node k node that are class c . A completely *pure* node in a binary tree will have $\hat{p} \in [0, 1]$ and $G = D = 0$. A completely impure node in a binary tree will have $\hat{p} = 0.5$ and $G = 0.5^2 \cdot 2 = 0.25$ and $D = -(0.5 \log(0.5)) \cdot 2 = 0.69$.

CART repeats the splitting process for each of the child nodes until a *stopping criterion* is satisfied, usually when no node size surpasses a predefined maximum, or continued splitting does not improve the model significantly. CART may also impose a minimum number of observations in each node.

The resulting tree likely over-fits the training data and therefore does not generalize well to test data, so CART *prunes* the tree, minimizing the cross-validated prediction error. Rather than cross-validating every possible subtree to find the one with minimum error, CART uses *cost-complexity pruning*. Cost-complexity is the tradeoff between error (cost) and tree size (complexity) where the tradeoff is quantified with cost-complexity parameter c_p . In the equation below, the cost complexity of the tree $R_{c_p}(T)$ is the sum of its risk (error) plus a “cost complexity” factor c_p multiple of the tree size $|T|$.

$$R_{c_p}(T) = R(T) + c_p |T|$$

c_p can take on any value from $[0, \infty]$, but it turns out there is an optimal tree for *ranges* of c_p values, so there are only a finite set of *interesting* values for c_p (James et al., 2013) (Therneau and Atkinson, 2019) (Kuhn and Johnson, 2016). A parametric algorithm identifies the interesting c_p values and their associated pruned trees, T_{c_p} .

CART uses cross-validation to determine which c_p is optimal.

8.1 Classification Tree

A simple classification tree is rarely performed on its own; the bagged, random forest, and gradient boosting methods build on this logic. However, it is good to start here to build understanding. I’ll learn by example. Using the `ISLR::OJ` data set, I will predict which brand of orange juice, Citrus Hill (CH) or Minute Maid = (MM), customers `Purchase` using from the 17 feature variables. Load the libraries and data.

```

library(ISLR) # OJ dataset
library(rpart) # classification and regression trees
library(caret) # modeling workflow
library(rpart.plot) # better formatted plots than the ones in rpart
library(plotROC) # ROC curves
library(ROCR)
library(tidyverse)
library(skimr) # neat alternative to glance & summary

oj_dat <- OJ
skim_with(numeric = list(p0 = NULL, p25 = NULL, p50 = NULL, p75 = NULL,
                        p100 = NULL, hist = NULL))

skim(oj_dat)

```

```

## Skim summary statistics
## n obs: 1070
## n variables: 18
##
## -- Variable type:factor -----
## variable missing complete      n n_unique      top_counts ordered
## Purchase          0      1070 1070          2 CH: 653, MM: 417, NA: 0  FALSE
##   Store7           0      1070 1070          2 No: 714, Yes: 356, NA: 0  FALSE
##
## -- Variable type:numeric -----
##      variable missing complete      n      mean      sd
##      DiscCH          0      1070 1070      0.052  0.12
##      DiscMM          0      1070 1070      0.12   0.21
##   ListPriceDiff      0      1070 1070      0.22   0.11
##      LoyalCH          0      1070 1070      0.57   0.31
##      PctDiscCH        0      1070 1070      0.027  0.062
##      PctDiscMM        0      1070 1070      0.059  0.1
##      PriceCH          0      1070 1070      1.87   0.1
##      PriceDiff        0      1070 1070      0.15   0.27
##      PriceMM          0      1070 1070      2.09   0.13
##   SalePriceCH        0      1070 1070      1.82   0.14
##   SalePriceMM        0      1070 1070      1.96   0.25
##      SpecialCH        0      1070 1070      0.15   0.35
##      SpecialMM        0      1070 1070      0.16   0.37
##      STORE           0      1070 1070      1.63   1.43
##      StoreID          0      1070 1070      3.96   2.31
##   WeekofPurchase      0      1070 1070     254.38  15.56

```

I'll split `oj_dat` ($n = 1,070$) into `oj_train` (80%, $n = 857$) and `oj_test` (20%, $n = 213$). I'll fit a simple decision tree with `oj_train`, then later a bagged tree,

a random forest, and a gradient boosting tree. I'll compare their predictive performance with `oj_test`.

```
set.seed(12345)
partition <- createDataPartition(y = oj_dat$Purchase, p = 0.8, list = FALSE)
oj_train <- oj_dat[partition, ]
oj_test <- oj_dat[-partition, ]
```

Function `rpart::rpart()` builds a full tree, minimizing the Gini index G by default (`parms = list(split = "gini")`), until the stopping criterion is satisfied. The default stopping criterion is

- only attempt a split if the current node has at least `minsplit = 20` observations,
- only accept a split if each of the two resulting nodes have at least `minbucket = round(minsplit/3)` observations, and
- only accept a split if the resulting overall fit improves by `cp = 0.01` (i.e., $\Delta G \leq 0.01$).


```
set.seed(123)
oj_model_1 <- rpart(
  formula = Purchase ~ .,
  data = oj_train,
  method = "class" # "class" for classification, "anova" for regression
)
print(oj_model_1)
```

```
## n= 857
##
## node), split, n, loss, yval, (yprob)
##      * denotes terminal node
##
##  1) root 857 334 CH (0.61026838 0.38973162)
##    2) LoyalCH>=0.48285 537 94 CH (0.82495345 0.17504655)
##      4) LoyalCH>=0.7648795 271 13 CH (0.95202952 0.04797048) *
##      5) LoyalCH< 0.7648795 266 81 CH (0.69548872 0.30451128)
##        10) PriceDiff>=-0.165 226 50 CH (0.77876106 0.22123894) *
##        11) PriceDiff< -0.165 40 9 MM (0.22500000 0.77500000) *
##    3) LoyalCH< 0.48285 320 80 MM (0.25000000 0.75000000)
##      6) LoyalCH>=0.2761415 146 58 MM (0.39726027 0.60273973)
##        12) SalePriceMM>=2.04 71 31 CH (0.56338028 0.43661972) *
##        13) SalePriceMM< 2.04 75 18 MM (0.24000000 0.76000000) *
##      7) LoyalCH< 0.2761415 174 22 MM (0.12643678 0.87356322) *
```

The output starts with the root node. The predicted class at the root is CH and this prediction produces 334 errors on the 857 observations for a success rate of 0.61026838 and an error rate of 0.38973162. The child nodes of node “x” are labeled 2x) and 2x+1), so the child nodes of 1) are 2) and 3), and the child nodes of 2) are 4) and 5). Terminal nodes are labeled with an asterisk (*).

Surprisingly, only 3 of the 17 features were used in the full tree: **LoyalCH** (Customer brand loyalty for CH), **PriceDiff** (relative price of MM over CH), and **SalePriceMM** (absolute price of MM). The first split is at **LoyalCH** = 0.48285. Here is what the full (unpruned) tree looks like.

```
rpart.plot(oj_model_1, yesno = TRUE)
```



The boxes show the node classification (based on mode), the proportion of observations that are *not* CH, and the proportion of observations included in the node.

rpart() not only grew the full tree, it identified the set of cost complexity parameters, and measured the model performance of each corresponding tree using cross-validation. **printcp()** displays the candidate c_p values. You can use this table to decide how to prune the tree.

```
printcp(oj_model_1)
```

```
##
```

```
## Classification tree:
## rpart(formula = Purchase ~ ., data = oj_train, method = "class")
##
## Variables actually used in tree construction:
## [1] LoyalCH      PriceDiff    SalePriceMM
##
## Root node error: 334/857 = 0.38973
##
## n= 857
##
##          CP nsplit rel error  xerror    xstd
## 1 0.479042     0  1.00000 1.00000 0.042745
## 2 0.032934     1  0.52096 0.54192 0.035775
## 3 0.013473     3  0.45509 0.47006 0.033905
## 4 0.010000     5  0.42814 0.46407 0.033736
```

There are 4 c_p values in this model. The model with the smallest complexity parameter allows the most splits (`nsplit`). The highest complexity parameter corresponds to a tree with just a root node. `rel error` is the error rate relative to the root node. The root node absolute error is 0.38973162 (the proportion of MM), so its `rel error` is $0.38973162/0.38973162 = 1.0$. That means the absolute error of the full tree (at $CP = 0.01$) is $0.42814 * 0.38973162 = 0.1669$. You can verify that by calculating the error rate of the predicted values:

```
data.frame(pred = predict(oj_model_1, newdata = oj_train, type = "class")) %>%
  mutate(obs = oj_train$Purchase,
         err = if_else(pred != obs, 1, 0)) %>%
  summarize(mean_err = mean(err))

##      mean_err
## 1 0.1668611
```

Finishing the CP table tour, `xerror` is the relative cross-validated error rate and `xstd` is its standard error. If you want the lowest possible error, then prune to the tree with the smallest relative CV error (`xerror`) ($c_p = 0.01$, CV error = 0.1809). If you want to balance predictive power with simplicity, prune to the smallest tree within 1 SE of the one with the smallest relative error. The CP table is not super-helpful for finding that tree. I'll add a column to find it.

```
oj_model_1$cptable %>%
  data.frame() %>%
  mutate(min_xerror_idx = which.min(oj_model_1$cptable[, "xerror"]),
         rownum = row_number(),
         xerror_cap = oj_model_1$cptable[min_xerror_idx, "xerror"] +
```



```


oj_model_1$cptable[min_xerror_idx, "xstd"],
eval = case_when(rownum == min_xerror_idx ~ "min xerror",
  xerror < xerror_cap ~ "under cap",
  TRUE ~ "") %>%
select(-rownum, -min_xerror_idx)

```

##	CP	nsplit	rel.error	xerror	xstd	xerror_cap	eval
## 1	0.47904192	0	1.0000000	1.0000000	0.04274518	0.4978082	
## 2	0.03293413	1	0.5209581	0.5419162	0.03577468	0.4978082	
## 3	0.01347305	3	0.4550898	0.4700599	0.03390486	0.4978082	under cap
## 4	0.01000000	5	0.4281437	0.4640719	0.03373631	0.4978082	min xerror

The simplest tree using the 1-SE rule is $\text{sc}_p = 0.01347305$, CV error = 0.1832). Fortunately, `plotcp()` presents a nice graphical representation of the relationship between `xerror` and `cp`.

```
plotcp(oj_model_1, upper = "splits")
```



The dashed line is set at the minimum `xerror` + `xstd`. The top axis shows the number of splits in the tree. I'm not sure why the CP values are not the same as in the table (they are close, but not the same). The figure suggests I should prune to 5 or 3 splits. I see this curve never really hits a minimum - it is still decreasing at 5 splits. The default tuning parameter value `cp = 0.01` may be too small, so I'll set it to `cp = 0.001` and start over.

```

set.seed(123)
oj_model_1b <- rpart(
  formula = Purchase ~ .,
  data = oj_train,
  method = "class",
  cp = 0.001
)
print(oj_model_1b)

```

```


## n= 857
##
## node), split, n, loss, yval, (yprob)
##      * denotes terminal node
##
## 1) root 857 334 CH (0.61026838 0.38973162)
##    2) LoyalCH>=0.48285 537 94 CH (0.82495345 0.17504655)
##      4) LoyalCH>=0.7648795 271 13 CH (0.95202952 0.04797048) *
##      5) LoyalCH< 0.7648795 266 81 CH (0.69548872 0.30451128)
##        10) PriceDiff>=-0.165 226 50 CH (0.77876106 0.22123894)
##          20) ListPriceDiff>=0.255 115 11 CH (0.90434783 0.09565217) *
##          21) ListPriceDiff< 0.255 111 39 CH (0.64864865 0.35135135)
##            42) PriceMM>=2.155 19 2 CH (0.89473684 0.10526316) *
##            43) PriceMM< 2.155 92 37 CH (0.59782609 0.40217391)
##              86) DiscCH>=0.115 7 0 CH (1.00000000 0.00000000) *
##              87) DiscCH< 0.115 85 37 CH (0.56470588 0.43529412)
##                174) ListPriceDiff>=0.215 45 15 CH (0.66666667 0.33333333) *
##                175) ListPriceDiff< 0.215 40 18 MM (0.45000000 0.55000000)
##                  350) LoyalCH>=0.527571 28 13 CH (0.53571429 0.46428571)
##                    700) WeekofPurchase< 266.5 21 8 CH (0.61904762 0.38095238) *
##                    701) WeekofPurchase>=266.5 7 2 MM (0.28571429 0.71428571) *
##                  351) LoyalCH< 0.527571 12 3 MM (0.25000000 0.75000000) *
##                11) PriceDiff< -0.165 40 9 MM (0.22500000 0.77500000) *
##      3) LoyalCH< 0.48285 320 80 MM (0.25000000 0.75000000)
##        6) LoyalCH>=0.2761415 146 58 MM (0.39726027 0.60273973)
##          12) SalePriceMM>=2.04 71 31 CH (0.56338028 0.43661972)
##            24) LoyalCH< 0.303104 7 0 CH (1.00000000 0.00000000) *
##            25) LoyalCH>=0.303104 64 31 CH (0.51562500 0.48437500)
##              50) WeekofPurchase>=246.5 52 22 CH (0.57692308 0.42307692)
##                100) PriceCH< 1.94 35 11 CH (0.68571429 0.31428571)
##                  200) StoreID< 1.5 9 1 CH (0.88888889 0.11111111) *
##                  201) StoreID>=1.5 26 10 CH (0.61538462 0.38461538)
##                    402) LoyalCH< 0.410969 17 4 CH (0.76470588 0.23529412) *
##                    403) LoyalCH>=0.410969 9 3 MM (0.33333333 0.66666667) *
##                  101) PriceCH>=1.94 17 6 MM (0.35294118 0.64705882) *
##                51) WeekofPurchase< 246.5 12 3 MM (0.25000000 0.75000000) *

```

```
##      13) SalePriceMM< 2.04 75  18 MM (0.24000000 0.76000000)
##      26) SpecialCH>=0.5 14   6 CH (0.57142857 0.42857143) *
##      27) SpecialCH< 0.5 61  10 MM (0.16393443 0.83606557) *
##      7) LoyalCH< 0.2761415 174 22 MM (0.12643678 0.87356322)
##      14) LoyalCH>=0.035047 117 21 MM (0.17948718 0.82051282)
##      28) WeekofPurchase< 273.5 104 21 MM (0.20192308 0.79807692)
##      56) PriceCH>=1.875 20   9 MM (0.45000000 0.55000000)
##      112) WeekofPurchase>=252.5 12   5 CH (0.58333333 0.41666667) *
##      113) WeekofPurchase< 252.5 8    2 MM (0.25000000 0.75000000) *
##      57) PriceCH< 1.875 84  12 MM (0.14285714 0.85714286) *
##      29) WeekofPurchase>=273.5 13   0 MM (0.00000000 1.00000000) *
##      15) LoyalCH< 0.035047 57   1 MM (0.01754386 0.98245614) *
```

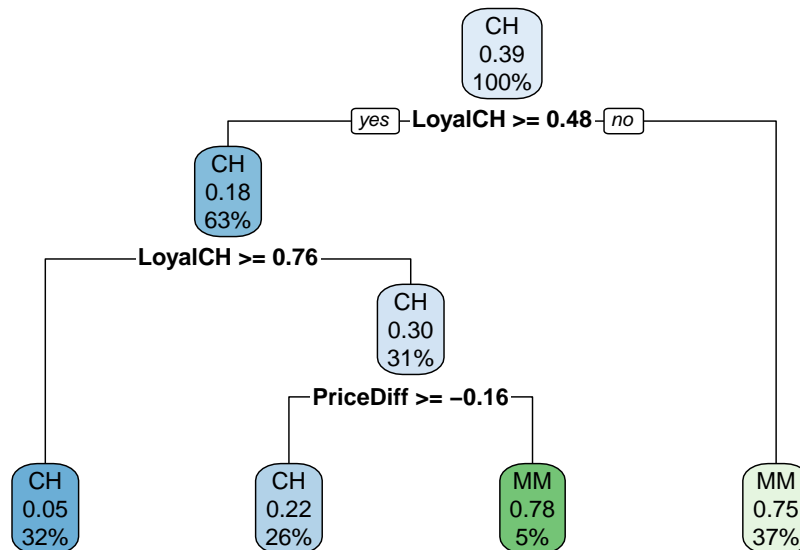
This is a much larger tree. Did I find a `cp` value that produces a local min?

```
plotcp(oj_model_1b, upper = "splits")
```



Yes, the min is at $CP = 0.011$ with 5 splits. The min + 1 SE is at $CP = 0.021$ with 3 splits. I'll prune the tree to 3 splits.

```
oj_model_1b_pruned <- prune(
  oj_model_1b,
  cp = oj_model_1b$cptable[oj_model_1b$cptable[, 2] == 3, "CP"]
)
rpart.plot(oj_model_1b_pruned, yesno = TRUE)
```



The most “important” indicator of `Purchase` appears to be `LoyalCH`. From the `rpart` vignette (page 12),

“An overall measure of variable importance is the sum of the goodness of split measures for each split for which it was the primary variable, plus goodness (adjusted agreement) for all splits in which it was a surrogate.”

Surrogates refer to alternative features for a node to handle missing data. For each split, CART evaluates a variety of alternative “surrogate” splits to use when the feature value for the primary split is NA. Surrogate splits are splits that produce results similar to the original split.

A variable’s importance is the sum of the improvement in the overall Gini (or RMSE) measure produced by the nodes in which it appears. Here is the variable importance for this model.

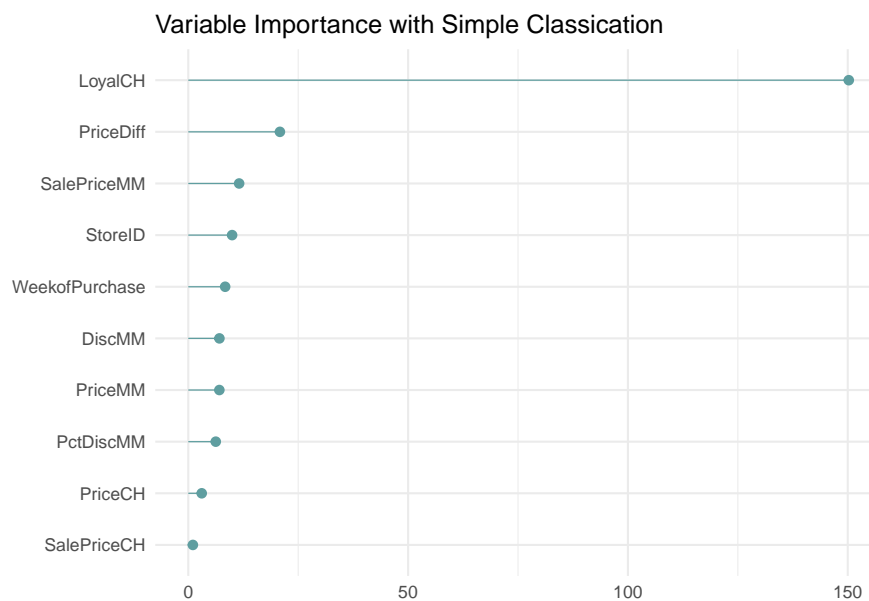
```
oj_model_1b_pruned$variable.importance
```

##	LoyalCH	PriceDiff	SalePriceMM	StoreID	WeekofPurchase
##	150.237336	20.843067	11.567443	9.965419	8.386282
##	DiscMM	PriceMM	PctDiscMM	PriceCH	SalePriceCH
##	7.081470	7.065493	6.252920	3.055594	1.042153

```

oj_model_1b_pruned$variable.importance %>%
  data.frame() %>%
  rownames_to_column(var = "Feature") %>%
  rename(Overall = '.') %>%
  ggplot(aes(x = fct_reorder(Feature, Overall), y = Overall)) +
  geom_pointrange(aes(ymin = 0, ymax = Overall), color = "cadetblue", size = .3) +
  theme_minimal() +
  coord_flip() +
  labs(x = "", y = "", title = "Variable Importance with Simple Classification")

```



LoyalCH is by far the most important variable, as expected from its position at the top of the tree, and one level down.

You can see how the surrogates appear in the model with the `summary()` function.

```
summary(oj_model_1b_pruned)
```

```

## Call:
## rpart(formula = Purchase ~ ., data = oj_train, method = "class",
##       cp = 0.001)
##   n= 857
##
##           CP nsplit rel error   xerror   xstd

```

```

## 1 0.47904192      0 1.0000000 1.0000000 0.04274518
## 2 0.03293413      1 0.5209581 0.5419162 0.03577468
## 3 0.01347305      3 0.4550898 0.4700599 0.03390486
##
## Variable importance
##      LoyalCH      PriceDiff      SalePriceMM      StoreID      WeekofPurchase
##           67           9           5           4           4
##      DiscMM      PriceMM      PctDiscMM      PriceCH
##           3           3           3           1
##
## Node number 1: 857 observations,      complexity param=0.4790419
##   predicted class=CH   expected loss=0.3897316   P(node) =1
##   class counts:      523      334
##   probabilities: 0.610 0.390
##   left son=2 (537 obs) right son=3 (320 obs)
##   Primary splits:
##     LoyalCH      < 0.48285   to the right, improve=132.56800, (0 missing)
##     StoreID      < 3.5       to the right, improve= 40.12097, (0 missing)
##     PriceDiff    < 0.015     to the right, improve= 24.26552, (0 missing)
##     ListPriceDiff < 0.255     to the right, improve= 22.79117, (0 missing)
##     SalePriceMM  < 1.84      to the right, improve= 20.16447, (0 missing)
##   Surrogate splits:
##     StoreID      < 3.5       to the right, agree=0.646, adj=0.053, (0 split)
##     PriceMM      < 1.89      to the right, agree=0.638, adj=0.031, (0 split)
##     WeekofPurchase < 229.5   to the right, agree=0.632, adj=0.016, (0 split)
##     DiscMM       < 0.77      to the left,  agree=0.629, adj=0.006, (0 split)
##     SalePriceMM  < 1.385     to the right, agree=0.629, adj=0.006, (0 split)
##
## Node number 2: 537 observations,      complexity param=0.03293413
##   predicted class=CH   expected loss=0.1750466   P(node) =0.6266044
##   class counts:      443      94
##   probabilities: 0.825 0.175
##   left son=4 (271 obs) right son=5 (266 obs)
##   Primary splits:
##     LoyalCH      < 0.7648795 to the right, improve=17.669310, (0 missing)
##     PriceDiff    < 0.015     to the right, improve=15.475200, (0 missing)
##     SalePriceMM  < 1.84      to the right, improve=13.951730, (0 missing)
##     ListPriceDiff < 0.255     to the right, improve=11.407560, (0 missing)
##     DiscMM       < 0.15      to the left,  improve= 7.795122, (0 missing)
##   Surrogate splits:
##     WeekofPurchase < 257.5   to the right, agree=0.594, adj=0.180, (0 split)
##     PriceCH       < 1.775     to the right, agree=0.590, adj=0.173, (0 split)
##     StoreID       < 3.5       to the right, agree=0.587, adj=0.165, (0 split)
##     PriceMM       < 2.04      to the right, agree=0.587, adj=0.165, (0 split)
##     SalePriceMM  < 2.04      to the right, agree=0.587, adj=0.165, (0 split)
##

```

```
## Node number 3: 320 observations
## predicted class=MM expected loss=0.25 P(node) =0.3733956
## class counts: 80 240
## probabilities: 0.250 0.750
##
## Node number 4: 271 observations
## predicted class=CH expected loss=0.04797048 P(node) =0.3162194
## class counts: 258 13
## probabilities: 0.952 0.048
##
## Node number 5: 266 observations, complexity param=0.03293413
## predicted class=CH expected loss=0.3045113 P(node) =0.3103851
## class counts: 185 81
## probabilities: 0.695 0.305
## left son=10 (226 obs) right son=11 (40 obs)
## Primary splits:
## PriceDiff < -0.165 to the right, improve=20.84307, (0 missing)
## ListPriceDiff < 0.235 to the right, improve=20.82404, (0 missing)
## SalePriceMM < 1.84 to the right, improve=16.80587, (0 missing)
## DiscMM < 0.15 to the left, improve=10.05120, (0 missing)
## PctDiscMM < 0.0729725 to the left, improve=10.05120, (0 missing)
## Surrogate splits:
## SalePriceMM < 1.585 to the right, agree=0.906, adj=0.375, (0 split)
## DiscMM < 0.57 to the left, agree=0.895, adj=0.300, (0 split)
## PctDiscMM < 0.264375 to the left, agree=0.895, adj=0.300, (0 split)
## WeekofPurchase < 274.5 to the left, agree=0.872, adj=0.150, (0 split)
## SalePriceCH < 2.075 to the left, agree=0.857, adj=0.050, (0 split)
##
## Node number 10: 226 observations
## predicted class=CH expected loss=0.2212389 P(node) =0.2637106
## class counts: 176 50
## probabilities: 0.779 0.221
##
## Node number 11: 40 observations
## predicted class=MM expected loss=0.225 P(node) =0.04667445
## class counts: 9 31
## probabilities: 0.225 0.775
```

The last step is to make predictions on the validation data set. For a classification tree, set argument `type = "class"`.

```
oj_model_1b_preds <- predict(oj_model_1b_pruned, oj_test, type = "class")
```

I'll evaluate the predictions and record the accuracy (correct classification percentage) for comparison to other models. Two ways to evaluate the model are the confusion matrix, and the ROC curve.

8.1.1 Confusion Matrix

Print the confusion matrix with `caret::confusionMatrix()` to see how well does this model performs against the test data set.

```
oj_model_1b_cm <- confusionMatrix(data = oj_model_1b_preds, reference = oj_test$Purchase,
oj_model_1b_cm
```

```
## Confusion Matrix and Statistics
##
##           Reference
## Prediction  CH  MM
##           CH 113  13
##           MM  17  70
##
##           Accuracy : 0.8592
##           95% CI : (0.8051, 0.9029)
##    No Information Rate : 0.6103
##    P-Value [Acc > NIR] : 1.265e-15
##
##           Kappa : 0.7064
##
##  McNemar's Test P-Value : 0.5839
##
##           Sensitivity : 0.8692
##           Specificity : 0.8434
##           Pos Pred Value : 0.8968
##           Neg Pred Value : 0.8046
##           Prevalence : 0.6103
##           Detection Rate : 0.5305
##    Detection Prevalence : 0.5915
##           Balanced Accuracy : 0.8563
##
##           'Positive' Class : CH
##
```

The confusion matrix is at the top. It also includes a lot of statistics. It's worth getting familiar with the stats. The model accuracy and 95% CI are calculated from the binomial test.

```
binom.test(x = 113 + 70, n = 213)
```

```
##
## Exact binomial test
```



```
##
## data: 113 + 70 and 213
## number of successes = 183, number of trials = 213, p-value <
## 2.2e-16
## alternative hypothesis: true probability of success is not equal to 0.5
## 95 percent confidence interval:
## 0.8050785 0.9029123
## sample estimates:
## probability of success
## 0.8591549
```

The “No Information Rate” (NIR) statistic is the class rate for the largest class. In this case CH is the largest class, so $NIR = 130/213 = 0.6103$. “P-Value [Acc > NIR]” is the binomial test that the model accuracy is significantly better than the NIR (i.e., significantly better than just always guessing CH).

```
binom.test(x = 113 + 70, n = 213, p = 130/213, alternative = "greater")
```

```
##
## Exact binomial test
##
## data: 113 + 70 and 213
## number of successes = 183, number of trials = 213, p-value =
## 1.265e-15
## alternative hypothesis: true probability of success is greater than 0.6103286
## 95 percent confidence interval:
## 0.8138446 1.0000000
## sample estimates:
## probability of success
## 0.8591549
```

The “Accuracy” statistic indicates the model predicts 0.8590 of the observations correctly. That’s good, but less impressive when you consider the prevalence of CH is 0.6103 - you could achieve 61% accuracy just by predicting CH every time. A measure that controls for the prevalence is Cohen’s kappa statistic. The kappa statistic is explained here. It compares the accuracy to the accuracy of a “random system”. It is defined as

$$\kappa = \frac{Acc - RA}{1 - RA}$$

where

$$RA = \frac{ActFalse \times PredFalse + ActTrue \times PredTrue}{Total \times Total}$$

is the hypotheical probability of a chance agreement. `ActFalse` will be the number of “MM” ($13 + 70 = 83$) and actual true will be the number of “CH” ($113 + 17 = 130$). The predicted counts are

```
table(oj_model_1b_preds)
```

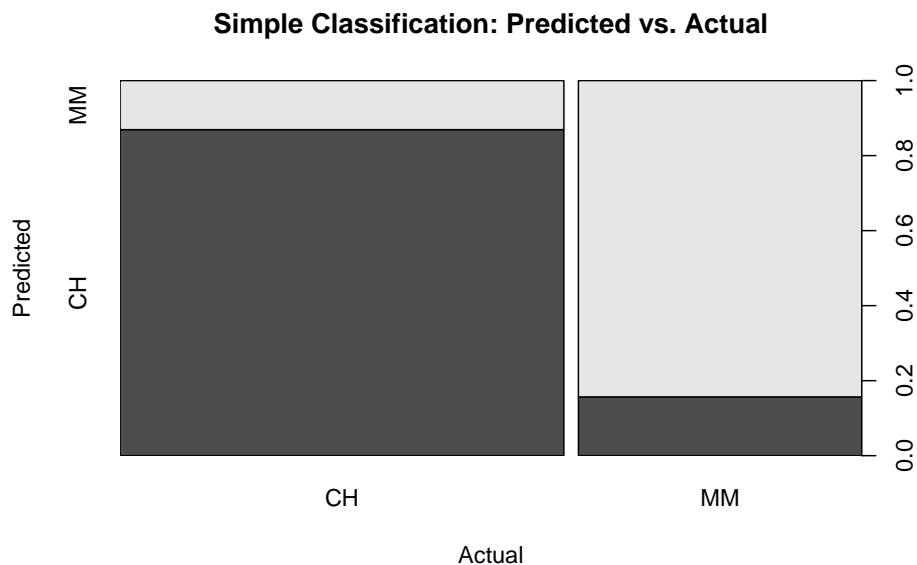
```
## oj_model_1b_preds
##  CH  MM
## 126  87
```

So, $RA = (83 * 87 + 130 * 126) / 213^2 = 0.5202$ and $\kappa = (0.8592 - 0.5202) / (1 - 0.5202) = 0.7064$. The kappa statistic varies from 0 to 1 where 0 means accurate predictions occur merely by chance, and 1 means the predictions are in perfect agreement with the observations. In this case, a kappa statistic of 0.7064 is “substantial”. See chart here.

The other measures from the `confusionMatrix()` output are various proportions and you can remind yourself of their definitions in the documentation with `?confusionMatrix`.

Visuals are almost always helpful. Here is a plot of the confusion matrix.

```
plot(oj_test$Purchase, oj_model_1b_preds,
     main = "Simple Classification: Predicted vs. Actual",
     xlab = "Actual",
     ylab = "Predicted")
```



By the way, how does the validation set accuracy ()

```
oj_model_1b_train_preds <- predict(oj_model_1b_pruned, oj_train, type = "class")
oj_model_1b_train_cm <- confusionMatrix(data = oj_model_1b_train_preds, reference = oj_train$Purchase,
oj_model_1b_train_cm$overall
```

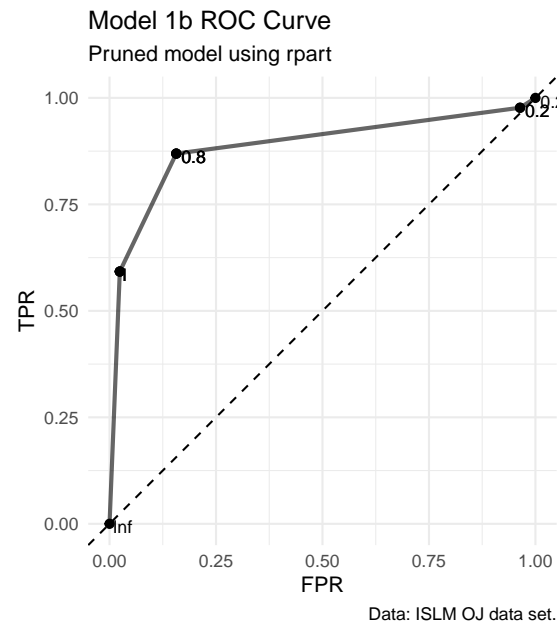
```
##      Accuracy      Kappa AccuracyLower AccuracyUpper AccuracyNull
## 8.226371e-01 6.323113e-01 7.953840e-01 8.476497e-01 6.102684e-01
## AccuracyPValue McNemarPValue
## 1.859617e-41 4.258396e-02
```

The accuracy on the training data set was a little lower than on the test data set. I thought it would be higher, not lower.

8.1.2 ROC Curve


Another measure of accuracy is the ROC (receiver operating characteristics) curve (Fawcett, 2005). The ROC curve is a plot of the true positive rate (TPR, sensitivity) versus the false positive rate (FPR, 1 - specificity) for a set of thresholds. By default, the threshold for predicting the default classification is 0.50, but it could be any threshold. The ROC curves varies the thresholds. (I'll use the `geom_roc` geom from **plotROC**).

```
data.frame(M = predict(oj_model_1b_pruned, oj_test, "prob")[, 1],
           D = if_else(oj_test$Purchase == "CH", 1, 0)) %>%
  ggplot() +
  geom_roc(aes(m = M, d = D), hjust = -0.4, vjust = 1.5, linealpha = 0.6, labelsize = 3, n.cuts = 10) +
  geom_abline(intercept = 0, slope = 1, linetype = 2) +
  coord_equal() +
  theme_minimal() +
  labs(x = "FPR", y = "TPR",
       title = "Model 1b ROC Curve",
       subtitle = "Pruned model using rpart",
       caption = "Data: ISLM OJ data set.")
```



You can also use `prediction()` and `plot.prediction()` from the **ROCR** package.

```
pred <- prediction(predict(oj_model_1b_pruned, newdata = oj_test, type = "prob"),
  plot(performance(pred, "tpr", "fpr"))
abline(0, 1, lty = 2)
```



Hmm, not quite the same...

A few points on the ROC space are helpful for understanding how to use it.

- The lower left point (0, 0) is the result of *always* predicting “negative” or in this case “MM” if “CH” is taken as the default class. Sure, your false positive rate is zero, but since you never predict a positive, your true positive rate is also zero.
- The upper right point (1, 1) is the results of *always* predicting “positive” (or “CH” here). You catch all the positives, but you miss all the negatives.
- The upper left point (0, 1) is the result of perfect accuracy. You catch all the positives and all the negatives.
- The lower right point (1, 0) is the result of perfect imbecility. You made the exact wrong prediction every time.
- The 45 degree diagonal is the result of randomly guessing positive (CH) X percent of the time. If you guess positive 90% of the time and the prevalence is 50%, your TPR will be 90% and your FPR will also be 90%, etc.

From the last bullet, it is evident that any point below and to the right of the 45 degree diagonal represents an instance where the model would have been better off just predicting entirely one way or the other. The goal is for all nodes to bunch up in the upper left.

Points to the left of the diagonal with a low TPR can be thought of as “conservative” predictors - they only make positive (CH) predictions with strong

evidence. Points to the left of the diagonal with a high TPR can be thought of as “liberal” predictors - they make positive (CH) predictions with weak evidence.

8.1.3 Caret Approach

I can also fit the model with `caret::train()`. There are two ways to tune hyperparameters in `train()`:

- set the number of tuning parameter values to consider by setting `tuneLength`, or
- set particular values to consider for each parameter by defining a `tuneGrid`.

I’ll build the model using 10-fold cross-validation to optimize the hyperparameter CP. If you don’t have any idea what the tuning parameter ought to look like, use `tuneLength` to get close, then fine-tune with `tuneGrid`. That’s what I’ll do. I’ll create a training control object that I can re-use in other model builds.

```
oj_trControl = trainControl(
  method = "cv", # k-fold cross validation
  number = 10, # 10 folds
  savePredictions = "final", # save predictions for the optimal tuning parameter
  classProbs = TRUE # return class probabilities in addition to predicted values
# summaryFunction = twoClassSummary # computes sensitivity, specificity and the area
)
```

Now fit the model.

```
set.seed(1234)
oj_model_2 = train(
  Purchase ~ .,
  data = oj_train,
  method = "rpart",
  tuneLength = 5,
  metric = "Accuracy",
  trControl = oj_trControl
)
```

`caret` built a full tree using `rpart`’s default parameters: gini splitting index, at least 20 observations in a node in order to consider splitting it, and at least 6 observations in each node. Caret then calculated the accuracy for each candidate value of α . Here is the results.

```
print(oj_model_2)

## CART
##
## 857 samples
## 17 predictor
## 2 classes: 'CH', 'MM'
##
## No pre-processing
## Resampling: Cross-Validated (10 fold)
## Summary of sample sizes: 772, 772, 771, 770, 771, 771, ...
## Resampling results across tuning parameters:
##
##      cp          Accuracy   Kappa
## 0.005988024 0.8085999 0.5931149
## 0.008982036 0.8086267 0.5943277
## 0.013473054 0.8051657 0.5885521
## 0.032934132 0.7841798 0.5371171
## 0.479041916 0.6603904 0.1774773
##
## Accuracy was used to select the optimal model using the largest value.
## The final value used for the model was cp = 0.008982036.
```

The second cp (0.008982036) produced the highest accuracy. I can drill into the best value of cp using a tuning grid. I'll try that now.

```
set.seed(1234)
oj_model_3 = train(
  Purchase ~ .,
  data = oj_train,
  method = "rpart",
  tuneGrid = expand.grid(cp = seq(from = 0.001, to = 0.010, length = 11)),
  metric = 'Accuracy',
  trControl = oj_trControl
)
print(oj_model_3)
```

```
## CART
##
## 857 samples
## 17 predictor
## 2 classes: 'CH', 'MM'
##
## No pre-processing
```

```
## Resampling: Cross-Validated (10 fold)
## Summary of sample sizes: 772, 772, 771, 770, 771, 771, ...
## Resampling results across tuning parameters:
##
##   cp      Accuracy   Kappa
##   0.0010  0.8004874  0.5753480
##   0.0019  0.8016502  0.5785232
##   0.0028  0.8039758  0.5845653
##   0.0037  0.8085999  0.5955198
##   0.0046  0.8039351  0.5851273
##   0.0055  0.8085863  0.5937949
##   0.0064  0.8085999  0.5931149
##   0.0073  0.8120883  0.6011446
##   0.0082  0.8120883  0.6011446
##   0.0091  0.8086267  0.5943277
##   0.0100  0.8086540  0.5953150
##
## Accuracy was used to select the optimal model using the largest value.
## The final value used for the model was cp = 0.0082.
```

The beset model is at $cp = 0.009$. Here are the rules in the final model.

```
oj_model_3$finalModel
```

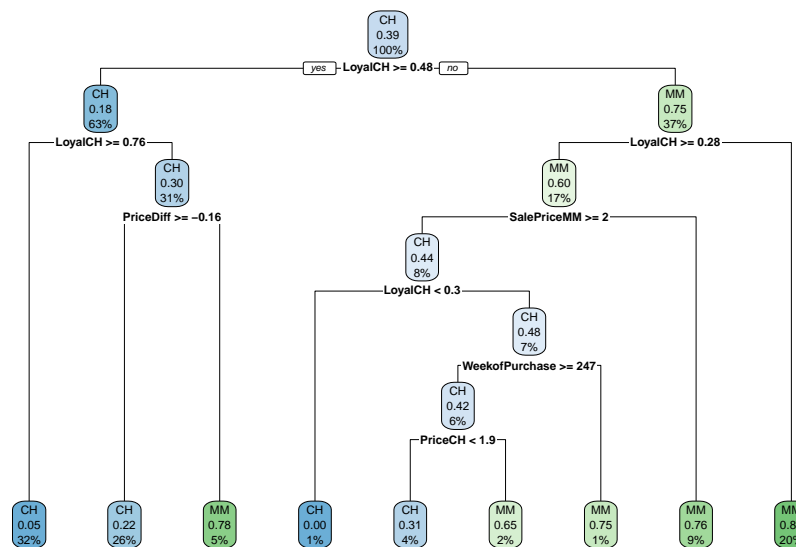
```
## n= 857
##
## node), split, n, loss, yval, (yprob)
##      * denotes terminal node
##
## 1) root 857 334 CH (0.61026838 0.38973162)
##    2) LoyalCH>=0.48285 537 94 CH (0.82495345 0.17504655)
##      4) LoyalCH>=0.7648795 271 13 CH (0.95202952 0.04797048) *
##      5) LoyalCH< 0.7648795 266 81 CH (0.69548872 0.30451128)
##        10) PriceDiff>=-0.165 226 50 CH (0.77876106 0.22123894) *
##        11) PriceDiff< -0.165 40 9 MM (0.22500000 0.77500000) *
##    3) LoyalCH< 0.48285 320 80 MM (0.25000000 0.75000000)
##      6) LoyalCH>=0.2761415 146 58 MM (0.39726027 0.60273973)
##        12) SalePriceMM>=2.04 71 31 CH (0.56338028 0.43661972)
##          24) LoyalCH< 0.303104 7 0 CH (1.00000000 0.00000000) *
##          25) LoyalCH>=0.303104 64 31 CH (0.51562500 0.48437500)
##            50) WeekofPurchase>=246.5 52 22 CH (0.57692308 0.42307692)
##              100) PriceCH< 1.94 35 11 CH (0.68571429 0.31428571) *
##              101) PriceCH>=1.94 17 6 MM (0.35294118 0.64705882) *
##            51) WeekofPurchase< 246.5 12 3 MM (0.25000000 0.75000000) *
##    13) SalePriceMM< 2.04 75 18 MM (0.24000000 0.76000000) *
```



```
##          7) LoyalCH< 0.2761415 174 22 MM (0.12643678 0.87356322) *
```

Here is the tree.

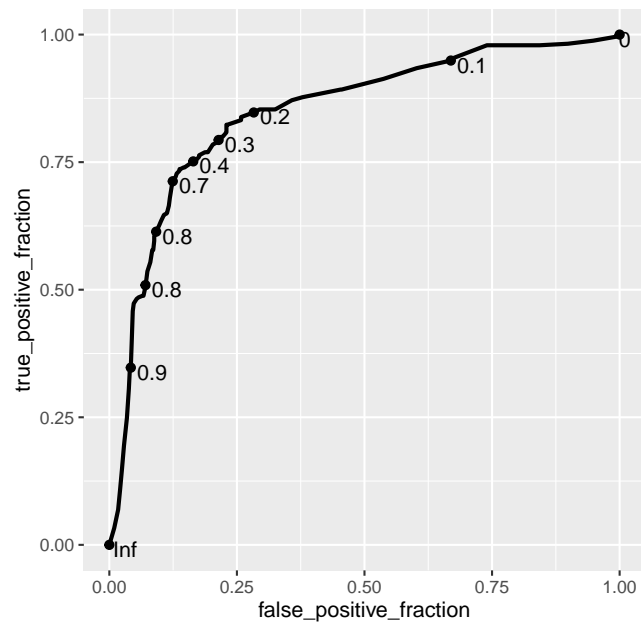
```
rpart.plot(oj_model_3$finalModel)
```



Here is the ROC curve.


```
library(plotROC)
ggplot(oj_model_3$pred) +
  geom_roc(
    aes(
      m = MM,
      d = factor(obs, levels = c("CH", "MM"))
    ),
    hjust = -0.4, vjust = 1.5
  ) +
  coord_equal()
```

```
## Warning in verify_d(data$d): D not labeled 0/1, assuming CH = 0 and MM = 1!
```



Here are the cross-validated Accuracy for each candidate cp value.

```
plot(oj_model_3)
```



Evaluate the model by making predictions with the test data set.

```
oj_model_3_preds <- predict(oj_model_3, oj_test, type = "raw")
```

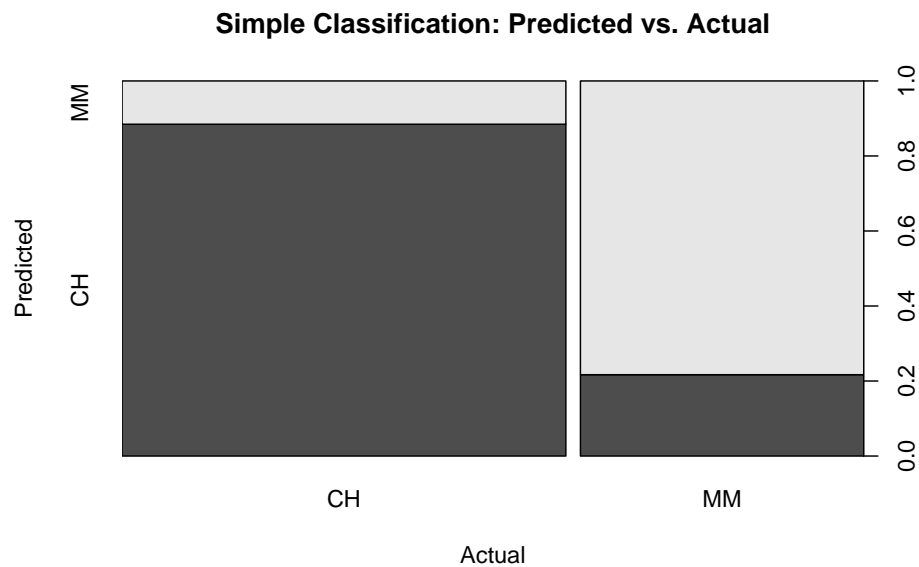
The confusion matrix shows the true positives and true negatives.

```
oj_model_3_cm <- confusionMatrix(
  data = oj_model_3_preds,
  reference = oj_test$Purchase
)
oj_model_3_cm

## Confusion Matrix and Statistics
##
##           Reference
## Prediction CH  MM
##           CH 115  18
##           MM  15  65
##
##              Accuracy : 0.8451
##              95% CI : (0.7894, 0.8909)
##    No Information Rate : 0.6103
##    P-Value [Acc > NIR] : 6.311e-14
##
##              Kappa : 0.6721
##
##  Mcnemar's Test P-Value : 0.7277
##
##              Sensitivity : 0.8846
##              Specificity : 0.7831
##              Pos Pred Value : 0.8647
##              Neg Pred Value : 0.8125
##              Prevalence : 0.6103
##              Detection Rate : 0.5399
##    Detection Prevalence : 0.6244
##              Balanced Accuracy : 0.8339
##
##              'Positive' Class : CH
##
```

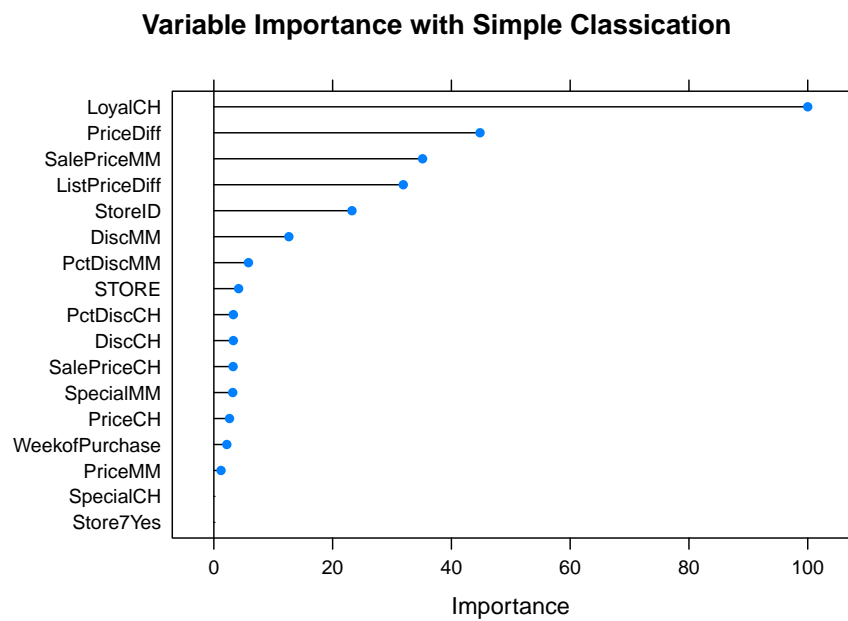
The accuracy metric is the slightly worse than in my previous model. Here is a graphical representation of the confusion matrix.

```
plot(oj_test$Purchase, oj_model_3_preds,
     main = "Simple Classification: Predicted vs. Actual",
     xlab = "Actual",
     ylab = "Predicted")
```



Finally, here is the variable importance plot.

```
plot(varImp(oj_model_3), main="Variable Importance with Simple Classification")
```



Looks like the manual effort faired best. Here is a summary the accuracy rates of the three models.

```
rbind(data.frame(model = "Manual Class", Acc = round(oj_model_1b_cm$overall["Accuracy"], 5)),
      data.frame(model = "Caret w/tuneGrid", Acc = round(oj_model_3_cm$overall["Accuracy"], 5))
)
```

```
##              model      Acc
## Accuracy      Manual Class 0.85915
## Accuracy1 Caret w/tuneGrid 0.84507
```

8.2 Regression Trees

A simple regression tree is built in a manner similar to a simple classification tree, and like the simple classification tree, it is rarely invoked on its own; the bagged, random forest, and gradient boosting methods build on this logic. I'll learn by example again. Using the `ISLR::Carseats` data set, I will predict `Sales` using from the 10 feature variables. Load the data.

```
carseats_dat <- Carseats
skim_with(numeric = list(p0 = NULL, p25 = NULL, p50 = NULL, p75 = NULL,
                        p100 = NULL, hist = NULL))
skim(carseats_dat)
```

```
## Skim summary statistics
##   n obs: 400
##   n variables: 11
##
## -- Variable type:factor -----
##   variable missing complete   n n_unique      top_counts
##   ShelfLoc      0      400 400         3 Med: 219, Bad: 96, Goo: 85, NA: 0
##     Urban      0      400 400         2      Yes: 282, No: 118, NA: 0
##        US      0      400 400         2      Yes: 258, No: 142, NA: 0
##   ordered
##   FALSE
##   FALSE
##   FALSE
##
## -- Variable type:numeric -----
##   variable missing complete   n   mean    sd
##   Advertising      0      400 400   6.63   6.65
##        Age      0      400 400  53.32  16.2
##    CompPrice      0      400 400 124.97  15.33
```

```
##      Education      0      400 400   13.9    2.62
##      Income        0      400 400   68.66   27.99
##      Population    0      400 400  264.84  147.38
##      Price         0      400 400   115.8   23.68
##      Sales         0      400 400    7.5    2.82
```

I'll split `careseats_dat` ($n = 400$) into `careseats_train` (80%, $n = 321$) and `careseats_test` (20%, $n = 79$). I'll fit a simple decision tree with `careseats_train`, then later a bagged tree, a random forest, and a gradient boosting tree. I'll compare their predictive performance with `careseats_test`.

```
set.seed(12345)
partition <- createDataPartition(y = careseats_dat$Sales, p = 0.8, list = FALSE)
careseats_train <- careseats_dat[partition, ]
careseats_test <- careseats_dat[-partition, ]
```

The first step is to build a full tree, then perform k-fold cross-validation to help select the optimal cost complexity (cp). The only difference here is the `rpart()` parameter `method = "anova"` to produce a regression tree.

```
set.seed(1234)
careseats_model_1 <- rpart(
  formula = Sales ~ .,
  data = careseats_train,
  method = "anova",
  xval = 10,
  model = TRUE # to plot splits with factor variables.
)
print(careseats_model_1)
```

```
## n= 321
##
## node), split, n, deviance, yval
##      * denotes terminal node
##
## 1) root 321 2567.76800  7.535950
##    2) ShelfLoc=Bad,Medium 251 1474.14100  6.770359
##      4) Price>=105.5 168  719.70630  5.987024
##        8) ShelfLoc=Bad 50  165.70160  4.693600
##          16) Population< 201.5 20  48.35505  3.646500 *
##          17) Population>=201.5 30  80.79922  5.391667 *
##          9) ShelfLoc=Medium 118  434.91370  6.535085
##            18) Advertising< 11.5 88  290.05490  6.113068
##              36) CompPrice< 142 69  193.86340  5.769420
```

```

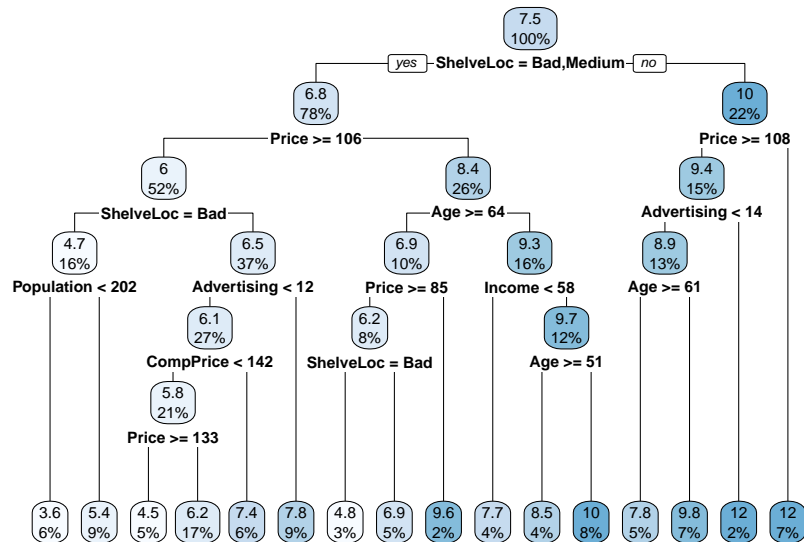
##          72) Price>=132.5 16    50.75440  4.455000 *
##          73) Price< 132.5 53   107.12060  6.166226 *
##          37) CompPrice>=142 19    58.45118  7.361053 *
##          19) Advertising>=11.5 30    83.21323  7.773000 *
##    5) Price< 105.5 83   442.68920  8.355904
##          10) Age>=63.5 32   153.42300  6.922500
##          20) Price>=85 25    66.89398  6.160800
##          40) ShelveLoc=Bad 9    18.39396  4.772222 *
##          41) ShelveLoc=Medium 16    21.38544  6.941875 *
##          21) Price< 85 7    20.22194  9.642857 *
##          11) Age< 63.5 51   182.26350  9.255294
##          22) Income< 57.5 12    28.03042  7.707500 *
##          23) Income>=57.5 39   116.63950  9.731538
##          46) Age>=50.5 14    21.32597  8.451429 *
##          47) Age< 50.5 25    59.52474  10.448400 *
##    3) ShelveLoc=Good 70   418.98290  10.281140
##          6) Price>=107.5 49   242.58730  9.441633
##          12) Advertising< 13.5 41   162.47820  8.926098
##          24) Age>=61 17    53.37051  7.757647 *
##          25) Age< 61 24    69.45776  9.753750 *
##          13) Advertising>=13.5 8    13.36599  12.083750 *
##          7) Price< 107.5 21    61.28200  12.240000 *

```

The output starts with the root node. The predicted **Sales** at the root is the mean **Sales** for the training data set, 7.535950 (values are \$000s). The deviance at the root is the SSE, 2567.768. The child nodes of node “x” are labeled 2x) and 2x+1), so the child nodes of 1) are 2) and 3), and the child nodes of 2) are 4) and 5). Terminal nodes are labeled with an asterisk (*).

The first split is at **ShelveLoc** = [Bad, Medium] vs Good. Here is what the full (unpruned) tree looks like.

```
rpart.plot(carseats_model_1, yesno = TRUE)
```



The boxes show the node predicted value (mean) and the proportion of observations that are in the node (or child nodes).

`rpart()` not only grew the full tree, it also used cross-validation to test the performance of the possible complexity hyperparameters. `printcp()` displays the candidate cp values. You can use this table to decide how to prune the tree.

```
printcp(carseats_model_1)
```

```
##
## Regression tree:
## rpart(formula = Sales ~ ., data = carseats_train, method = "anova",
##       model = TRUE, xval = 10)
##
## Variables actually used in tree construction:
## [1] Advertising Age          CompPrice  Income      Population Price
## [7] ShelveLoc
##
## Root node error: 2567.8/321 = 7.9993
##
## n= 321
##
##      CP nsplit rel error  xerror    xstd
## 1 0.262736      0  1.00000 1.00635 0.076664
## 2 0.121407      1  0.73726 0.74888 0.058981
```



```
## 3 0.046379      2 0.61586 0.65278 0.050839
## 4 0.044830      3 0.56948 0.67245 0.051638
## 5 0.041671      4 0.52465 0.66230 0.051065
## 6 0.025993      5 0.48298 0.62345 0.049368
## 7 0.025823      6 0.45698 0.61980 0.048026
## 8 0.024007      7 0.43116 0.62058 0.048213
## 9 0.015441      8 0.40715 0.58061 0.041738
## 10 0.014698     9 0.39171 0.56413 0.041368
## 11 0.014641    10 0.37701 0.56277 0.041271
## 12 0.014233    11 0.36237 0.56081 0.041097
## 13 0.014015    12 0.34814 0.55647 0.038308
## 14 0.013938    13 0.33413 0.55647 0.038308
## 15 0.010560    14 0.32019 0.57110 0.038872
## 16 0.010000    15 0.30963 0.56676 0.038090
```

There are 16 possible cp values in this model. The model with the smallest complexity parameter allows the most splits (`nsplit`). The highest complexity parameter corresponds to a tree with just a root node. `rel error` is the SSE relative to the root node. The root node SSE is 2567.76800, so its `rel error` is $2567.76800/2567.76800 = 1.0$. That means the absolute error of the full tree (at $CP = 0.01$) is $0.30963 * 2567.76800 = 795.058$. You can verify that by calculating the SSE of the model predicted values:

```
data.frame(pred = predict(carseats_model_1, newdata = carseats_train)) %>%
  mutate(obs = carseats_train$Sales,
         sq_err = (obs - pred)^2) %>%
  summarize(sse = sum(sq_err))
```

```
##           sse
## 1 795.0525
```

Finishing the CP table tour, `xerror` is the cross-validated SSE and `xstd` is its standard error. If you want the lowest possible error, then prune to the tree with the smallest relative SSE (`xerror`). If you want to balance predictive power with simplicity, prune to the smallest tree within 1 SE of the one with the smallest relative SSE. The CP table is not super-helpful for finding that tree. I'll add a column to find it.


```
carseats_model_1$cptable %>%
  data.frame() %>%
  mutate(min_xerror_idx = which.min(carseats_model_1$cptable[, "xerror"]),
         rownum = row_number(),
         xerror_cap = carseats_model_1$cptable[min_xerror_idx, "xerror"] +
           carseats_model_1$cptable[min_xerror_idx, "xstd"],
```

```
eval = case_when(rownum == min_xerror_idx ~ "min xerror",
                  xerror < xerror_cap ~ "under cap",
                  TRUE ~ "") %>%
select(-rownum, -min_xerror_idx)
```

##	CP	nsplit	rel.error	xerror	xstd	xerror_cap	eval
## 1	0.26273578	0	1.0000000	1.0063530	0.07666355	0.5947744	
## 2	0.12140705	1	0.7372642	0.7488767	0.05898146	0.5947744	
## 3	0.04637919	2	0.6158572	0.6527823	0.05083938	0.5947744	
## 4	0.04483023	3	0.5694780	0.6724529	0.05163819	0.5947744	
## 5	0.04167149	4	0.5246478	0.6623028	0.05106530	0.5947744	
## 6	0.02599265	5	0.4829763	0.6234457	0.04936799	0.5947744	
## 7	0.02582284	6	0.4569836	0.6198034	0.04802643	0.5947744	
## 8	0.02400748	7	0.4311608	0.6205756	0.04821332	0.5947744	
## 9	0.01544139	8	0.4071533	0.5806072	0.04173785	0.5947744	under cap
## 10	0.01469771	9	0.3917119	0.5641331	0.04136793	0.5947744	under cap
## 11	0.01464055	10	0.3770142	0.5627713	0.04127139	0.5947744	under cap
## 12	0.01423309	11	0.3623736	0.5608073	0.04109662	0.5947744	under cap
## 13	0.01401541	12	0.3481405	0.5564663	0.03830810	0.5947744	min xerror
## 14	0.01393771	13	0.3341251	0.5564663	0.03830810	0.5947744	under cap
## 15	0.01055959	14	0.3201874	0.5710951	0.03887227	0.5947744	under cap
## 16	0.01000000	15	0.3096278	0.5667561	0.03808991	0.5947744	under cap


Okay, so the simplest tree is the one with $CP = 0.01544139$ (8 splits). Fortunately, `plotcp()` presents a nice graphical representation of the relationship between `xerror` and `cp`.

```
plotcp(carseats_model_1, upper = "splits")
```



The dashed line is set at the minimum `xerror` + `xstd`. The top axis shows the number of splits in the tree. I'm not sure why the CP values are not the same as in the table (they are close, but not the same). The smallest relative error is at 0.0140154, but the maximum CP below the dashed line (one standard deviation above the minimum error) is at CP = .019 (8 splits). Use the `prune()` function to prune the tree by specifying the associated cost-complexity `cp`.

```
carseats_model_1_pruned <- prune(
  carseats_model_1,
  cp = carseats_model_1$cptable[carseats_model_1$cptable[, 2] == 8, "CP"]
)
rpart.plot(carseats_model_1_pruned, yesno = TRUE)
```

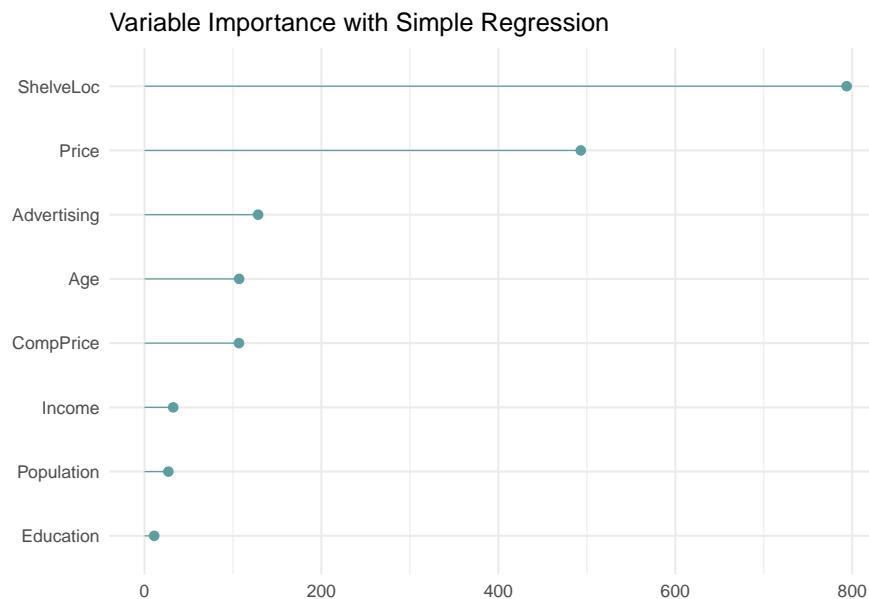


The most “important” indicator of Sales is `ShelveLoc`. Here are the importance values from the model.

```

carseats_model_1_pruned$variable.importance %>%
  data.frame() %>%
  rownames_to_column(var = "Feature") %>%
  rename(Overall = '.') %>%
  ggplot(aes(x = fct_reorder(Feature, Overall), y = Overall)) +
  geom_pointrange(aes(ymin = 0, ymax = Overall), color = "cadetblue", size = .3) +
  theme_minimal() +
  coord_flip() +
  labs(x = "", y = "", title = "Variable Importance with Simple Regression")

```



The most important indicator of **Sales** is **ShelveLoc**, then **Price**, then **Age**, all of which appear in the final model. **CompPrice** was also important.

The last step is to make predictions on the validation data set. The root mean squared error ($RMSE = \sqrt{(1/2) \sum (actual - pred)^2}$) and mean absolute error ($MAE = (1/n) \sum |actual - pred|$) are the two most common measures of predictive accuracy. The key difference is that RMSE punishes large errors more harshly. For a regression tree, set argument `type = "vector"` (or do not specify at all).

```
carseats_model_1_preds <- predict(
  carseats_model_1_pruned,
  carseats_test,
  type = "vector"
)


carseats_model_1_pruned_rmse <- RMSE(
  pred = carseats_model_1_preds,
  obs = carseats_test$Sales
)
carseats_model_1_pruned_rmse
```

```
## [1] 2.388059
```

The pruning process leads to an average prediction error of 2.388 in the test

data set. Not too bad considering the standard deviation of `Sales` is 2.801. Here is a predicted vs actual plot.

```
plot(carseats_test$Sales, carseats_model_1_preds,
     main = "Simple Regression: Predicted vs. Actual",
     xlab = "Actual",
     ylab = "Predicted")
abline(0, 1)
```



The 6 possible predicted values do a decent job of binning the observations.

8.2.1 Caret Approach

I can also fit the model with `caret::train()`, specifying `method = "rpart"`.

I'll build the model using 10-fold cross-validation to optimize the hyperparameter `CP`.

```
carseats_trControl = trainControl(
  method = "cv", # k-fold cross validation
  number = 10, # 10 folds
  savePredictions = "final" # save predictions for the optimal tuning parameter
)
```

I'll let the model look for the best CP tuning parameter with `tuneLength` to get close, then fine-tune with `tuneGrid`.

```
set.seed(1234)
carseats_model_2 = train(
  Sales ~ .,
  data = carseats_train,
  method = "rpart", # for classification tree
  tuneLength = 5, # choose up to 5 combinations of tuning parameters (cp)
  metric = "RMSE", # evaluate hyperparameter combinations with RMSE
  trControl = carseats_trControl
)
```

```
## Warning in nominalTrainWorkflow(x = x, y = y, wts = weights, info =
## trainInfo, : There were missing values in resampled performance measures.
```

```
print(carseats_model_2)
```

```
## CART
##
## 321 samples
## 10 predictor
##
## No pre-processing
## Resampling: Cross-Validated (10 fold)
## Summary of sample sizes: 289, 289, 289, 289, 289, 289, ...
## Resampling results across tuning parameters:
##
##   cp          RMSE      Rsquared    MAE
## 0.04167149  2.209383  0.4065251  1.778797
## 0.04483023  2.243618  0.3849728  1.805027
## 0.04637919  2.275563  0.3684309  1.808814
## 0.12140705  2.400455  0.2942663  1.936927
## 0.26273578  2.692867  0.1898998  2.192774
##
## RMSE was used to select the optimal model using the smallest value.
## The final value used for the model was cp = 0.04167149.
```

The first cp (0.04167149) produced the smallest RMSE. I can drill into the best value of cp using a tuning grid. I'll try that now.

```
myGrid <- expand.grid(cp = seq(from = 0, to = 0.1, by = 0.01))
carseats_model_3 = train(
  Sales ~ .,
```

```

data = carseats_train,
method = "rpart", # for classification tree
tuneGrid = myGrid, # choose up to 5 combinations of tuning parameters (cp)
metric = "RMSE", # evaluate hyperparameter combinations with RMSE
trControl = carseats_trControl
)
print(carseats_model_3)

```

```

## CART
##
## 321 samples
## 10 predictor
##
## No pre-processing
## Resampling: Cross-Validated (10 fold)
## Summary of sample sizes: 289, 289, 289, 289, 288, 289, ...
## Resampling results across tuning parameters:
##
##   cp      RMSE      Rsquared    MAE
##   0.00  2.131814  0.4578761  1.725960
##   0.01  2.203111  0.4294647  1.790050
##   0.02  2.240209  0.3948080  1.834786
##   0.03  2.206168  0.4139717  1.762170
##   0.04  2.274313  0.3686176  1.795154
##   0.05  2.309746  0.3405228  1.830556
##   0.06  2.246757  0.3703977  1.780266
##   0.07  2.253725  0.3679986  1.794485
##   0.08  2.253725  0.3679986  1.794485
##   0.09  2.253725  0.3679986  1.794485
##   0.10  2.253725  0.3679986  1.794485
##
## RMSE was used to select the optimal model using the smallest value.
## The final value used for the model was cp = 0.


```

It looks like the best performing tree is the unpruned one.

```

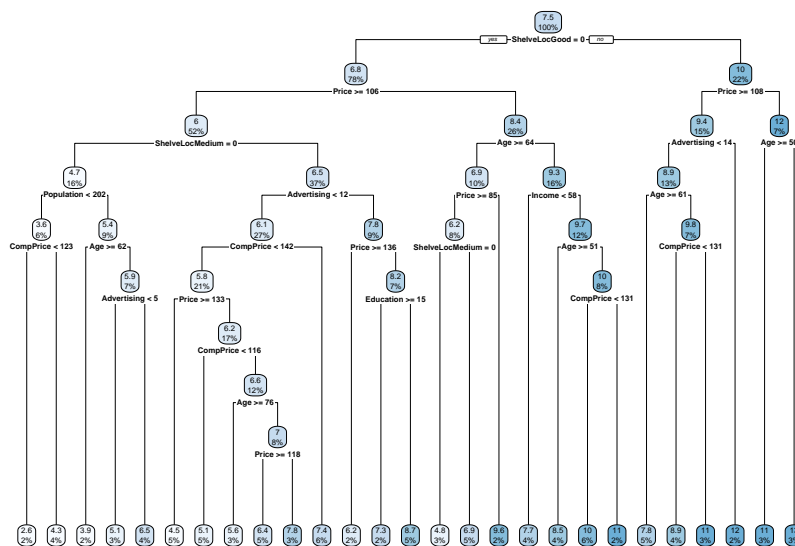
plot(carseats_model_3)

```


Lets's see the final model.

```
rpart.plot(carseats_model_3$finalModel)
```



What were the most important variables?


```
plot(varImp(carseats_model_3), main="Variable Importance with Simple Regression")
```



Evaluate the model by making predictions with the test data set.

```
carseats_model_3_preds <- predict(carseats_model_3, carseats_test, type = "raw")
data.frame(Actual = carseats_test$Sales, Predicted = carseats_model_3_preds) %>%
ggplot(aes(x = Actual, y = Predicted)) +
  geom_point() +
  geom_smooth() +
  geom_abline(slope = 1, intercept = 0) +
  scale_y_continuous(limits = c(0, 15)) +
  labs(title = "Simple Regression: Predicted vs. Actual")
```

```
## `geom_smooth()` using method = 'loess' and formula 'y ~ x'
```



Looks like the model over-estimates at the low end and underestimates at the high end. Calculate the test data set RMSE.

```
carseats_model_3_pruned_rmse <- RMSE(
  pred = carseats_model_3_preds,
  obs = carseats_test$Sales
)
carseats_model_3_pruned_rmse
```

```
## [1] 2.298331
```

Caret faired better in this model. Here is a summary the RMSE values of the two models.

```
rbind(data.frame(model = "Manual ANOVA",
  RMSE = round(carseats_model_1_pruned_rmse, 5)),
  data.frame(model = "Caret",
  RMSE = round(carseats_model_3_pruned_rmse, 5))
)
```

```
##      model      RMSE
## 1 Manual ANOVA 2.38806
## 2      Caret 2.29833
```

8.3 Bagging

Bootstrap aggregation, or *bagging*, is a general-purpose procedure for reducing the variance of a statistical learning method. The algorithm constructs B regression trees using B bootstrapped training sets, and averages the resulting predictions. These trees are grown deep, and are not pruned. Hence each individual tree has high variance, but low bias. Averaging these B trees reduces the variance. For classification trees, bagging takes the “majority vote” for the prediction. Use a value of B sufficiently large that the error has settled down.

To test the model accuracy, the out-of-bag observations are predicted from the models that do not use them. If $B/3$ of observations are in-bag, there are $B/3$ predictions per observation. These predictions are averaged for the test prediction. Again, for classification trees, a majority vote is taken.

The downside to bagging is that it improves accuracy at the expense of interpretability. There is no longer a single tree to interpret, so it is no longer clear which variables are more important than others.

Bagged trees are a special case of random forests, so see the next section for an example.

8.4 Random Forests

Random forests improve bagged trees by way of a small tweak that de-correlates the trees. As in bagging, the algorithm builds a number of decision trees on bootstrapped training samples. But when building these decision trees, each time a split in a tree is considered, a random sample of $mtry$ predictors is chosen as split candidates from the full set of p predictors. A fresh sample of $mtry$ predictors is taken at each split. Typically $mtry \sim \sqrt{p}$. Bagged trees are thus a special case of random forests where $mtry = p$.

8.4.0.1 Bagging Classification Example

Again using the OJ data set to predict `Purchase`, this time I'll use the bagging method by specifying `method = "treebag"`. I'll use `tuneLength = 5` and not worry about `tuneGrid` anymore. Caret has no hyperparameters to tune with this model.

```
oj.bag = train(Purchase ~ .,
               data = oj_train,
               method = "treebag", # for bagging
               tuneLength = 5, # choose up to 5 combinations of tuning parameters
               metric = "ROC", # evaluate hyperparameter combinations with ROC
```

```

trControl = trainControl(
  method = "cv", # k-fold cross validation
  number = 10, # k=10 folds
  savePredictions = "final", # save predictions for the optimal tuning parameter
  classProbs = TRUE, # return class probabilities in addition to predicted values
  summaryFunction = twoClassSummary # for binary response variable
)
oj.bag

```

```


## Bagged CART
##
## 857 samples
## 17 predictor
## 2 classes: 'CH', 'MM'
##
## No pre-processing
## Resampling: Cross-Validated (10 fold)
## Summary of sample sizes: 771, 772, 771, 771, 771, 772, ...
## Resampling results:
##
##      ROC      Sens      Spec
## 0.8524038 0.8165094 0.7217469

```

```

#plot(oj.bag$)
oj.pred <- predict(oj.bag, oj_test, type = "raw")
plot(oj_test$Purchase, oj.pred,
     main = "Bagging Classification: Predicted vs. Actual",
     xlab = "Actual",
     ylab = "Predicted")

```




```
(oj.conf <- confusionMatrix(data = oj.pred,
                             reference = oj_test$Purchase))
```

```
## Confusion Matrix and Statistics
##
##           Reference
## Prediction CH  MM
##      CH 110  17
##      MM  20  66
##
##           Accuracy : 0.8263
##           95% CI : (0.7686, 0.8746)
##      No Information Rate : 0.6103
##      P-Value [Acc > NIR] : 7.121e-12
##
##           Kappa : 0.6372
##
##  McNemar's Test P-Value : 0.7423
##
##           Sensitivity : 0.8462
##           Specificity : 0.7952
##      Pos Pred Value : 0.8661
##      Neg Pred Value : 0.7674
##           Prevalence : 0.6103
```

```
##          Detection Rate : 0.5164
##    Detection Prevalence : 0.5962
##          Balanced Accuracy : 0.8207
##
##          'Positive' Class : CH
##
```

```
oj.bag.acc <- as.numeric(oj.conf$overall[1])
rm(oj.pred)
rm(oj.conf)
#plot(oj.bag$, oj.bag$finalModel$y)
plot(varImp(oj.bag), main="Variable Importance with Simple Classification")
```



8.4.0.2 Random Forest Classification Example

Now I'll try it with the random forest method by specifying `method = "ranger"`. I'll stick with `tuneLength = 5`. Caret tunes three hyperparameters:

- `mtry`: number of randomly selected predictors. Default is `sqrt(p)`.
- `splitrule`: splitting rule. For classification, options are "gini" (default) and "extratrees".
- `min.node.size`: minimal node size. Default is 1 for classification.

```

oj.frst = train(Purchase ~ .,
               data = oj_train,
               method = "ranger", # for random forest
               tuneLength = 5, # choose up to 5 combinations of tuning parameters
               metric = "ROC", # evaluate hyperparameter combinations with ROC
               trControl = trainControl(
                 method = "cv", # k-fold cross validation
                 number = 10, # 10 folds
                 savePredictions = "final", # save predictions for the optimal t
                 classProbs = TRUE, # return class probabilities in addition to predi
                 summaryFunction = twoClassSummary # for binary response variable
               )
oj.frst

```


```

## Random Forest
##
## 857 samples
## 17 predictor
## 2 classes: 'CH', 'MM'
##
## No pre-processing
## Resampling: Cross-Validated (10 fold)
## Summary of sample sizes: 772, 771, 772, 770, 772, 772, ...
## Resampling results across tuning parameters:
##
##  mtry  splitrule  ROC          Sens          Spec
##  2      gini      0.8603930  0.8719158  0.6946524
##  2      extratrees 0.8499806  0.8814586  0.6287879
##  5      gini      0.8683505  0.8470247  0.7246881
##  5      extratrees 0.8642275  0.8584543  0.6886809
##  9      gini      0.8687568  0.8374456  0.7272727
##  9      extratrees 0.8670702  0.8451379  0.6858289
##  13     gini      0.8642114  0.8297896  0.7361854
##  13     extratrees 0.8680705  0.8298258  0.7064171
##  17     gini      0.8642378  0.8145501  0.7423351
##  17     extratrees 0.8634162  0.8260160  0.7093583
##
## Tuning parameter 'min.node.size' was held constant at a value of 1
## ROC was used to select the optimal model using the largest value.
## The final values used for the model were mtry = 9, splitrule = gini
## and min.node.size = 1.


```



```
plot(oj.first)
```



```
oj.pred <- predict(oj.first, oj_test, type = "raw")
plot(oj_test$Purchase, oj.pred,
     main = "Random Forest Classification: Predicted vs. Actual",
     xlab = "Actual",
     ylab = "Predicted")
```



```
(oj.conf <- confusionMatrix(data = oj.pred,
                             reference = oj_test$Purchase))
```

```
## Confusion Matrix and Statistics
##
##           Reference
## Prediction  CH  MM
##      CH 109  16
##      MM  21  67
##
##           Accuracy : 0.8263
##           95% CI : (0.7686, 0.8746)
##      No Information Rate : 0.6103
##      P-Value [Acc > NIR] : 7.121e-12
##
##           Kappa : 0.6387
##
##  McNemar's Test P-Value : 0.5108
##
##           Sensitivity : 0.8385
##           Specificity : 0.8072
##      Pos Pred Value : 0.8720
##      Neg Pred Value : 0.7614
##           Prevalence : 0.6103
```

```
##           Detection Rate : 0.5117
##      Detection Prevalence : 0.5869
##           Balanced Accuracy : 0.8228
##
##           'Positive' Class : CH
##
```

```
oj.first.acc <- as.numeric(oj.conf$overall[1])
rm(oj.pred)
rm(oj.conf)
#plot(oj.bag$, oj.bag$finalModel$y)
#plot(varImp(oj.frst), main="Variable Importance with Simple Classification")
```

The model algorithm explains “ROC was used to select the optimal model using the largest value. The final values used for the model were `mtry = 9`, `splitrule = extratrees` and `min.node.size = 1`.” You can see the results of tuning grid combinations in the associated plot of ROC AUC vs `mtry` grouped by splitting rule.

The bagging (accuracy = 0.80751) and random forest (accuracy = 0.81690) models fared pretty well, but the manual classification tree is still in first place. There’s still gradient boosting to investigate!

```
rbind(data.frame(model = "Manual Class", Accuracy = round(oj_model_1b_cm$overall["Accuracy"], 5)),
      data.frame(model = "Caret w.tuneGrid", Accuracy = round(oj_model_3_cm$overall["Accuracy"], 5)),
      data.frame(model = "Bagging", Accuracy = round(oj.bag.acc, 5)),
      data.frame(model = "Random Forest", Accuracy = round(oj.first.acc, 5))
) %>% arrange(desc(Accuracy))
```

```
##           model Accuracy
## 1      Manual Class 0.85915
## 2 Caret w.tuneGrid 0.84507
## 3           Bagging 0.82629
## 4      Random Forest 0.82629
```

8.4.0.3 Bagging Regression Example

Again using the `Carseats` data set to predict `Sales`, this time I’ll use the bagging method by specifying `method = "treebag"`. I’ll use `tuneLength = 5` and not worry about `tuneGrid` anymore. Caret has no hyperparameters to tune with this model.

```
carseats.bag = train(Sales ~ .,
                     data = carseats_train,
```

```

        method = "treebag", # for bagging
        tuneLength = 5, # choose up to 5 combinations of tuning parameters
        metric = "RMSE", # evaluate hyperparameter combinations with RMSE
        trControl = trainControl(
            method = "cv", # k-fold cross validation
            number = 10, # 10 folds
            savePredictions = "final" # save predictions for the optimal tuning
        )
    )
carseats.bag

```

```


## Bagged CART
##
## 321 samples
## 10 predictor
##
## No pre-processing
## Resampling: Cross-Validated (10 fold)
## Summary of sample sizes: 289, 289, 289, 288, 289, 289, ...
## Resampling results:
##
##    RMSE      Rsquared    MAE
## 1.709371 0.6532837 1.374155

```

```

#plot(carseats.bag$finalModel)
carseats.pred <- predict(carseats.bag, carseats_test, type = "raw")
plot(carseats_test$Sales, carseats.pred,
     main = "Bagging Regression: Predicted vs. Actual",
     xlab = "Actual",
     ylab = "Predicted")
abline(0, 1)

```




```
(carseats.bag.rmse <- RMSE(pred = carseats.pred,  
                           obs = carseats_test$Sales))
```

```
## [1] 1.932792
```

```
rm(carseats.pred)  
plot(varImp(carseats.bag), main="Variable Importance with Regression Bagging")
```

Variable Importance with Regression Bagging



8.4.0.4 Random Forest Regression Example

Now I'll try it with the random forest method by specifying `method = "ranger"`. I'll stick with `tuneLength = 5`. Caret tunes three hyperparameters:


- `mtry`: number of randomly selected predictors
- `splitrule`: splitting rule. For regression, options are “variance” (default), “extratrees”, and “maxstat”.
- `min.node.size`: minimal node size

```
carseats.frst = train(Sales ~ .,
  data = carseats_train,
  method = "ranger", # for random forest
  tuneLength = 5, # choose up to 5 combinations of tuning parameters
  metric = "RMSE", # evaluate hyperparamter combinations with RMSE
  trControl = trainControl(
    method = "cv", # k-fold cross validation
    number = 10, # 10 folds
    savePredictions = "final" # save predictions for the optimal tuning
  )
)
carseats.frst
```

```
## Random Forest
```

```
##
## 321 samples
## 10 predictor
##
## No pre-processing
## Resampling: Cross-Validated (10 fold)
## Summary of sample sizes: 289, 289, 289, 289, 289, 288, ...
## Resampling results across tuning parameters:
##
##   mtry  splitrule  RMSE      Rsquared  MAE
##   2     variance  1.806943  0.6957452  1.446420
##   2     extratrees 1.905011  0.6466527  1.539096
##   4     variance  1.600763  0.7288625  1.266868
##   4     extratrees 1.702009  0.6862545  1.357981
##   6     variance  1.541675  0.7336448  1.217061
##   6     extratrees 1.639248  0.6966549  1.302159
##   8     variance  1.542806  0.7236085  1.221671
##   8     extratrees 1.601484  0.7053834  1.269484
##  11     variance  1.555271  0.7168108  1.230252
##  11     extratrees 1.589152  0.7058982  1.255090
##
## Tuning parameter 'min.node.size' was held constant at a value of 5
## RMSE was used to select the optimal model using the smallest value.
## The final values used for the model were mtry = 6, splitrule =
##   variance and min.node.size = 5.
```


```
plot(carseats.frst)
```



```

carseats.pred <- predict(carseats.frst, carseats_test, type = "raw")
plot(carseats_test$Sales, carseats.pred,
     main = "Random Forest Regression: Predicted vs. Actual",
     xlab = "Actual",
     ylab = "Predicted")
abline(0, 1)

```

```
(carseats.frst.rmse <- RMSE(pred = carseats.pred,
  obs = carseats_test$Sales))
```

```
## [1] 1.758112
```

```
rm(carseats.pred)
#plot(varImp(carseats.frst), main="Variable Importance with Regression Random Forest")
```

The model algorithm explains “*RMSE was used to select the optimal model using the smallest value. The final values used for the model were `mtry = 11`, `splitrule = variance` and `min.node.size = 5`.*” You can see the results of tuning grid combinations in the associated plot of ROC AUC vs `mtry` grouped by splitting rule.

The bagging and random forest models faired very well - they took over the first and second place!

```
rbind(data.frame(model = "Manual ANOVA", RMSE = round(carseats_model_1_pruned_rmse, 5)),
  data.frame(model = "ANOVA w.tuneGrid", RMSE = round(carseats_model_3_pruned_rmse, 5)),
  data.frame(model = "Bagging", RMSE = round(carseats.bag.rmse, 5)),
  data.frame(model = "Random Forest", RMSE = round(carseats.frst.rmse, 5))
) %>% arrange(RMSE)
```

```
##           model      RMSE
```

```
## 1    Random Forest 1.75811
## 2          Bagging 1.93279
## 3 ANOVA w.tuneGrid 2.29833
## 4    Manual ANOVA 2.38806
```

8.5 Gradient Boosting

Boosting is a method to improve (boost) the weak learners sequentially and increase the model accuracy with a combined model. There are several boosting algorithms. One of the earliest was AdaBoost (adaptive boost). A more recent innovation is gradient boosting.

Adaboost creates a single split tree (decision stump) then weights the observations by how well the initial tree performed, putting more weight on the difficult observations. It then creates a second tree using the weights so that it focuses on the difficult observations. Observations that are difficult to classify receive increasing larger weights until the algorithm identifies a model that correctly classifies them. The final model returns predictions that are a majority vote. (*I think Adaboost applies only to classification problems, not regressions*).

Gradient boosting generalizes the AdaBoost method, so that the object is to minimize a loss function. In the case of classification problems, the loss function is the log-loss; for regression problems, the loss function is mean squared error. The regression trees are additive, so that the successive models can be added together to correct the residuals in the earlier models. Gradient boosting constructs its trees in a “greedy” manner, meaning it chooses the best splits based on purity scores like Gini or minimizing the loss. It is common to constrain the weak learners by setting maximum tree size parameters. Gradient boosting continues until it reaches maximum number of trees or an acceptable error level. This can result in overfitting, so it is common to employ regularization methods that penalize aspects of the model.

Tree Constraints. In general the more constrained the tree, the more trees need to be grown. Parameters to optimize include number of trees, tree depth, number of nodes, minimum observations per split, and minimum improvement to loss.

Learning Rate. Each successive tree can be weighted to slow down the learning rate. Decreasing the learning rate increases the number of required trees. Common growth rates are 0.1 to 0.3.

The gradient boosting algorithm fits a shallow tree T_1 to the data, $M_1 = T_1$. Then it fits a tree T_2 to the residuals and adds a weighted sum of the tree to the original tree as $M_2 = M_1 + \gamma T_2$. For regularized boosting, include a learning rate factor $\eta \in (0..1)$, $M_2 = M_1 + \eta \gamma T_2$. A larger η produces faster learning, but risks overfitting. The process repeats until the residuals are small enough, or until it reaches the maximum iterations. Because overfitting is a risk, use

cross-validation to select the appropriate number of trees (the number of trees producing the lowest RMSE).

8.5.0.1 Gradient Boosting Classification Example

Again using the OJ data set to predict `Purchase`, this time I'll use the gradient boosting method by specifying `method = "gbm"`. I'll use `tuneLength = 5` and not worry about `tuneGrid` anymore. `Caret` tunes the following hyperparameters (see `modelLookup("gbm")`).

- `n.trees`: number of boosting iterations
- `interaction.depth`: maximum tree depth
- `shrinkage`: shrinkage
- `n.minobsinnode`: minimum terminal node size

```
oj.gbm <- train(Purchase ~ .,
  data = oj_train,
  method = "gbm", # for bagged tree
  tuneLength = 5, # choose up to 5 combinations of tuning parameters
  metric = "ROC", # evaluate hyperparameter combinations with ROC
  trControl = trainControl(
    method = "cv", # k-fold cross validation
    number = 10, # 10 folds
    savePredictions = "final", # save predictions for the optimal tuning parameters
    classProbs = TRUE, # return class probabilities in addition to predicted values
    summaryFunction = twoClassSummary # for binary response variable
  )
)
```

## Iter	TrainDeviance	ValidDeviance	StepSize	Improve
## 1	1.2789	nan	0.1000	0.0273
## 2	1.2286	nan	0.1000	0.0245
## 3	1.1929	nan	0.1000	0.0175
## 4	1.1613	nan	0.1000	0.0148
## 5	1.1263	nan	0.1000	0.0146
## 6	1.0991	nan	0.1000	0.0105
## 7	1.0752	nan	0.1000	0.0102
## 8	1.0579	nan	0.1000	0.0087
## 9	1.0433	nan	0.1000	0.0047
## 10	1.0280	nan	0.1000	0.0082
## 20	0.9233	nan	0.1000	0.0026
## 40	0.8226	nan	0.1000	0.0010
## 60	0.7809	nan	0.1000	-0.0001
## 80	0.7595	nan	0.1000	-0.0002

##	100	0.7506	nan	0.1000	-0.0008
##	120	0.7407	nan	0.1000	-0.0005
##	140	0.7317	nan	0.1000	-0.0005
##	160	0.7277	nan	0.1000	-0.0009
##	180	0.7232	nan	0.1000	-0.0004
##	200	0.7181	nan	0.1000	-0.0007
##	220	0.7115	nan	0.1000	-0.0008
##	240	0.7096	nan	0.1000	-0.0010
##	250	0.7081	nan	0.1000	-0.0015

##

##	Iter	TrainDeviance	ValidDeviance	StepSize	Improve
##	1	1.2695	nan	0.1000	0.0319
##	2	1.2150	nan	0.1000	0.0260
##	3	1.1702	nan	0.1000	0.0225
##	4	1.1260	nan	0.1000	0.0186
##	5	1.0913	nan	0.1000	0.0147
##	6	1.0586	nan	0.1000	0.0160
##	7	1.0276	nan	0.1000	0.0146
##	8	1.0045	nan	0.1000	0.0109
##	9	0.9836	nan	0.1000	0.0099
##	10	0.9624	nan	0.1000	0.0068
##	20	0.8337	nan	0.1000	0.0027
##	40	0.7525	nan	0.1000	-0.0005
##	60	0.7240	nan	0.1000	-0.0005
##	80	0.7063	nan	0.1000	-0.0006
##	100	0.6879	nan	0.1000	-0.0011
##	120	0.6751	nan	0.1000	-0.0018
##	140	0.6605	nan	0.1000	-0.0012
##	160	0.6477	nan	0.1000	-0.0013
##	180	0.6359	nan	0.1000	-0.0010
##	200	0.6274	nan	0.1000	-0.0018
##	220	0.6166	nan	0.1000	-0.0005
##	240	0.6078	nan	0.1000	-0.0011
##	250	0.6014	nan	0.1000	-0.0019

##

##	Iter	TrainDeviance	ValidDeviance	StepSize	Improve
##	1	1.2548	nan	0.1000	0.0377
##	2	1.1905	nan	0.1000	0.0294
##	3	1.1343	nan	0.1000	0.0258
##	4	1.0935	nan	0.1000	0.0180
##	5	1.0529	nan	0.1000	0.0168
##	6	1.0172	nan	0.1000	0.0159
##	7	0.9824	nan	0.1000	0.0151
##	8	0.9534	nan	0.1000	0.0127
##	9	0.9277	nan	0.1000	0.0109
##	10	0.9066	nan	0.1000	0.0088

```

##      20      0.7870      nan    0.1000    0.0023
##      40      0.7150      nan    0.1000   -0.0008
##      60      0.6799      nan    0.1000   -0.0023
##      80      0.6520      nan    0.1000   -0.0012
##     100      0.6298      nan    0.1000   -0.0005
##     120      0.6117      nan    0.1000   -0.0024
##     140      0.5973      nan    0.1000   -0.0016
##     160      0.5849      nan    0.1000   -0.0023
##     180      0.5670      nan    0.1000   -0.0015
##     200      0.5548      nan    0.1000   -0.0006
##     220      0.5440      nan    0.1000   -0.0024
##     240      0.5290      nan    0.1000   -0.0020
##     250      0.5228      nan    0.1000   -0.0016
##
## Iter  TrainDeviance  ValidDeviance  StepSize  Improve
##      1          1.2554          nan    0.1000    0.0399
##      2          1.1847          nan    0.1000    0.0346
##      3          1.1321          nan    0.1000    0.0199
##      4          1.0823          nan    0.1000    0.0224
##      5          1.0392          nan    0.1000    0.0208
##      6          1.0067          nan    0.1000    0.0145
##      7          0.9768          nan    0.1000    0.0139
##      8          0.9462          nan    0.1000    0.0123
##      9          0.9238          nan    0.1000    0.0095
##     10          0.8966          nan    0.1000    0.0090
##     20          0.7681          nan    0.1000    0.0007
##     40          0.6937          nan    0.1000   -0.0004
##     60          0.6552          nan    0.1000   -0.0017
##     80          0.6202          nan    0.1000   -0.0018
##    100          0.5887          nan    0.1000   -0.0027
##    120          0.5653          nan    0.1000   -0.0012
##    140          0.5434          nan    0.1000   -0.0017
##    160          0.5275          nan    0.1000   -0.0008
##    180          0.5068          nan    0.1000   -0.0012
##    200          0.4935          nan    0.1000   -0.0016
##    220          0.4801          nan    0.1000   -0.0018
##    240          0.4665          nan    0.1000   -0.0010
##    250          0.4603          nan    0.1000   -0.0012
##
## Iter  TrainDeviance  ValidDeviance  StepSize  Improve
##      1          1.2513          nan    0.1000    0.0375
##      2          1.1795          nan    0.1000    0.0320
##      3          1.1264          nan    0.1000    0.0254
##      4          1.0742          nan    0.1000    0.0225
##      5          1.0282          nan    0.1000    0.0196
##      6          0.9888          nan    0.1000    0.0177

```

##	7	0.9547	nan	0.1000	0.0136
##	8	0.9303	nan	0.1000	0.0103
##	9	0.9008	nan	0.1000	0.0121
##	10	0.8803	nan	0.1000	0.0073
##	20	0.7563	nan	0.1000	0.0003
##	40	0.6715	nan	0.1000	-0.0012
##	60	0.6253	nan	0.1000	-0.0016
##	80	0.5868	nan	0.1000	-0.0021
##	100	0.5538	nan	0.1000	-0.0015
##	120	0.5285	nan	0.1000	-0.0034
##	140	0.5070	nan	0.1000	-0.0025
##	160	0.4872	nan	0.1000	-0.0012
##	180	0.4736	nan	0.1000	-0.0023
##	200	0.4566	nan	0.1000	-0.0015
##	220	0.4407	nan	0.1000	-0.0011
##	240	0.4262	nan	0.1000	-0.0013
##	250	0.4186	nan	0.1000	-0.0024

##	Iter	TrainDeviance	ValidDeviance	StepSize	Improve
##	1	1.2743	nan	0.1000	0.0307
##	2	1.2220	nan	0.1000	0.0240
##	3	1.1885	nan	0.1000	0.0165
##	4	1.1522	nan	0.1000	0.0177
##	5	1.1186	nan	0.1000	0.0136
##	6	1.0912	nan	0.1000	0.0111
##	7	1.0693	nan	0.1000	0.0106
##	8	1.0492	nan	0.1000	0.0089
##	9	1.0309	nan	0.1000	0.0093
##	10	1.0172	nan	0.1000	0.0069
##	20	0.9206	nan	0.1000	0.0030
##	40	0.8357	nan	0.1000	-0.0002
##	60	0.7936	nan	0.1000	-0.0000
##	80	0.7764	nan	0.1000	-0.0009
##	100	0.7682	nan	0.1000	-0.0004
##	120	0.7620	nan	0.1000	-0.0008
##	140	0.7582	nan	0.1000	-0.0011
##	160	0.7536	nan	0.1000	-0.0005
##	180	0.7501	nan	0.1000	-0.0006
##	200	0.7448	nan	0.1000	-0.0008
##	220	0.7409	nan	0.1000	-0.0006
##	240	0.7385	nan	0.1000	-0.0011
##	250	0.7368	nan	0.1000	-0.0007

##	Iter	TrainDeviance	ValidDeviance	StepSize	Improve
##	1	1.2697	nan	0.1000	0.0323
##	2	1.2121	nan	0.1000	0.0276

```

##      3      1.1636      nan      0.1000      0.0251
##      4      1.1220      nan      0.1000      0.0166
##      5      1.0826      nan      0.1000      0.0131
##      6      1.0537      nan      0.1000      0.0134
##      7      1.0269      nan      0.1000      0.0104
##      8      1.0061      nan      0.1000      0.0084
##      9      0.9858      nan      0.1000      0.0082
##     10      0.9678      nan      0.1000      0.0066
##     20      0.8429      nan      0.1000      0.0024
##     40      0.7685      nan      0.1000     -0.0010
##     60      0.7422      nan      0.1000     -0.0006
##     80      0.7228      nan      0.1000     -0.0009
##    100      0.7073      nan      0.1000     -0.0013
##    120      0.6937      nan      0.1000     -0.0024
##    140      0.6836      nan      0.1000     -0.0014
##    160      0.6703      nan      0.1000     -0.0022
##    180      0.6607      nan      0.1000     -0.0009
##    200      0.6529      nan      0.1000     -0.0011
##    220      0.6438      nan      0.1000     -0.0017
##    240      0.6370      nan      0.1000     -0.0015
##    250      0.6311      nan      0.1000     -0.0011
##
## Iter  TrainDeviance  ValidDeviance  StepSize  Improve
##      1      1.2569      nan      0.1000      0.0361
##      2      1.1946      nan      0.1000      0.0301
##      3      1.1386      nan      0.1000      0.0266
##      4      1.0954      nan      0.1000      0.0205
##      5      1.0524      nan      0.1000      0.0204
##      6      1.0186      nan      0.1000      0.0149
##      7      0.9847      nan      0.1000      0.0126
##      8      0.9618      nan      0.1000      0.0086
##      9      0.9344      nan      0.1000      0.0114
##     10      0.9135      nan      0.1000      0.0095
##     20      0.8003      nan      0.1000      0.0027
##     40      0.7353      nan      0.1000     -0.0011
##     60      0.7042      nan      0.1000     -0.0027
##     80      0.6800      nan      0.1000     -0.0017
##    100      0.6602      nan      0.1000     -0.0008
##    120      0.6393      nan      0.1000     -0.0017
##    140      0.6231      nan      0.1000     -0.0016
##    160      0.6077      nan      0.1000     -0.0028
##    180      0.5977      nan      0.1000     -0.0012
##    200      0.5863      nan      0.1000     -0.0014
##    220      0.5749      nan      0.1000     -0.0013
##    240      0.5618      nan      0.1000     -0.0022
##    250      0.5577      nan      0.1000     -0.0022

```

```

##
## Iter    TrainDeviance    ValidDeviance    StepSize    Improve
##      1          1.2495          nan      0.1000      0.0419
##      2          1.1884          nan      0.1000      0.0265
##      3          1.1279          nan      0.1000      0.0248
##      4          1.0820          nan      0.1000      0.0208
##      5          1.0426          nan      0.1000      0.0166
##      6          1.0089          nan      0.1000      0.0146
##      7          0.9799          nan      0.1000      0.0126
##      8          0.9474          nan      0.1000      0.0140
##      9          0.9225          nan      0.1000      0.0079
##     10          0.9066          nan      0.1000      0.0048
##     20          0.7815          nan      0.1000      0.0014
##     40          0.7028          nan      0.1000     -0.0019
##     60          0.6661          nan      0.1000     -0.0011
##     80          0.6386          nan      0.1000     -0.0006
##    100          0.6075          nan      0.1000     -0.0005
##    120          0.5861          nan      0.1000     -0.0019
##    140          0.5674          nan      0.1000     -0.0020
##    160          0.5467          nan      0.1000     -0.0016
##    180          0.5318          nan      0.1000     -0.0020
##    200          0.5200          nan      0.1000     -0.0025
##    220          0.5050          nan      0.1000     -0.0009
##    240          0.4930          nan      0.1000     -0.0020
##    250          0.4883          nan      0.1000     -0.0016
##
## Iter    TrainDeviance    ValidDeviance    StepSize    Improve
##      1          1.2558          nan      0.1000      0.0336
##      2          1.1856          nan      0.1000      0.0326
##      3          1.1172          nan      0.1000      0.0305
##      4          1.0713          nan      0.1000      0.0222
##      5          1.0313          nan      0.1000      0.0171
##      6          0.9965          nan      0.1000      0.0164
##      7          0.9613          nan      0.1000      0.0156
##      8          0.9354          nan      0.1000      0.0103
##      9          0.9089          nan      0.1000      0.0111
##     10          0.8859          nan      0.1000      0.0059
##     20          0.7690          nan      0.1000      0.0006
##     40          0.6889          nan      0.1000     -0.0004
##     60          0.6452          nan      0.1000     -0.0021
##     80          0.6127          nan      0.1000     -0.0021
##    100          0.5811          nan      0.1000     -0.0032
##    120          0.5557          nan      0.1000     -0.0011
##    140          0.5332          nan      0.1000     -0.0012
##    160          0.5118          nan      0.1000     -0.0014
##    180          0.4879          nan      0.1000     -0.0012

```



```

##      200      0.4737      nan      0.1000     -0.0020
##      220      0.4591      nan      0.1000     -0.0024
##      240      0.4460      nan      0.1000     -0.0030
##      250      0.4386      nan      0.1000     -0.0013
##
## Iter   TrainDeviance   ValidDeviance   StepSize   Improve
##      1          1.2717          nan      0.1000     0.0320
##      2          1.2209          nan      0.1000     0.0250
##      3          1.1751          nan      0.1000     0.0208
##      4          1.1404          nan      0.1000     0.0149
##      5          1.1056          nan      0.1000     0.0119
##      6          1.0782          nan      0.1000     0.0125
##      7          1.0569          nan      0.1000     0.0081
##      8          1.0356          nan      0.1000     0.0098
##      9          1.0176          nan      0.1000     0.0080
##     10          1.0042          nan      0.1000     0.0066
##     20          0.9023          nan      0.1000     0.0021
##     40          0.8156          nan      0.1000     0.0006
##     60          0.7793          nan      0.1000     0.0004
##     80          0.7609          nan      0.1000    -0.0014
##    100          0.7514          nan      0.1000    -0.0006
##    120          0.7448          nan      0.1000    -0.0005
##    140          0.7406          nan      0.1000    -0.0008
##    160          0.7353          nan      0.1000    -0.0007
##    180          0.7329          nan      0.1000    -0.0008
##    200          0.7281          nan      0.1000    -0.0014
##    220          0.7239          nan      0.1000    -0.0008
##    240          0.7211          nan      0.1000    -0.0011
##    250          0.7203          nan      0.1000    -0.0010
##
## Iter   TrainDeviance   ValidDeviance   StepSize   Improve
##      1          1.2647          nan      0.1000     0.0368
##      2          1.2060          nan      0.1000     0.0275
##      3          1.1558          nan      0.1000     0.0232
##      4          1.1207          nan      0.1000     0.0178
##      5          1.0787          nan      0.1000     0.0172
##      6          1.0478          nan      0.1000     0.0141
##      7          1.0178          nan      0.1000     0.0117
##      8          0.9963          nan      0.1000     0.0099
##      9          0.9749          nan      0.1000     0.0107
##     10          0.9514          nan      0.1000     0.0095
##     20          0.8341          nan      0.1000     0.0029
##     40          0.7601          nan      0.1000    -0.0015
##     60          0.7335          nan      0.1000    -0.0010
##     80          0.7131          nan      0.1000    -0.0011
##    100          0.7006          nan      0.1000    -0.0018

```

##	120	0.6886	nan	0.1000	-0.0008
##	140	0.6740	nan	0.1000	-0.0014
##	160	0.6628	nan	0.1000	-0.0017
##	180	0.6522	nan	0.1000	-0.0010
##	200	0.6423	nan	0.1000	-0.0005
##	220	0.6353	nan	0.1000	-0.0020
##	240	0.6251	nan	0.1000	-0.0017
##	250	0.6211	nan	0.1000	-0.0013

##

##	Iter	TrainDeviance	ValidDeviance	StepSize	Improve
##	1	1.2533	nan	0.1000	0.0369
##	2	1.1914	nan	0.1000	0.0304
##	3	1.1304	nan	0.1000	0.0280
##	4	1.0808	nan	0.1000	0.0214
##	5	1.0411	nan	0.1000	0.0177
##	6	1.0095	nan	0.1000	0.0131
##	7	0.9791	nan	0.1000	0.0118
##	8	0.9487	nan	0.1000	0.0146
##	9	0.9234	nan	0.1000	0.0114
##	10	0.9050	nan	0.1000	0.0081
##	20	0.7868	nan	0.1000	-0.0018
##	40	0.7190	nan	0.1000	-0.0002
##	60	0.6929	nan	0.1000	-0.0013
##	80	0.6706	nan	0.1000	-0.0012
##	100	0.6511	nan	0.1000	-0.0008
##	120	0.6313	nan	0.1000	-0.0028
##	140	0.6161	nan	0.1000	-0.0010
##	160	0.6016	nan	0.1000	-0.0019
##	180	0.5875	nan	0.1000	-0.0013
##	200	0.5754	nan	0.1000	-0.0021
##	220	0.5613	nan	0.1000	-0.0011
##	240	0.5456	nan	0.1000	-0.0014
##	250	0.5423	nan	0.1000	-0.0016

##

##	Iter	TrainDeviance	ValidDeviance	StepSize	Improve
##	1	1.2526	nan	0.1000	0.0406
##	2	1.1845	nan	0.1000	0.0346
##	3	1.1309	nan	0.1000	0.0246
##	4	1.0809	nan	0.1000	0.0250
##	5	1.0380	nan	0.1000	0.0205
##	6	1.0017	nan	0.1000	0.0164
##	7	0.9661	nan	0.1000	0.0146
##	8	0.9372	nan	0.1000	0.0130
##	9	0.9115	nan	0.1000	0.0105
##	10	0.8921	nan	0.1000	0.0084
##	20	0.7698	nan	0.1000	0.0014

```

##      40      0.6902      nan      0.1000     -0.0016
##      60      0.6501      nan      0.1000     -0.0019
##      80      0.6200      nan      0.1000     -0.0006
##     100      0.5995      nan      0.1000     -0.0026
##     120      0.5766      nan      0.1000     -0.0019
##     140      0.5576      nan      0.1000     -0.0020
##     160      0.5428      nan      0.1000     -0.0027
##     180      0.5276      nan      0.1000     -0.0026
##     200      0.5091      nan      0.1000     -0.0011
##     220      0.4954      nan      0.1000     -0.0014
##     240      0.4801      nan      0.1000     -0.0028
##     250      0.4748      nan      0.1000     -0.0021
##
## Iter   TrainDeviance   ValidDeviance   StepSize   Improve
##      1           1.2578           nan      0.1000    0.0379
##      2           1.1876           nan      0.1000    0.0342
##      3           1.1214           nan      0.1000    0.0297
##      4           1.0650           nan      0.1000    0.0254
##      5           1.0182           nan      0.1000    0.0188
##      6           0.9812           nan      0.1000    0.0172
##      7           0.9484           nan      0.1000    0.0116
##      8           0.9182           nan      0.1000    0.0116
##      9           0.8929           nan      0.1000    0.0068
##     10           0.8704           nan      0.1000    0.0085
##     20           0.7522           nan      0.1000    0.0002
##     40           0.6778           nan      0.1000   -0.0019
##     60           0.6318           nan      0.1000   -0.0015
##     80           0.5982           nan      0.1000   -0.0011
##    100           0.5669           nan      0.1000   -0.0032
##    120           0.5451           nan      0.1000   -0.0012
##    140           0.5224           nan      0.1000   -0.0002
##    160           0.5005           nan      0.1000   -0.0027
##    180           0.4834           nan      0.1000   -0.0015
##    200           0.4693           nan      0.1000   -0.0009
##    220           0.4530           nan      0.1000   -0.0016
##    240           0.4419           nan      0.1000   -0.0035
##    250           0.4324           nan      0.1000   -0.0016
##
## Iter   TrainDeviance   ValidDeviance   StepSize   Improve
##      1           1.2761           nan      0.1000    0.0283
##      2           1.2243           nan      0.1000    0.0231
##      3           1.1816           nan      0.1000    0.0192
##      4           1.1477           nan      0.1000    0.0158
##      5           1.1174           nan      0.1000    0.0122
##      6           1.0930           nan      0.1000    0.0116
##      7           1.0704           nan      0.1000    0.0102

```

##	8	1.0499	nan	0.1000	0.0083
##	9	1.0340	nan	0.1000	0.0078
##	10	1.0167	nan	0.1000	0.0082
##	20	0.9152	nan	0.1000	0.0021
##	40	0.8226	nan	0.1000	0.0007
##	60	0.7895	nan	0.1000	0.0000
##	80	0.7690	nan	0.1000	-0.0008
##	100	0.7612	nan	0.1000	-0.0010
##	120	0.7541	nan	0.1000	-0.0004
##	140	0.7491	nan	0.1000	-0.0012
##	160	0.7443	nan	0.1000	-0.0006
##	180	0.7405	nan	0.1000	-0.0009
##	200	0.7369	nan	0.1000	-0.0009
##	220	0.7329	nan	0.1000	-0.0007
##	240	0.7287	nan	0.1000	-0.0014
##	250	0.7268	nan	0.1000	-0.0011
##					
##	Iter	TrainDeviance	ValidDeviance	StepSize	Improve
##	1	1.2642	nan	0.1000	0.0359
##	2	1.2066	nan	0.1000	0.0279
##	3	1.1624	nan	0.1000	0.0237
##	4	1.1213	nan	0.1000	0.0202
##	5	1.0853	nan	0.1000	0.0162
##	6	1.0560	nan	0.1000	0.0116
##	7	1.0275	nan	0.1000	0.0125
##	8	1.0040	nan	0.1000	0.0109
##	9	0.9823	nan	0.1000	0.0077
##	10	0.9612	nan	0.1000	0.0105
##	20	0.8409	nan	0.1000	0.0026
##	40	0.7578	nan	0.1000	-0.0007
##	60	0.7294	nan	0.1000	-0.0007
##	80	0.7095	nan	0.1000	-0.0026
##	100	0.6965	nan	0.1000	-0.0012
##	120	0.6857	nan	0.1000	-0.0022
##	140	0.6751	nan	0.1000	-0.0004
##	160	0.6650	nan	0.1000	-0.0018
##	180	0.6581	nan	0.1000	-0.0017
##	200	0.6520	nan	0.1000	-0.0009
##	220	0.6436	nan	0.1000	-0.0011
##	240	0.6351	nan	0.1000	-0.0009
##	250	0.6312	nan	0.1000	-0.0016
##					
##	Iter	TrainDeviance	ValidDeviance	StepSize	Improve
##	1	1.2577	nan	0.1000	0.0367
##	2	1.1908	nan	0.1000	0.0288
##	3	1.1314	nan	0.1000	0.0257

```

##      4      1.0832      nan      0.1000      0.0231
##      5      1.0473      nan      0.1000      0.0147
##      6      1.0104      nan      0.1000      0.0174
##      7      0.9743      nan      0.1000      0.0139
##      8      0.9460      nan      0.1000      0.0113
##      9      0.9236      nan      0.1000      0.0105
##     10      0.9026      nan      0.1000      0.0077
##     20      0.7942      nan      0.1000      0.0009
##     40      0.7292      nan      0.1000     -0.0018
##     60      0.6933      nan      0.1000     -0.0013
##     80      0.6650      nan      0.1000     -0.0008
##    100      0.6458      nan      0.1000     -0.0020
##    120      0.6252      nan      0.1000     -0.0018
##    140      0.6106      nan      0.1000     -0.0010
##    160      0.5953      nan      0.1000     -0.0009
##    180      0.5810      nan      0.1000     -0.0014
##    200      0.5683      nan      0.1000     -0.0016
##    220      0.5544      nan      0.1000     -0.0009
##    240      0.5425      nan      0.1000     -0.0010
##    250      0.5367      nan      0.1000     -0.0012
##
## Iter  TrainDeviance  ValidDeviance  StepSize  Improve
##      1      1.2542      nan      0.1000      0.0373
##      2      1.1874      nan      0.1000      0.0310
##      3      1.1277      nan      0.1000      0.0279
##      4      1.0765      nan      0.1000      0.0231
##      5      1.0393      nan      0.1000      0.0179
##      6      1.0012      nan      0.1000      0.0141
##      7      0.9658      nan      0.1000      0.0139
##      8      0.9421      nan      0.1000      0.0101
##      9      0.9187      nan      0.1000      0.0088
##     10      0.8966      nan      0.1000      0.0095
##     20      0.7715      nan      0.1000      0.0019
##     40      0.6966      nan      0.1000     -0.0017
##     60      0.6516      nan      0.1000     -0.0010
##     80      0.6206      nan      0.1000     -0.0015
##    100      0.5991      nan      0.1000     -0.0028
##    120      0.5818      nan      0.1000     -0.0019
##    140      0.5644      nan      0.1000     -0.0018
##    160      0.5476      nan      0.1000     -0.0016
##    180      0.5329      nan      0.1000     -0.0016
##    200      0.5212      nan      0.1000     -0.0016
##    220      0.5055      nan      0.1000     -0.0032
##    240      0.4926      nan      0.1000     -0.0013
##    250      0.4850      nan      0.1000     -0.0006
##

```

##	Iter	TrainDeviance	ValidDeviance	StepSize	Improve
##	1	1.2523	nan	0.1000	0.0397
##	2	1.1844	nan	0.1000	0.0301
##	3	1.1269	nan	0.1000	0.0263
##	4	1.0773	nan	0.1000	0.0216
##	5	1.0310	nan	0.1000	0.0211
##	6	0.9902	nan	0.1000	0.0176
##	7	0.9582	nan	0.1000	0.0119
##	8	0.9338	nan	0.1000	0.0106
##	9	0.9054	nan	0.1000	0.0120
##	10	0.8869	nan	0.1000	0.0083
##	20	0.7620	nan	0.1000	-0.0001
##	40	0.6734	nan	0.1000	-0.0010
##	60	0.6333	nan	0.1000	-0.0004
##	80	0.5982	nan	0.1000	-0.0023
##	100	0.5685	nan	0.1000	-0.0013
##	120	0.5436	nan	0.1000	-0.0024
##	140	0.5196	nan	0.1000	-0.0012
##	160	0.5010	nan	0.1000	-0.0031
##	180	0.4832	nan	0.1000	-0.0032
##	200	0.4677	nan	0.1000	-0.0014
##	220	0.4504	nan	0.1000	-0.0016
##	240	0.4333	nan	0.1000	-0.0013
##	250	0.4273	nan	0.1000	-0.0022
##					
##	Iter	TrainDeviance	ValidDeviance	StepSize	Improve
##	1	1.2726	nan	0.1000	0.0321
##	2	1.2186	nan	0.1000	0.0259
##	3	1.1770	nan	0.1000	0.0190
##	4	1.1350	nan	0.1000	0.0186
##	5	1.1017	nan	0.1000	0.0147
##	6	1.0810	nan	0.1000	0.0096
##	7	1.0570	nan	0.1000	0.0105
##	8	1.0375	nan	0.1000	0.0098
##	9	1.0188	nan	0.1000	0.0081
##	10	1.0020	nan	0.1000	0.0072
##	20	0.9034	nan	0.1000	0.0023
##	40	0.8070	nan	0.1000	-0.0001
##	60	0.7671	nan	0.1000	-0.0003
##	80	0.7500	nan	0.1000	-0.0007
##	100	0.7415	nan	0.1000	-0.0008
##	120	0.7349	nan	0.1000	-0.0008
##	140	0.7282	nan	0.1000	-0.0006
##	160	0.7222	nan	0.1000	-0.0006
##	180	0.7190	nan	0.1000	-0.0003
##	200	0.7160	nan	0.1000	-0.0013

```

##      220      0.7128      nan      0.1000     -0.0008
##      240      0.7092      nan      0.1000     -0.0012
##      250      0.7066      nan      0.1000     -0.0005
##
## Iter   TrainDeviance   ValidDeviance   StepSize   Improve
##      1          1.2597          nan      0.1000     0.0362
##      2          1.2019          nan      0.1000     0.0279
##      3          1.1504          nan      0.1000     0.0230
##      4          1.1084          nan      0.1000     0.0168
##      5          1.0702          nan      0.1000     0.0184
##      6          1.0378          nan      0.1000     0.0136
##      7          1.0090          nan      0.1000     0.0131
##      8          0.9848          nan      0.1000     0.0115
##      9          0.9632          nan      0.1000     0.0078
##     10          0.9410          nan      0.1000     0.0094
##     20          0.8215          nan      0.1000     0.0037
##     40          0.7478          nan      0.1000    -0.0001
##     60          0.7184          nan      0.1000    -0.0008
##     80          0.6985          nan      0.1000    -0.0007
##    100          0.6799          nan      0.1000    -0.0013
##    120          0.6660          nan      0.1000     0.0002
##    140          0.6534          nan      0.1000    -0.0009
##    160          0.6423          nan      0.1000    -0.0016
##    180          0.6325          nan      0.1000    -0.0016
##    200          0.6226          nan      0.1000    -0.0011
##    220          0.6139          nan      0.1000    -0.0010
##    240          0.6043          nan      0.1000    -0.0017
##    250          0.6016          nan      0.1000    -0.0012
##
## Iter   TrainDeviance   ValidDeviance   StepSize   Improve
##      1          1.2549          nan      0.1000     0.0413
##      2          1.1843          nan      0.1000     0.0345
##      3          1.1274          nan      0.1000     0.0245
##      4          1.0809          nan      0.1000     0.0213
##      5          1.0418          nan      0.1000     0.0178
##      6          1.0056          nan      0.1000     0.0128
##      7          0.9797          nan      0.1000     0.0122
##      8          0.9562          nan      0.1000     0.0083
##      9          0.9288          nan      0.1000     0.0116
##     10          0.9070          nan      0.1000     0.0088
##     20          0.7819          nan      0.1000     0.0013
##     40          0.7119          nan      0.1000    -0.0008
##     60          0.6706          nan      0.1000    -0.0013
##     80          0.6482          nan      0.1000    -0.0018
##    100          0.6271          nan      0.1000    -0.0027
##    120          0.6065          nan      0.1000    -0.0013

```

##	140	0.5901	nan	0.1000	-0.0021
##	160	0.5709	nan	0.1000	-0.0009
##	180	0.5563	nan	0.1000	-0.0005
##	200	0.5427	nan	0.1000	-0.0025
##	220	0.5333	nan	0.1000	-0.0012
##	240	0.5239	nan	0.1000	-0.0015
##	250	0.5186	nan	0.1000	-0.0018

##	Iter	TrainDeviance	ValidDeviance	StepSize	Improve
##	1	1.2503	nan	0.1000	0.0405
##	2	1.1806	nan	0.1000	0.0343
##	3	1.1211	nan	0.1000	0.0276
##	4	1.0676	nan	0.1000	0.0242
##	5	1.0226	nan	0.1000	0.0193
##	6	0.9842	nan	0.1000	0.0164
##	7	0.9549	nan	0.1000	0.0120
##	8	0.9274	nan	0.1000	0.0125
##	9	0.8994	nan	0.1000	0.0100
##	10	0.8810	nan	0.1000	0.0068
##	20	0.7641	nan	0.1000	0.0015
##	40	0.6902	nan	0.1000	-0.0013
##	60	0.6497	nan	0.1000	-0.0026
##	80	0.6225	nan	0.1000	-0.0020
##	100	0.5898	nan	0.1000	-0.0009
##	120	0.5662	nan	0.1000	-0.0014
##	140	0.5460	nan	0.1000	-0.0015
##	160	0.5293	nan	0.1000	-0.0029
##	180	0.5142	nan	0.1000	-0.0018
##	200	0.5021	nan	0.1000	-0.0021
##	220	0.4850	nan	0.1000	-0.0025
##	240	0.4738	nan	0.1000	-0.0010
##	250	0.4697	nan	0.1000	-0.0025

##	Iter	TrainDeviance	ValidDeviance	StepSize	Improve
##	1	1.2510	nan	0.1000	0.0400
##	2	1.1788	nan	0.1000	0.0356
##	3	1.1185	nan	0.1000	0.0284
##	4	1.0652	nan	0.1000	0.0255
##	5	1.0228	nan	0.1000	0.0209
##	6	0.9819	nan	0.1000	0.0175
##	7	0.9467	nan	0.1000	0.0146
##	8	0.9149	nan	0.1000	0.0133
##	9	0.8872	nan	0.1000	0.0105
##	10	0.8681	nan	0.1000	0.0067
##	20	0.7435	nan	0.1000	0.0013
##	40	0.6620	nan	0.1000	-0.0012


```

##      60      0.6096      nan      0.1000     -0.0011
##      80      0.5775      nan      0.1000     -0.0017
##     100      0.5491      nan      0.1000     -0.0016
##     120      0.5248      nan      0.1000     -0.0030
##     140      0.5082      nan      0.1000     -0.0016
##     160      0.4873      nan      0.1000     -0.0022
##     180      0.4700      nan      0.1000     -0.0020
##     200      0.4543      nan      0.1000     -0.0013
##     220      0.4388      nan      0.1000     -0.0031
##     240      0.4269      nan      0.1000     -0.0013
##     250      0.4237      nan      0.1000     -0.0015
##
## Iter   TrainDeviance   ValidDeviance   StepSize   Improve
##      1          1.2737           nan      0.1000    0.0306
##      2          1.2188           nan      0.1000    0.0241
##      3          1.1752           nan      0.1000    0.0184
##      4          1.1384           nan      0.1000    0.0168
##      5          1.1052           nan      0.1000    0.0134
##      6          1.0800           nan      0.1000    0.0103
##      7          1.0601           nan      0.1000    0.0106
##      8          1.0416           nan      0.1000    0.0097
##      9          1.0234           nan      0.1000    0.0070
##     10          1.0079           nan      0.1000    0.0055
##     20          0.9090           nan      0.1000    0.0034
##     40          0.8145           nan      0.1000    0.0004
##     60          0.7708           nan      0.1000    0.0003
##     80          0.7528           nan      0.1000   -0.0009
##    100          0.7433           nan      0.1000   -0.0010
##    120          0.7366           nan      0.1000   -0.0002
##    140          0.7317           nan      0.1000   -0.0007
##    160          0.7262           nan      0.1000   -0.0009
##    180          0.7217           nan      0.1000   -0.0005
##    200          0.7172           nan      0.1000   -0.0013
##    220          0.7141           nan      0.1000   -0.0012
##    240          0.7102           nan      0.1000   -0.0002
##    250          0.7086           nan      0.1000   -0.0012
##
## Iter   TrainDeviance   ValidDeviance   StepSize   Improve
##      1          1.2660           nan      0.1000    0.0361
##      2          1.2076           nan      0.1000    0.0275
##      3          1.1575           nan      0.1000    0.0247
##      4          1.1197           nan      0.1000    0.0178
##      5          1.0845           nan      0.1000    0.0173
##      6          1.0529           nan      0.1000    0.0131
##      7          1.0259           nan      0.1000    0.0123
##      8          0.9969           nan      0.1000    0.0117

```

##	9	0.9748	nan	0.1000	0.0082
##	10	0.9535	nan	0.1000	0.0076
##	20	0.8305	nan	0.1000	0.0033
##	40	0.7475	nan	0.1000	-0.0002
##	60	0.7207	nan	0.1000	-0.0018
##	80	0.7037	nan	0.1000	-0.0017
##	100	0.6882	nan	0.1000	-0.0014
##	120	0.6802	nan	0.1000	-0.0010
##	140	0.6709	nan	0.1000	-0.0015
##	160	0.6613	nan	0.1000	-0.0014
##	180	0.6523	nan	0.1000	-0.0004
##	200	0.6449	nan	0.1000	-0.0011
##	220	0.6367	nan	0.1000	-0.0014
##	240	0.6286	nan	0.1000	-0.0013
##	250	0.6233	nan	0.1000	-0.0019

##	Iter	TrainDeviance	ValidDeviance	StepSize	Improve
##	1	1.2551	nan	0.1000	0.0346
##	2	1.1931	nan	0.1000	0.0315
##	3	1.1316	nan	0.1000	0.0272
##	4	1.0872	nan	0.1000	0.0204
##	5	1.0437	nan	0.1000	0.0187
##	6	1.0098	nan	0.1000	0.0139
##	7	0.9818	nan	0.1000	0.0107
##	8	0.9579	nan	0.1000	0.0111
##	9	0.9321	nan	0.1000	0.0129
##	10	0.9121	nan	0.1000	0.0066
##	20	0.7838	nan	0.1000	0.0022
##	40	0.7153	nan	0.1000	-0.0001
##	60	0.6832	nan	0.1000	-0.0011
##	80	0.6612	nan	0.1000	-0.0019
##	100	0.6415	nan	0.1000	-0.0020
##	120	0.6252	nan	0.1000	-0.0020
##	140	0.6084	nan	0.1000	-0.0017
##	160	0.5936	nan	0.1000	-0.0009
##	180	0.5834	nan	0.1000	-0.0021
##	200	0.5669	nan	0.1000	-0.0009
##	220	0.5568	nan	0.1000	-0.0011
##	240	0.5467	nan	0.1000	-0.0027
##	250	0.5422	nan	0.1000	-0.0015

##	Iter	TrainDeviance	ValidDeviance	StepSize	Improve
##	1	1.2574	nan	0.1000	0.0391
##	2	1.1858	nan	0.1000	0.0330
##	3	1.1308	nan	0.1000	0.0237
##	4	1.0768	nan	0.1000	0.0205

```

##      5      1.0311      nan      0.1000      0.0203
##      6      0.9934      nan      0.1000      0.0175
##      7      0.9643      nan      0.1000      0.0121
##      8      0.9342      nan      0.1000      0.0118
##      9      0.9061      nan      0.1000      0.0104
##     10      0.8869      nan      0.1000      0.0078
##     20      0.7608      nan      0.1000      0.0000
##     40      0.6861      nan      0.1000     -0.0012
##     60      0.6452      nan      0.1000     -0.0007
##     80      0.6147      nan      0.1000     -0.0020
##    100      0.5919      nan      0.1000     -0.0015
##    120      0.5685      nan      0.1000     -0.0015
##    140      0.5516      nan      0.1000     -0.0007
##    160      0.5322      nan      0.1000     -0.0024
##    180      0.5188      nan      0.1000     -0.0013
##    200      0.5045      nan      0.1000     -0.0011
##    220      0.4913      nan      0.1000     -0.0012
##    240      0.4791      nan      0.1000     -0.0021
##    250      0.4735      nan      0.1000     -0.0011
##
## Iter   TrainDeviance   ValidDeviance   StepSize   Improve
##      1      1.2526      nan      0.1000      0.0413
##      2      1.1759      nan      0.1000      0.0342
##      3      1.1130      nan      0.1000      0.0286
##      4      1.0692      nan      0.1000      0.0198
##      5      1.0244      nan      0.1000      0.0208
##      6      0.9858      nan      0.1000      0.0160
##      7      0.9512      nan      0.1000      0.0156
##      8      0.9230      nan      0.1000      0.0086
##      9      0.8945      nan      0.1000      0.0128
##     10      0.8741      nan      0.1000      0.0071
##     20      0.7489      nan      0.1000      0.0012
##     40      0.6645      nan      0.1000     -0.0017
##     60      0.6174      nan      0.1000     -0.0018
##     80      0.5875      nan      0.1000     -0.0022
##    100      0.5620      nan      0.1000     -0.0009
##    120      0.5374      nan      0.1000     -0.0026
##    140      0.5202      nan      0.1000     -0.0017
##    160      0.5003      nan      0.1000     -0.0014
##    180      0.4829      nan      0.1000     -0.0021
##    200      0.4660      nan      0.1000     -0.0024
##    220      0.4560      nan      0.1000     -0.0024
##    240      0.4425      nan      0.1000     -0.0020
##    250      0.4349      nan      0.1000     -0.0021
##
## Iter   TrainDeviance   ValidDeviance   StepSize   Improve

```

##	1	1.2722	nan	0.1000	0.0322
##	2	1.2256	nan	0.1000	0.0203
##	3	1.1806	nan	0.1000	0.0216
##	4	1.1444	nan	0.1000	0.0179
##	5	1.1141	nan	0.1000	0.0147
##	6	1.0890	nan	0.1000	0.0118
##	7	1.0707	nan	0.1000	0.0085
##	8	1.0493	nan	0.1000	0.0108
##	9	1.0332	nan	0.1000	0.0069
##	10	1.0157	nan	0.1000	0.0086
##	20	0.9131	nan	0.1000	0.0039
##	40	0.8200	nan	0.1000	0.0010
##	60	0.7820	nan	0.1000	0.0004
##	80	0.7654	nan	0.1000	-0.0014
##	100	0.7561	nan	0.1000	-0.0009
##	120	0.7476	nan	0.1000	-0.0011
##	140	0.7412	nan	0.1000	-0.0003
##	160	0.7365	nan	0.1000	-0.0004
##	180	0.7340	nan	0.1000	-0.0005
##	200	0.7299	nan	0.1000	-0.0007
##	220	0.7266	nan	0.1000	-0.0008
##	240	0.7237	nan	0.1000	-0.0008
##	250	0.7234	nan	0.1000	-0.0005
##					
##	Iter	TrainDeviance	ValidDeviance	StepSize	Improve
##	1	1.2651	nan	0.1000	0.0334
##	2	1.2070	nan	0.1000	0.0269
##	3	1.1557	nan	0.1000	0.0231
##	4	1.1153	nan	0.1000	0.0171
##	5	1.0815	nan	0.1000	0.0155
##	6	1.0482	nan	0.1000	0.0152
##	7	1.0172	nan	0.1000	0.0130
##	8	0.9922	nan	0.1000	0.0112
##	9	0.9735	nan	0.1000	0.0078
##	10	0.9545	nan	0.1000	0.0083
##	20	0.8345	nan	0.1000	0.0033
##	40	0.7573	nan	0.1000	-0.0001
##	60	0.7333	nan	0.1000	-0.0011
##	80	0.7105	nan	0.1000	-0.0006
##	100	0.6966	nan	0.1000	-0.0011
##	120	0.6815	nan	0.1000	-0.0012
##	140	0.6641	nan	0.1000	-0.0008
##	160	0.6482	nan	0.1000	-0.0016
##	180	0.6413	nan	0.1000	-0.0012
##	200	0.6329	nan	0.1000	-0.0008
##	220	0.6240	nan	0.1000	-0.0015

```

##      240      0.6160      nan      0.1000     -0.0011
##      250      0.6111      nan      0.1000     -0.0017
##
## Iter   TrainDeviance   ValidDeviance   StepSize   Improve
##      1           1.2637           nan      0.1000     0.0352
##      2           1.1913           nan      0.1000     0.0330
##      3           1.1333           nan      0.1000     0.0254
##      4           1.0844           nan      0.1000     0.0231
##      5           1.0391           nan      0.1000     0.0195
##      6           1.0021           nan      0.1000     0.0152
##      7           0.9729           nan      0.1000     0.0118
##      8           0.9513           nan      0.1000     0.0086
##      9           0.9272           nan      0.1000     0.0096
##     10           0.9064           nan      0.1000     0.0085
##     20           0.7931           nan      0.1000     0.0041
##     40           0.7237           nan      0.1000    -0.0004
##     60           0.6981           nan      0.1000    -0.0012
##     80           0.6718           nan      0.1000    -0.0026
##    100           0.6529           nan      0.1000    -0.0016
##    120           0.6396           nan      0.1000    -0.0010
##    140           0.6233           nan      0.1000    -0.0013
##    160           0.6101           nan      0.1000    -0.0011
##    180           0.5932           nan      0.1000    -0.0015
##    200           0.5794           nan      0.1000    -0.0010
##    220           0.5645           nan      0.1000    -0.0011
##    240           0.5498           nan      0.1000    -0.0018
##    250           0.5448           nan      0.1000    -0.0012
##
## Iter   TrainDeviance   ValidDeviance   StepSize   Improve
##      1           1.2583           nan      0.1000     0.0380
##      2           1.1835           nan      0.1000     0.0345
##      3           1.1231           nan      0.1000     0.0268
##      4           1.0730           nan      0.1000     0.0218
##      5           1.0329           nan      0.1000     0.0173
##      6           0.9952           nan      0.1000     0.0161
##      7           0.9630           nan      0.1000     0.0130
##      8           0.9348           nan      0.1000     0.0126
##      9           0.9109           nan      0.1000     0.0108
##     10           0.8940           nan      0.1000     0.0058
##     20           0.7714           nan      0.1000     0.0009
##     40           0.6915           nan      0.1000    -0.0022
##     60           0.6527           nan      0.1000    -0.0006
##     80           0.6232           nan      0.1000    -0.0011
##    100           0.6008           nan      0.1000    -0.0018
##    120           0.5783           nan      0.1000    -0.0017
##    140           0.5562           nan      0.1000    -0.0023

```

##	160	0.5375	nan	0.1000	-0.0014
##	180	0.5165	nan	0.1000	-0.0019
##	200	0.5043	nan	0.1000	-0.0021
##	220	0.4907	nan	0.1000	-0.0022
##	240	0.4783	nan	0.1000	-0.0019
##	250	0.4728	nan	0.1000	-0.0022

##

##	Iter	TrainDeviance	ValidDeviance	StepSize	Improve
##	1	1.2435	nan	0.1000	0.0421
##	2	1.1748	nan	0.1000	0.0315
##	3	1.1147	nan	0.1000	0.0294
##	4	1.0672	nan	0.1000	0.0236
##	5	1.0265	nan	0.1000	0.0188
##	6	0.9873	nan	0.1000	0.0165
##	7	0.9555	nan	0.1000	0.0151
##	8	0.9289	nan	0.1000	0.0100
##	9	0.9030	nan	0.1000	0.0085
##	10	0.8836	nan	0.1000	0.0077
##	20	0.7615	nan	0.1000	0.0016
##	40	0.6861	nan	0.1000	-0.0016
##	60	0.6393	nan	0.1000	-0.0017
##	80	0.6022	nan	0.1000	-0.0021
##	100	0.5787	nan	0.1000	-0.0023
##	120	0.5532	nan	0.1000	-0.0024
##	140	0.5303	nan	0.1000	-0.0017
##	160	0.5079	nan	0.1000	-0.0015
##	180	0.4894	nan	0.1000	-0.0017
##	200	0.4716	nan	0.1000	-0.0019
##	220	0.4537	nan	0.1000	-0.0017
##	240	0.4389	nan	0.1000	-0.0007
##	250	0.4327	nan	0.1000	-0.0013

##

##	Iter	TrainDeviance	ValidDeviance	StepSize	Improve
##	1	1.2786	nan	0.1000	0.0297
##	2	1.2244	nan	0.1000	0.0249
##	3	1.1873	nan	0.1000	0.0203
##	4	1.1524	nan	0.1000	0.0140
##	5	1.1216	nan	0.1000	0.0141
##	6	1.0961	nan	0.1000	0.0125
##	7	1.0730	nan	0.1000	0.0103
##	8	1.0536	nan	0.1000	0.0086
##	9	1.0376	nan	0.1000	0.0067
##	10	1.0208	nan	0.1000	0.0083
##	20	0.9191	nan	0.1000	0.0025
##	40	0.8304	nan	0.1000	0.0006
##	60	0.7967	nan	0.1000	-0.0009

```

##      80      0.7809      nan      0.1000     -0.0007
##     100      0.7724      nan      0.1000     -0.0012
##     120      0.7661      nan      0.1000     -0.0004
##     140      0.7613      nan      0.1000     -0.0012
##     160      0.7558      nan      0.1000     -0.0004
##     180      0.7511      nan      0.1000     -0.0004
##     200      0.7460      nan      0.1000     -0.0006
##     220      0.7421      nan      0.1000     -0.0013
##     240      0.7388      nan      0.1000     -0.0007
##     250      0.7365      nan      0.1000     -0.0009
##
## Iter   TrainDeviance   ValidDeviance   StepSize   Improve
##      1           1.2662           nan      0.1000     0.0324
##      2           1.2096           nan      0.1000     0.0223
##      3           1.1565           nan      0.1000     0.0234
##      4           1.1199           nan      0.1000     0.0184
##      5           1.0844           nan      0.1000     0.0144
##      6           1.0520           nan      0.1000     0.0153
##      7           1.0254           nan      0.1000     0.0130
##      8           1.0023           nan      0.1000     0.0091
##      9           0.9789           nan      0.1000     0.0100
##     10           0.9608           nan      0.1000     0.0074
##     20           0.8450           nan      0.1000     0.0015
##     40           0.7728           nan      0.1000    -0.0010
##     60           0.7478           nan      0.1000    -0.0011
##     80           0.7312           nan      0.1000    -0.0005
##    100           0.7193           nan      0.1000    -0.0008
##    120           0.7053           nan      0.1000    -0.0015
##    140           0.6952           nan      0.1000    -0.0006
##    160           0.6855           nan      0.1000    -0.0007
##    180           0.6758           nan      0.1000    -0.0013
##    200           0.6663           nan      0.1000    -0.0011
##    220           0.6599           nan      0.1000    -0.0008
##    240           0.6511           nan      0.1000    -0.0015
##    250           0.6471           nan      0.1000    -0.0012
##
## Iter   TrainDeviance   ValidDeviance   StepSize   Improve
##      1           1.2605           nan      0.1000     0.0359
##      2           1.1990           nan      0.1000     0.0295
##      3           1.1393           nan      0.1000     0.0250
##      4           1.0933           nan      0.1000     0.0198
##      5           1.0564           nan      0.1000     0.0142
##      6           1.0189           nan      0.1000     0.0153
##      7           0.9922           nan      0.1000     0.0114
##      8           0.9677           nan      0.1000     0.0095
##      9           0.9451           nan      0.1000     0.0105

```

##	10	0.9247	nan	0.1000	0.0070
##	20	0.8103	nan	0.1000	0.0028
##	40	0.7386	nan	0.1000	-0.0008
##	60	0.7043	nan	0.1000	-0.0029
##	80	0.6768	nan	0.1000	-0.0006
##	100	0.6509	nan	0.1000	-0.0012
##	120	0.6340	nan	0.1000	-0.0015
##	140	0.6159	nan	0.1000	-0.0020
##	160	0.6042	nan	0.1000	-0.0015
##	180	0.5892	nan	0.1000	-0.0009
##	200	0.5800	nan	0.1000	-0.0015
##	220	0.5693	nan	0.1000	-0.0015
##	240	0.5583	nan	0.1000	-0.0014
##	250	0.5545	nan	0.1000	-0.0009

##	Iter	TrainDeviance	ValidDeviance	StepSize	Improve
##	1	1.2535	nan	0.1000	0.0373
##	2	1.1854	nan	0.1000	0.0315
##	3	1.1321	nan	0.1000	0.0243
##	4	1.0809	nan	0.1000	0.0232
##	5	1.0425	nan	0.1000	0.0193
##	6	1.0039	nan	0.1000	0.0155
##	7	0.9741	nan	0.1000	0.0124
##	8	0.9470	nan	0.1000	0.0120
##	9	0.9198	nan	0.1000	0.0119
##	10	0.9026	nan	0.1000	0.0068
##	20	0.7913	nan	0.1000	0.0012
##	40	0.7186	nan	0.1000	-0.0003
##	60	0.6837	nan	0.1000	-0.0016
##	80	0.6483	nan	0.1000	-0.0009
##	100	0.6236	nan	0.1000	-0.0014
##	120	0.5982	nan	0.1000	-0.0024
##	140	0.5790	nan	0.1000	-0.0015
##	160	0.5596	nan	0.1000	-0.0035
##	180	0.5411	nan	0.1000	-0.0007
##	200	0.5267	nan	0.1000	-0.0020
##	220	0.5139	nan	0.1000	-0.0012
##	240	0.4973	nan	0.1000	-0.0021
##	250	0.4919	nan	0.1000	-0.0009

##	Iter	TrainDeviance	ValidDeviance	StepSize	Improve
##	1	1.2535	nan	0.1000	0.0388
##	2	1.1840	nan	0.1000	0.0318
##	3	1.1232	nan	0.1000	0.0262
##	4	1.0698	nan	0.1000	0.0244
##	5	1.0285	nan	0.1000	0.0185


```

##      6      0.9955      nan      0.1000      0.0134
##      7      0.9661      nan      0.1000      0.0146
##      8      0.9362      nan      0.1000      0.0139
##      9      0.9138      nan      0.1000      0.0090
##     10      0.8958      nan      0.1000      0.0069
##     20      0.7770      nan      0.1000      0.0004
##     40      0.7015      nan      0.1000     -0.0011
##     60      0.6610      nan      0.1000     -0.0024
##     80      0.6300      nan      0.1000     -0.0035
##    100      0.5995      nan      0.1000     -0.0016
##    120      0.5673      nan      0.1000     -0.0024
##    140      0.5459      nan      0.1000     -0.0038
##    160      0.5245      nan      0.1000     -0.0029
##    180      0.5053      nan      0.1000     -0.0016
##    200      0.4902      nan      0.1000     -0.0012
##    220      0.4748      nan      0.1000     -0.0026
##    240      0.4577      nan      0.1000     -0.0022
##    250      0.4513      nan      0.1000     -0.0027
##
## Iter   TrainDeviance   ValidDeviance   StepSize   Improve
##      1          1.2725           nan      0.1000    0.0315
##      2          1.2219           nan      0.1000    0.0246
##      3          1.1809           nan      0.1000    0.0195
##      4          1.1501           nan      0.1000    0.0136
##      5          1.1172           nan      0.1000    0.0151
##      6          1.0930           nan      0.1000    0.0122
##      7          1.0669           nan      0.1000    0.0116
##      8          1.0449           nan      0.1000    0.0093
##      9          1.0285           nan      0.1000    0.0066
##     10          1.0180           nan      0.1000    0.0038
##     20          0.9148           nan      0.1000    0.0024
##     40          0.8255           nan      0.1000   -0.0002
##     60          0.7874           nan      0.1000    0.0005
##     80          0.7699           nan      0.1000   -0.0002
##    100          0.7627           nan      0.1000   -0.0013
##    120          0.7552           nan      0.1000   -0.0003
##    140          0.7490           nan      0.1000   -0.0009
##    160          0.7430           nan      0.1000   -0.0010
##    180          0.7382           nan      0.1000   -0.0008
##    200          0.7349           nan      0.1000   -0.0005
##    220          0.7314           nan      0.1000   -0.0013
##    240          0.7268           nan      0.1000   -0.0007
##    250          0.7253           nan      0.1000   -0.0006
##
## Iter   TrainDeviance   ValidDeviance   StepSize   Improve
##      1          1.2623           nan      0.1000    0.0345

```

##	2	1.2026	nan	0.1000	0.0294
##	3	1.1548	nan	0.1000	0.0218
##	4	1.1129	nan	0.1000	0.0177
##	5	1.0812	nan	0.1000	0.0157
##	6	1.0509	nan	0.1000	0.0147
##	7	1.0280	nan	0.1000	0.0099
##	8	1.0037	nan	0.1000	0.0104
##	9	0.9810	nan	0.1000	0.0097
##	10	0.9609	nan	0.1000	0.0089
##	20	0.8459	nan	0.1000	0.0019
##	40	0.7693	nan	0.1000	-0.0002
##	60	0.7363	nan	0.1000	-0.0008
##	80	0.7162	nan	0.1000	-0.0023
##	100	0.7052	nan	0.1000	-0.0015
##	120	0.6917	nan	0.1000	-0.0028
##	140	0.6756	nan	0.1000	-0.0020
##	160	0.6652	nan	0.1000	-0.0003
##	180	0.6557	nan	0.1000	-0.0014
##	200	0.6468	nan	0.1000	-0.0006
##	220	0.6353	nan	0.1000	-0.0006
##	240	0.6277	nan	0.1000	-0.0007
##	250	0.6244	nan	0.1000	-0.0011
##					
##	Iter	TrainDeviance	ValidDeviance	StepSize	Improve
##	1	1.2566	nan	0.1000	0.0386
##	2	1.1937	nan	0.1000	0.0297
##	3	1.1355	nan	0.1000	0.0258
##	4	1.0909	nan	0.1000	0.0201
##	5	1.0509	nan	0.1000	0.0179
##	6	1.0145	nan	0.1000	0.0162
##	7	0.9840	nan	0.1000	0.0123
##	8	0.9551	nan	0.1000	0.0122
##	9	0.9369	nan	0.1000	0.0088
##	10	0.9164	nan	0.1000	0.0072
##	20	0.8017	nan	0.1000	0.0007
##	40	0.7255	nan	0.1000	-0.0002
##	60	0.6940	nan	0.1000	-0.0009
##	80	0.6684	nan	0.1000	-0.0008
##	100	0.6504	nan	0.1000	-0.0011
##	120	0.6325	nan	0.1000	-0.0013
##	140	0.6140	nan	0.1000	-0.0013
##	160	0.5974	nan	0.1000	-0.0002
##	180	0.5833	nan	0.1000	-0.0010
##	200	0.5666	nan	0.1000	-0.0011
##	220	0.5550	nan	0.1000	-0.0020
##	240	0.5451	nan	0.1000	-0.0023

```

##      250      0.5372      nan      0.1000     -0.0015
##
## Iter   TrainDeviance   ValidDeviance   StepSize   Improve
##      1      1.2498      nan      0.1000     0.0417
##      2      1.1832      nan      0.1000     0.0281
##      3      1.1271      nan      0.1000     0.0244
##      4      1.0777      nan      0.1000     0.0218
##      5      1.0407      nan      0.1000     0.0151
##      6      1.0058      nan      0.1000     0.0141
##      7      0.9753      nan      0.1000     0.0148
##      8      0.9445      nan      0.1000     0.0133
##      9      0.9192      nan      0.1000     0.0115
##     10      0.9036      nan      0.1000     0.0060
##     20      0.7800      nan      0.1000     0.0017
##     40      0.7027      nan      0.1000    -0.0018
##     60      0.6571      nan      0.1000    -0.0015
##     80      0.6285      nan      0.1000    -0.0028
##    100      0.6088      nan      0.1000    -0.0028
##    120      0.5816      nan      0.1000    -0.0033
##    140      0.5649      nan      0.1000    -0.0020
##    160      0.5430      nan      0.1000    -0.0012
##    180      0.5294      nan      0.1000    -0.0016
##    200      0.5109      nan      0.1000    -0.0021
##    220      0.4943      nan      0.1000    -0.0027
##    240      0.4793      nan      0.1000    -0.0010
##    250      0.4732      nan      0.1000    -0.0016
##
## Iter   TrainDeviance   ValidDeviance   StepSize   Improve
##      1      1.2514      nan      0.1000     0.0374
##      2      1.1825      nan      0.1000     0.0323
##      3      1.1249      nan      0.1000     0.0244
##      4      1.0758      nan      0.1000     0.0248
##      5      1.0288      nan      0.1000     0.0205
##      6      0.9950      nan      0.1000     0.0151
##      7      0.9607      nan      0.1000     0.0120
##      8      0.9307      nan      0.1000     0.0132
##      9      0.9071      nan      0.1000     0.0100
##     10      0.8828      nan      0.1000     0.0094
##     20      0.7598      nan      0.1000     0.0013
##     40      0.6744      nan      0.1000    -0.0017
##     60      0.6343      nan      0.1000    -0.0018
##     80      0.5947      nan      0.1000    -0.0013
##    100      0.5667      nan      0.1000    -0.0019
##    120      0.5457      nan      0.1000    -0.0029
##    140      0.5235      nan      0.1000    -0.0024
##    160      0.5036      nan      0.1000    -0.0019

```

##	180	0.4856	nan	0.1000	-0.0030
##	200	0.4678	nan	0.1000	-0.0020
##	220	0.4534	nan	0.1000	-0.0023
##	240	0.4412	nan	0.1000	-0.0018
##	250	0.4336	nan	0.1000	-0.0019

##

##	Iter	TrainDeviance	ValidDeviance	StepSize	Improve
##	1	1.2803	nan	0.1000	0.0300
##	2	1.2282	nan	0.1000	0.0252
##	3	1.1859	nan	0.1000	0.0212
##	4	1.1515	nan	0.1000	0.0150
##	5	1.1195	nan	0.1000	0.0146
##	6	1.0928	nan	0.1000	0.0127
##	7	1.0679	nan	0.1000	0.0107
##	8	1.0480	nan	0.1000	0.0092
##	9	1.0320	nan	0.1000	0.0078
##	10	1.0185	nan	0.1000	0.0060
##	20	0.9144	nan	0.1000	0.0035
##	40	0.8227	nan	0.1000	0.0009
##	60	0.7776	nan	0.1000	-0.0003
##	80	0.7595	nan	0.1000	-0.0005
##	100	0.7498	nan	0.1000	-0.0002
##	120	0.7441	nan	0.1000	-0.0002
##	140	0.7392	nan	0.1000	-0.0006
##	160	0.7347	nan	0.1000	-0.0004
##	180	0.7323	nan	0.1000	-0.0013
##	200	0.7259	nan	0.1000	-0.0003
##	220	0.7226	nan	0.1000	-0.0003
##	240	0.7204	nan	0.1000	-0.0004
##	250	0.7180	nan	0.1000	-0.0010

##

##	Iter	TrainDeviance	ValidDeviance	StepSize	Improve
##	1	1.2639	nan	0.1000	0.0365
##	2	1.2041	nan	0.1000	0.0277
##	3	1.1555	nan	0.1000	0.0224
##	4	1.1135	nan	0.1000	0.0190
##	5	1.0741	nan	0.1000	0.0174
##	6	1.0444	nan	0.1000	0.0128
##	7	1.0172	nan	0.1000	0.0120
##	8	0.9928	nan	0.1000	0.0114
##	9	0.9698	nan	0.1000	0.0105
##	10	0.9523	nan	0.1000	0.0071
##	20	0.8290	nan	0.1000	0.0023
##	40	0.7563	nan	0.1000	-0.0012
##	60	0.7274	nan	0.1000	-0.0011
##	80	0.7078	nan	0.1000	-0.0014

```

##      100      0.6940      nan      0.1000     -0.0008
##      120      0.6795      nan      0.1000     -0.0015
##      140      0.6698      nan      0.1000     -0.0008
##      160      0.6585      nan      0.1000     -0.0011
##      180      0.6453      nan      0.1000     -0.0002
##      200      0.6358      nan      0.1000     -0.0007
##      220      0.6313      nan      0.1000     -0.0019
##      240      0.6263      nan      0.1000     -0.0020
##      250      0.6222      nan      0.1000     -0.0013
##
## Iter   TrainDeviance   ValidDeviance   StepSize   Improve
##      1          1.2566           nan      0.1000     0.0352
##      2          1.1939           nan      0.1000     0.0309
##      3          1.1369           nan      0.1000     0.0252
##      4          1.0871           nan      0.1000     0.0217
##      5          1.0427           nan      0.1000     0.0188
##      6          1.0083           nan      0.1000     0.0132
##      7          0.9772           nan      0.1000     0.0138
##      8          0.9496           nan      0.1000     0.0146
##      9          0.9264           nan      0.1000     0.0103
##     10          0.9027           nan      0.1000     0.0097
##     20          0.7804           nan      0.1000     0.0026
##     40          0.7150           nan      0.1000    -0.0014
##     60          0.6814           nan      0.1000    -0.0023
##     80          0.6575           nan      0.1000    -0.0021
##    100          0.6395           nan      0.1000    -0.0021
##    120          0.6198           nan      0.1000    -0.0014
##    140          0.6030           nan      0.1000    -0.0016
##    160          0.5852           nan      0.1000    -0.0013
##    180          0.5747           nan      0.1000    -0.0017
##    200          0.5616           nan      0.1000    -0.0014
##    220          0.5461           nan      0.1000    -0.0019
##    240          0.5350           nan      0.1000    -0.0025
##    250          0.5289           nan      0.1000    -0.0008
##
## Iter   TrainDeviance   ValidDeviance   StepSize   Improve
##      1          1.2508           nan      0.1000     0.0399
##      2          1.1799           nan      0.1000     0.0327
##      3          1.1202           nan      0.1000     0.0267
##      4          1.0681           nan      0.1000     0.0230
##      5          1.0267           nan      0.1000     0.0188
##      6          0.9903           nan      0.1000     0.0144
##      7          0.9579           nan      0.1000     0.0146
##      8          0.9311           nan      0.1000     0.0122
##      9          0.9093           nan      0.1000     0.0100
##     10          0.8920           nan      0.1000     0.0071

```

##	20	0.7669	nan	0.1000	0.0027
##	40	0.6912	nan	0.1000	-0.0021
##	60	0.6458	nan	0.1000	-0.0026
##	80	0.6124	nan	0.1000	-0.0004
##	100	0.5870	nan	0.1000	-0.0014
##	120	0.5656	nan	0.1000	-0.0013
##	140	0.5518	nan	0.1000	-0.0023
##	160	0.5341	nan	0.1000	-0.0011
##	180	0.5194	nan	0.1000	-0.0010
##	200	0.5076	nan	0.1000	-0.0026
##	220	0.4978	nan	0.1000	-0.0018
##	240	0.4803	nan	0.1000	-0.0012
##	250	0.4720	nan	0.1000	-0.0022

##	Iter	TrainDeviance	ValidDeviance	StepSize	Improve
##	1	1.2479	nan	0.1000	0.0415
##	2	1.1762	nan	0.1000	0.0296
##	3	1.1123	nan	0.1000	0.0268
##	4	1.0657	nan	0.1000	0.0197
##	5	1.0210	nan	0.1000	0.0199
##	6	0.9815	nan	0.1000	0.0159
##	7	0.9479	nan	0.1000	0.0105
##	8	0.9197	nan	0.1000	0.0126
##	9	0.8955	nan	0.1000	0.0088
##	10	0.8742	nan	0.1000	0.0090
##	20	0.7586	nan	0.1000	0.0012
##	40	0.6843	nan	0.1000	-0.0018
##	60	0.6355	nan	0.1000	-0.0027
##	80	0.5939	nan	0.1000	-0.0014
##	100	0.5648	nan	0.1000	-0.0014
##	120	0.5434	nan	0.1000	-0.0016
##	140	0.5223	nan	0.1000	-0.0021
##	160	0.5026	nan	0.1000	-0.0021
##	180	0.4889	nan	0.1000	-0.0021
##	200	0.4758	nan	0.1000	-0.0012
##	220	0.4520	nan	0.1000	-0.0008
##	240	0.4417	nan	0.1000	-0.0027
##	250	0.4358	nan	0.1000	-0.0023

##	Iter	TrainDeviance	ValidDeviance	StepSize	Improve
##	1	1.2643	nan	0.1000	0.0332
##	2	1.2100	nan	0.1000	0.0272
##	3	1.1593	nan	0.1000	0.0240
##	4	1.1180	nan	0.1000	0.0196
##	5	1.0825	nan	0.1000	0.0170
##	6	1.0508	nan	0.1000	0.0135


```
##      7      1.0219      nan    0.1000    0.0112
##      8      0.9984      nan    0.1000    0.0118
##      9      0.9778      nan    0.1000    0.0072
##     10      0.9565      nan    0.1000    0.0096
##     20      0.8389      nan    0.1000    0.0025
##     40      0.7624      nan    0.1000   -0.0009
##     50      0.7453      nan    0.1000   -0.0006
```

```
oj.gbm
```


```
## Stochastic Gradient Boosting
##
## 857 samples
## 17 predictor
## 2 classes: 'CH', 'MM'
##
## No pre-processing
## Resampling: Cross-Validated (10 fold)
## Summary of sample sizes: 772, 770, 771, 771, 772, 771, ...
## Resampling results across tuning parameters:
##
##  interaction.depth  n.trees  ROC      Sens      Spec
##  1                   50      0.8869715  0.8720247  0.7245098
##  1                   100     0.8889466  0.8738389  0.7302139
##  1                   150     0.8850612  0.8680697  0.7303030
##  1                   200     0.8838868  0.8738026  0.7213012
##  1                   250     0.8805090  0.8622642  0.7275401
##  2                    50     0.8895139  0.8719521  0.7451872
##  2                   100     0.8832208  0.8623367  0.7334225
##  2                   150     0.8838687  0.8679971  0.7394831
##  2                   200     0.8811028  0.8642598  0.7245989
##  2                   250     0.8780432  0.8489840  0.7336898
##  3                    50     0.8833620  0.8700290  0.7245989
##  3                   100     0.8785063  0.8604862  0.7454545
##  3                   150     0.8744782  0.8470972  0.7156863
##  3                   200     0.8723217  0.8414731  0.7246881
##  3                   250     0.8722653  0.8337446  0.7217469
##  4                    50     0.8854180  0.8604862  0.7605169
##  4                   100     0.8794721  0.8413280  0.7336898
##  4                   150     0.8741609  0.8432511  0.7308378
##  4                   200     0.8710638  0.8470972  0.7249554
##  4                   250     0.8674771  0.8490203  0.7218360
##  5                    50     0.8868076  0.8585994  0.7574866
##  5                   100     0.8786285  0.8471698  0.7486631
##  5                   150     0.8708328  0.8395138  0.7458111
```

```
##      5              200      0.8648622 0.8452467 0.7455437
##      5              250      0.8594578 0.8375544 0.7336007
##
## Tuning parameter 'shrinkage' was held constant at a value of 0.1
##
## Tuning parameter 'n.minobsinnode' was held constant at a value of 10
## ROC was used to select the optimal model using the largest value.
## The final values used for the model were n.trees = 50, interaction.depth
## = 2, shrinkage = 0.1 and n.minobsinnode = 10.
```

```
plot(oj.gbm)
```



```
oj.pred <- predict(oj.gbm, oj_test, type = "raw")
plot(oj_test$Purchase, oj.pred,
     main = "Gradient Boosting Classification: Predicted vs. Actual",
     xlab = "Actual",
     ylab = "Predicted")
```

```
(oj.conf <- confusionMatrix(data = oj.pred,
                             reference = oj_test$Purchase))
```

```
## Confusion Matrix and Statistics
##
##           Reference
## Prediction  CH  MM
##      CH 117  18
##      MM  13  65
##
##           Accuracy : 0.8545
##           95% CI : (0.7998, 0.8989)
##      No Information Rate : 0.6103
##      P-Value [Acc > NIR] : 4.83e-15
##
##           Kappa : 0.6907
##
##  McNemar's Test P-Value : 0.4725
##
##           Sensitivity : 0.9000
##           Specificity : 0.7831
##      Pos Pred Value : 0.8667
##      Neg Pred Value : 0.8333
##           Prevalence : 0.6103
```

```
##           Detection Rate : 0.5493
##      Detection Prevalence : 0.6338
##           Balanced Accuracy : 0.8416
##
##           'Positive' Class : CH
##

oj.gbm.acc <- as.numeric(oj.conf$overall[1])
rm(oj.pred)
rm(oj.conf)
#plot(oj.bag$, oj.bag$finalModel$y)
#plot(varImp(oj.gbm), main="Variable Importance with Gradient Boosting")
```

8.5.0.2 Gradient Boosting Regression Example

Again using the `Carseats` data set to predict `Sales`, this time I'll use the gradient boosting method by specifying `method = "gbm"`. I'll use `tuneLength = 5` and not worry about `tuneGrid` anymore. `Caret` tunes the following hyperparameters.

- `n.trees`: number of boosting iterations (increasing `n.trees` reduces the error on training set, but may lead to over-fitting)
- `interaction.depth`: maximum tree depth (the default six - node tree appears to do an excellent job)
- `shrinkage`: learning rate (reduces the impact of each additional fitted base-learner (tree) by reducing the size of incremental steps and thus penalizes the importance of each consecutive iteration. The intuition is that it is better to improve a model by taking many small steps than by taking fewer large steps. If one of the boosting iterations turns out to be erroneous, its negative impact can be easily corrected in subsequent steps.)
- `n.minobsinnode`: minimum terminal node size

```
carseats.gbm <- train(Sales ~ .,
  data = carseats_train,
  method = "gbm", # for bagged tree
  tuneLength = 5, # choose up to 5 combinations of tuning parameters
  metric = "RMSE", # evaluate hyperparameter combinations with ROC
  trControl = trainControl(
    method = "cv", # k-fold cross validation
    number = 10, # 10 folds
    savePredictions = "final", # save predictions for the op
    verboseIter = FALSE,
    returnData = FALSE
  )
)
```

##	Iter	TrainDeviance	ValidDeviance	StepSize	Improve
##	1	7.6155	nan	0.1000	0.3126
##	2	7.3677	nan	0.1000	0.2471
##	3	7.1383	nan	0.1000	0.1559
##	4	6.8796	nan	0.1000	0.3000
##	5	6.6084	nan	0.1000	0.2696
##	6	6.3846	nan	0.1000	0.1575
##	7	6.1551	nan	0.1000	0.2016
##	8	5.9837	nan	0.1000	0.1171
##	9	5.7969	nan	0.1000	0.1558
##	10	5.6503	nan	0.1000	0.1243
##	20	4.5758	nan	0.1000	0.0472
##	40	3.3276	nan	0.1000	0.0043
##	60	2.6161	nan	0.1000	0.0154
##	80	2.1215	nan	0.1000	-0.0029
##	100	1.7822	nan	0.1000	-0.0166
##	120	1.5354	nan	0.1000	-0.0016
##	140	1.3313	nan	0.1000	0.0067
##	160	1.2074	nan	0.1000	-0.0030
##	180	1.0966	nan	0.1000	0.0005
##	200	1.0083	nan	0.1000	-0.0019
##	220	0.9572	nan	0.1000	-0.0008
##	240	0.9048	nan	0.1000	-0.0052
##	250	0.8922	nan	0.1000	-0.0051
##					
##	Iter	TrainDeviance	ValidDeviance	StepSize	Improve
##	1	7.4939	nan	0.1000	0.5303
##	2	6.9609	nan	0.1000	0.4503
##	3	6.5646	nan	0.1000	0.3312
##	4	6.2135	nan	0.1000	0.2836
##	5	5.9347	nan	0.1000	0.2472
##	6	5.6654	nan	0.1000	0.2045
##	7	5.3757	nan	0.1000	0.2134
##	8	5.1883	nan	0.1000	0.1810
##	9	5.0431	nan	0.1000	0.1205
##	10	4.8440	nan	0.1000	0.0749
##	20	3.4421	nan	0.1000	0.1291
##	40	2.0285	nan	0.1000	0.0246
##	60	1.4136	nan	0.1000	0.0003
##	80	1.0839	nan	0.1000	0.0003
##	100	0.9011	nan	0.1000	0.0052
##	120	0.8195	nan	0.1000	-0.0075
##	140	0.7664	nan	0.1000	-0.0036
##	160	0.7243	nan	0.1000	-0.0010
##	180	0.6839	nan	0.1000	-0.0028
##	200	0.6448	nan	0.1000	-0.0048

##	220	0.6123	nan	0.1000	-0.0035
##	240	0.5897	nan	0.1000	-0.0038
##	250	0.5767	nan	0.1000	-0.0050

##

##	Iter	TrainDeviance	ValidDeviance	StepSize	Improve
##	1	7.3122	nan	0.1000	0.5956
##	2	6.7737	nan	0.1000	0.4295
##	3	6.3033	nan	0.1000	0.4215
##	4	5.9081	nan	0.1000	0.2966
##	5	5.5342	nan	0.1000	0.3320
##	6	5.1666	nan	0.1000	0.2748
##	7	4.9365	nan	0.1000	0.1688
##	8	4.6557	nan	0.1000	0.2144
##	9	4.4412	nan	0.1000	0.1072
##	10	4.2075	nan	0.1000	0.1861
##	20	2.7722	nan	0.1000	0.0107
##	40	1.4659	nan	0.1000	0.0203
##	60	1.0163	nan	0.1000	-0.0055
##	80	0.8430	nan	0.1000	-0.0074
##	100	0.7427	nan	0.1000	-0.0031
##	120	0.6731	nan	0.1000	-0.0078
##	140	0.6151	nan	0.1000	-0.0080
##	160	0.5814	nan	0.1000	-0.0074
##	180	0.5452	nan	0.1000	-0.0054
##	200	0.5023	nan	0.1000	-0.0071
##	220	0.4697	nan	0.1000	-0.0058
##	240	0.4340	nan	0.1000	-0.0022
##	250	0.4207	nan	0.1000	-0.0088

##

##	Iter	TrainDeviance	ValidDeviance	StepSize	Improve
##	1	7.3508	nan	0.1000	0.7012
##	2	6.7011	nan	0.1000	0.5533
##	3	6.1585	nan	0.1000	0.5138
##	4	5.7274	nan	0.1000	0.3542
##	5	5.2438	nan	0.1000	0.3597
##	6	4.9357	nan	0.1000	0.2545
##	7	4.6359	nan	0.1000	0.2195
##	8	4.3960	nan	0.1000	0.2294
##	9	4.1786	nan	0.1000	0.1539
##	10	3.9850	nan	0.1000	0.1414
##	20	2.4927	nan	0.1000	0.0532
##	40	1.2882	nan	0.1000	0.0112
##	60	0.8551	nan	0.1000	-0.0074
##	80	0.6752	nan	0.1000	-0.0044
##	100	0.5795	nan	0.1000	-0.0097
##	120	0.4983	nan	0.1000	-0.0038

```

##      140      0.4440      nan      0.1000     -0.0032
##      160      0.4044      nan      0.1000     -0.0056
##      180      0.3666      nan      0.1000     -0.0076
##      200      0.3337      nan      0.1000     -0.0065
##      220      0.3066      nan      0.1000     -0.0022
##      240      0.2736      nan      0.1000     -0.0039
##      250      0.2608      nan      0.1000     -0.0038
##
## Iter   TrainDeviance   ValidDeviance   StepSize   Improve
##      1           7.2719           nan      0.1000     0.7029
##      2           6.5623           nan      0.1000     0.5180
##      3           6.0426           nan      0.1000     0.4144
##      4           5.5613           nan      0.1000     0.4419
##      5           5.1070           nan      0.1000     0.2878
##      6           4.7641           nan      0.1000     0.2383
##      7           4.4677           nan      0.1000     0.2141
##      8           4.2399           nan      0.1000     0.0706
##      9           3.9751           nan      0.1000     0.1927
##     10           3.7100           nan      0.1000     0.1701
##     20           2.2565           nan      0.1000     0.0859
##     40           1.1000           nan      0.1000    -0.0080
##     60           0.7464           nan      0.1000    -0.0020
##     80           0.5790           nan      0.1000    -0.0024
##    100           0.4941           nan      0.1000    -0.0054
##    120           0.4290           nan      0.1000    -0.0030
##    140           0.3723           nan      0.1000    -0.0052
##    160           0.3330           nan      0.1000    -0.0042
##    180           0.2951           nan      0.1000    -0.0056
##    200           0.2627           nan      0.1000    -0.0038
##    220           0.2336           nan      0.1000    -0.0018
##    240           0.2089           nan      0.1000    -0.0017
##    250           0.1964           nan      0.1000    -0.0029
##
## Iter   TrainDeviance   ValidDeviance   StepSize   Improve
##      1           7.8758           nan      0.1000     0.1611
##      2           7.4638           nan      0.1000     0.4117
##      3           7.1509           nan      0.1000     0.2868
##      4           6.8550           nan      0.1000     0.2921
##      5           6.6144           nan      0.1000     0.2482
##      6           6.4519           nan      0.1000     0.1654
##      7           6.3108           nan      0.1000     0.0792
##      8           6.1631           nan      0.1000     0.1260
##      9           6.0160           nan      0.1000     0.1119
##     10           5.8043           nan      0.1000     0.1567
##     20           4.6275           nan      0.1000     0.0715
##     40           3.4435           nan      0.1000    -0.0080

```

##	60	2.6905	nan	0.1000	0.0048
##	80	2.1544	nan	0.1000	0.0127
##	100	1.7772	nan	0.1000	-0.0062
##	120	1.4927	nan	0.1000	-0.0058
##	140	1.3013	nan	0.1000	0.0010
##	160	1.1609	nan	0.1000	0.0023
##	180	1.0670	nan	0.1000	-0.0058
##	200	0.9890	nan	0.1000	-0.0088
##	220	0.9407	nan	0.1000	0.0000
##	240	0.9016	nan	0.1000	-0.0042
##	250	0.8853	nan	0.1000	-0.0037

##	Iter	TrainDeviance	ValidDeviance	StepSize	Improve
##	1	7.6412	nan	0.1000	0.5298
##	2	7.0819	nan	0.1000	0.4917
##	3	6.5521	nan	0.1000	0.3297
##	4	6.1791	nan	0.1000	0.2744
##	5	5.9412	nan	0.1000	0.1842
##	6	5.6434	nan	0.1000	0.2504
##	7	5.3703	nan	0.1000	0.2167
##	8	5.1224	nan	0.1000	0.1739
##	9	4.9715	nan	0.1000	0.1184
##	10	4.7654	nan	0.1000	0.1615
##	20	3.3795	nan	0.1000	0.0636
##	40	2.0395	nan	0.1000	0.0080
##	60	1.4605	nan	0.1000	-0.0029
##	80	1.1344	nan	0.1000	-0.0025
##	100	0.9495	nan	0.1000	-0.0087
##	120	0.8562	nan	0.1000	-0.0037
##	140	0.7855	nan	0.1000	-0.0043
##	160	0.7298	nan	0.1000	-0.0057
##	180	0.6813	nan	0.1000	-0.0010
##	200	0.6441	nan	0.1000	-0.0027
##	220	0.6118	nan	0.1000	-0.0040
##	240	0.5838	nan	0.1000	-0.0071
##	250	0.5734	nan	0.1000	-0.0007

##	Iter	TrainDeviance	ValidDeviance	StepSize	Improve
##	1	7.4922	nan	0.1000	0.6088
##	2	6.9029	nan	0.1000	0.5644
##	3	6.4669	nan	0.1000	0.4526
##	4	5.9980	nan	0.1000	0.4160
##	5	5.5505	nan	0.1000	0.2603
##	6	5.2643	nan	0.1000	0.2640
##	7	5.0169	nan	0.1000	0.1944
##	8	4.8024	nan	0.1000	0.1732

```

##      9      4.5720      nan      0.1000      0.1124
##     10      4.3508      nan      0.1000      0.1700
##     20      2.8015      nan      0.1000      0.0839
##     40      1.5058      nan      0.1000      0.0041
##     60      1.0433      nan      0.1000     -0.0028
##     80      0.8409      nan      0.1000     -0.0038
##    100      0.7262      nan      0.1000     -0.0091
##    120      0.6504      nan      0.1000     -0.0011
##    140      0.5944      nan      0.1000     -0.0082
##    160      0.5390      nan      0.1000     -0.0053
##    180      0.5004      nan      0.1000     -0.0058
##    200      0.4655      nan      0.1000     -0.0034
##    220      0.4306      nan      0.1000     -0.0051
##    240      0.4014      nan      0.1000     -0.0052
##    250      0.3922      nan      0.1000     -0.0045
##
## Iter  TrainDeviance  ValidDeviance  StepSize  Improve
##     1           7.5384           nan      0.1000    0.6216
##     2           6.8492           nan      0.1000    0.5668
##     3           6.2425           nan      0.1000    0.4236
##     4           5.7761           nan      0.1000    0.4459
##     5           5.3374           nan      0.1000    0.3815
##     6           5.0229           nan      0.1000    0.2852
##     7           4.7392           nan      0.1000    0.2818
##     8           4.4323           nan      0.1000    0.1682
##     9           4.2010           nan      0.1000    0.2173
##    10           3.9190           nan      0.1000    0.1539
##    20           2.3973           nan      0.1000    0.0126
##    40           1.2163           nan      0.1000    0.0078
##    60           0.8634           nan      0.1000   -0.0088
##    80           0.7010           nan      0.1000   -0.0174
##   100           0.6125           nan      0.1000   -0.0047
##   120           0.5409           nan      0.1000   -0.0043
##   140           0.4880           nan      0.1000   -0.0087
##   160           0.4416           nan      0.1000   -0.0077
##   180           0.4022           nan      0.1000   -0.0087
##   200           0.3700           nan      0.1000   -0.0066
##   220           0.3326           nan      0.1000   -0.0048
##   240           0.3010           nan      0.1000   -0.0020
##   250           0.2897           nan      0.1000   -0.0057
##
## Iter  TrainDeviance  ValidDeviance  StepSize  Improve
##     1           7.4641           nan      0.1000    0.6228
##     2           6.7842           nan      0.1000    0.6612
##     3           6.2574           nan      0.1000    0.5100
##     4           5.6693           nan      0.1000    0.5417

```

```

##      5      5.2506      nan      0.1000      0.3664
##      6      4.8195      nan      0.1000      0.3499
##      7      4.4803      nan      0.1000      0.2959
##      8      4.1807      nan      0.1000      0.1964
##      9      3.9058      nan      0.1000      0.1460
##     10      3.6831      nan      0.1000      0.1246
##     20      2.1373      nan      0.1000      0.0659
##     40      1.0923      nan      0.1000      0.0073
##     60      0.7575      nan      0.1000     -0.0129
##     80      0.6127      nan      0.1000     -0.0132
##    100      0.5130      nan      0.1000     -0.0123
##    120      0.4313      nan      0.1000     -0.0060
##    140      0.3732      nan      0.1000     -0.0102
##    160      0.3229      nan      0.1000     -0.0040
##    180      0.2862      nan      0.1000     -0.0015
##    200      0.2556      nan      0.1000     -0.0035
##    220      0.2289      nan      0.1000     -0.0049
##    240      0.2066      nan      0.1000     -0.0024
##    250      0.1957      nan      0.1000     -0.0030
##
## Iter   TrainDeviance   ValidDeviance   StepSize   Improve
##      1           8.0133           nan      0.1000    0.2136
##      2           7.6170           nan      0.1000    0.3813
##      3           7.2757           nan      0.1000    0.3566
##      4           6.9741           nan      0.1000    0.2780
##      5           6.7490           nan      0.1000    0.2272
##      6           6.6030           nan      0.1000    0.0955
##      7           6.3914           nan      0.1000    0.1949
##      8           6.2217           nan      0.1000    0.1345
##      9           6.0156           nan      0.1000    0.1629
##     10           5.8483           nan      0.1000    0.1217
##     20           4.7672           nan      0.1000    0.0084
##     40           3.5325           nan      0.1000    0.0072
##     60           2.7373           nan      0.1000    0.0204
##     80           2.2393           nan      0.1000    0.0177
##    100           1.8533           nan      0.1000   -0.0056
##    120           1.5588           nan      0.1000   -0.0006
##    140           1.3684           nan      0.1000   -0.0059
##    160           1.2137           nan      0.1000   -0.0029
##    180           1.0929           nan      0.1000    0.0053
##    200           1.0225           nan      0.1000   -0.0018
##    220           0.9612           nan      0.1000   -0.0014
##    240           0.9268           nan      0.1000   -0.0089
##    250           0.9073           nan      0.1000   -0.0005
##
## Iter   TrainDeviance   ValidDeviance   StepSize   Improve

```



```

##      1      7.6769      nan    0.1000    0.4910
##      2      7.1336      nan    0.1000    0.4624
##      3      6.7130      nan    0.1000    0.3977
##      4      6.3914      nan    0.1000    0.2512
##      5      6.0813      nan    0.1000    0.3267
##      6      5.8205      nan    0.1000    0.1912
##      7      5.5575      nan    0.1000    0.1760
##      8      5.3300      nan    0.1000    0.1566
##      9      5.1406      nan    0.1000    0.0895
##     10      4.9999      nan    0.1000    0.0993
##     20      3.4990      nan    0.1000    0.1416
##     40      2.1301      nan    0.1000    0.0205
##     60      1.4553      nan    0.1000   -0.0010
##     80      1.1386      nan    0.1000   -0.0129
##    100      0.9532      nan    0.1000   -0.0062
##    120      0.8461      nan    0.1000   -0.0020
##    140      0.7941      nan    0.1000   -0.0103
##    160      0.7479      nan    0.1000   -0.0069
##    180      0.7076      nan    0.1000   -0.0046
##    200      0.6764      nan    0.1000   -0.0034
##    220      0.6375      nan    0.1000   -0.0039
##    240      0.6118      nan    0.1000   -0.0084
##    250      0.5967      nan    0.1000   -0.0038
##
## Iter  TrainDeviance  ValidDeviance  StepSize  Improve
##      1      7.5573      nan    0.1000    0.6889
##      2      6.9522      nan    0.1000    0.5641
##      3      6.4789      nan    0.1000    0.3647
##      4      6.1178      nan    0.1000    0.1990
##      5      5.7289      nan    0.1000    0.3137
##      6      5.4295      nan    0.1000    0.1935
##      7      5.0791      nan    0.1000    0.2922
##      8      4.8212      nan    0.1000    0.2147
##      9      4.5435      nan    0.1000    0.1970
##     10      4.3324      nan    0.1000    0.1800
##     20      2.8598      nan    0.1000    0.0779
##     40      1.5947      nan    0.1000    0.0099
##     60      1.0802      nan    0.1000   -0.0061
##     80      0.8896      nan    0.1000   -0.0082
##    100      0.7621      nan    0.1000   -0.0069
##    120      0.6690      nan    0.1000   -0.0080
##    140      0.6158      nan    0.1000   -0.0050
##    160      0.5707      nan    0.1000   -0.0019
##    180      0.5336      nan    0.1000   -0.0046
##    200      0.4984      nan    0.1000   -0.0062
##    220      0.4607      nan    0.1000   -0.0045

```

```

##      240      0.4323      nan      0.1000     -0.0052
##      250      0.4196      nan      0.1000     -0.0041
##
## Iter   TrainDeviance   ValidDeviance   StepSize   Improve
##      1           7.4906           nan      0.1000     0.6789
##      2           6.9506           nan      0.1000     0.3624
##      3           6.3953           nan      0.1000     0.4020
##      4           5.9823           nan      0.1000     0.3891
##      5           5.5966           nan      0.1000     0.2061
##      6           5.1879           nan      0.1000     0.2865
##      7           4.8319           nan      0.1000     0.2808
##      8           4.5432           nan      0.1000     0.2153
##      9           4.2635           nan      0.1000     0.2355
##     10           3.9960           nan      0.1000     0.1091
##     20           2.4108           nan      0.1000     0.1056
##     40           1.2650           nan      0.1000     0.0002
##     60           0.8721           nan      0.1000     0.0013
##     80           0.6979           nan      0.1000    -0.0035
##    100           0.5984           nan      0.1000    -0.0062
##    120           0.5296           nan      0.1000    -0.0059
##    140           0.4772           nan      0.1000    -0.0068
##    160           0.4300           nan      0.1000    -0.0125
##    180           0.3866           nan      0.1000    -0.0057
##    200           0.3463           nan      0.1000    -0.0064
##    220           0.3136           nan      0.1000    -0.0027
##    240           0.2853           nan      0.1000    -0.0041
##    250           0.2719           nan      0.1000    -0.0052
##
## Iter   TrainDeviance   ValidDeviance   StepSize   Improve
##      1           7.4823           nan      0.1000     0.6726
##      2           6.8225           nan      0.1000     0.6054
##      3           6.2303           nan      0.1000     0.4431
##      4           5.7262           nan      0.1000     0.3136
##      5           5.2917           nan      0.1000     0.3020
##      6           4.9481           nan      0.1000     0.2905
##      7           4.5620           nan      0.1000     0.2910
##      8           4.3225           nan      0.1000     0.1791
##      9           4.0699           nan      0.1000     0.2349
##     10           3.8179           nan      0.1000     0.1849
##     20           2.3059           nan      0.1000     0.0418
##     40           1.1688           nan      0.1000    -0.0035
##     60           0.7851           nan      0.1000     0.0072
##     80           0.6170           nan      0.1000    -0.0038
##    100           0.5177           nan      0.1000    -0.0026
##    120           0.4381           nan      0.1000    -0.0056
##    140           0.3801           nan      0.1000    -0.0035

```

```

##      160      0.3340      nan      0.1000     -0.0031
##      180      0.2933      nan      0.1000     -0.0080
##      200      0.2558      nan      0.1000     -0.0058
##      220      0.2289      nan      0.1000     -0.0028
##      240      0.2021      nan      0.1000     -0.0050
##      250      0.1882      nan      0.1000     -0.0018
##
## Iter   TrainDeviance   ValidDeviance   StepSize   Improve
##      1           7.4102           nan      0.1000     0.3839
##      2           7.0494           nan      0.1000     0.3040
##      3           6.8614           nan      0.1000     0.1707
##      4           6.6267           nan      0.1000     0.2365
##      5           6.4159           nan      0.1000     0.2192
##      6           6.2072           nan      0.1000     0.1818
##      7           5.9995           nan      0.1000     0.1441
##      8           5.8238           nan      0.1000     0.1470
##      9           5.6784           nan      0.1000     0.0881
##     10           5.5111           nan      0.1000     0.1191
##     20           4.4453           nan      0.1000     0.0629
##     40           3.2797           nan      0.1000     0.0303
##     60           2.5893           nan      0.1000     0.0080
##     80           2.1025           nan      0.1000    -0.0054
##    100           1.7328           nan      0.1000     0.0016
##    120           1.5056           nan      0.1000     0.0039
##    140           1.3304           nan      0.1000    -0.0021
##    160           1.2081           nan      0.1000    -0.0085
##    180           1.1043           nan      0.1000     0.0002
##    200           1.0182           nan      0.1000     0.0011
##    220           0.9519           nan      0.1000    -0.0138
##    240           0.9161           nan      0.1000    -0.0036
##    250           0.9019           nan      0.1000    -0.0080
##
## Iter   TrainDeviance   ValidDeviance   StepSize   Improve
##      1           7.3296           nan      0.1000     0.5238
##      2           6.8236           nan      0.1000     0.4935
##      3           6.4024           nan      0.1000     0.3814
##      4           6.0460           nan      0.1000     0.3352
##      5           5.7799           nan      0.1000     0.2079
##      6           5.5073           nan      0.1000     0.2249
##      7           5.2509           nan      0.1000     0.1626
##      8           5.0841           nan      0.1000     0.0927
##      9           4.8799           nan      0.1000     0.1436
##     10           4.7477           nan      0.1000     0.0329
##     20           3.3287           nan      0.1000     0.0705
##     40           2.0283           nan      0.1000     0.0048
##     60           1.3816           nan      0.1000     0.0072

```

##	80	1.0728	nan	0.1000	-0.0016
##	100	0.9094	nan	0.1000	0.0076
##	120	0.8114	nan	0.1000	-0.0058
##	140	0.7446	nan	0.1000	-0.0069
##	160	0.7042	nan	0.1000	-0.0008
##	180	0.6679	nan	0.1000	-0.0038
##	200	0.6315	nan	0.1000	-0.0026
##	220	0.6063	nan	0.1000	-0.0076
##	240	0.5797	nan	0.1000	-0.0034
##	250	0.5662	nan	0.1000	-0.0046

##

##	Iter	TrainDeviance	ValidDeviance	StepSize	Improve
##	1	7.2648	nan	0.1000	0.6118
##	2	6.7163	nan	0.1000	0.5338
##	3	6.2474	nan	0.1000	0.4288
##	4	5.8703	nan	0.1000	0.2640
##	5	5.5532	nan	0.1000	0.3193
##	6	5.2463	nan	0.1000	0.3198
##	7	4.9416	nan	0.1000	0.1943
##	8	4.6990	nan	0.1000	0.1722
##	9	4.4949	nan	0.1000	0.1469
##	10	4.2995	nan	0.1000	0.1804
##	20	2.7721	nan	0.1000	0.0861
##	40	1.4803	nan	0.1000	0.0123
##	60	1.0423	nan	0.1000	-0.0095
##	80	0.8424	nan	0.1000	0.0072
##	100	0.7137	nan	0.1000	-0.0079
##	120	0.6440	nan	0.1000	-0.0074
##	140	0.5827	nan	0.1000	-0.0044
##	160	0.5432	nan	0.1000	-0.0097
##	180	0.5079	nan	0.1000	-0.0044
##	200	0.4749	nan	0.1000	-0.0006
##	220	0.4448	nan	0.1000	-0.0068
##	240	0.4197	nan	0.1000	-0.0041
##	250	0.4073	nan	0.1000	-0.0040

##

##	Iter	TrainDeviance	ValidDeviance	StepSize	Improve
##	1	7.1301	nan	0.1000	0.6518
##	2	6.4611	nan	0.1000	0.5540
##	3	5.8979	nan	0.1000	0.4264
##	4	5.4531	nan	0.1000	0.3997
##	5	5.0649	nan	0.1000	0.2551
##	6	4.7495	nan	0.1000	0.2099
##	7	4.4799	nan	0.1000	0.2317
##	8	4.1899	nan	0.1000	0.2269
##	9	4.0084	nan	0.1000	0.0679

```

##      10      3.8183      nan      0.1000      0.0965
##      20      2.3118      nan      0.1000      0.0409
##      40      1.2467      nan      0.1000      0.0054
##      60      0.8880      nan      0.1000     -0.0007
##      80      0.7226      nan      0.1000     -0.0048
##     100      0.6209      nan      0.1000     -0.0051
##     120      0.5444      nan      0.1000     -0.0087
##     140      0.4935      nan      0.1000     -0.0049
##     160      0.4445      nan      0.1000     -0.0053
##     180      0.3989      nan      0.1000     -0.0063
##     200      0.3638      nan      0.1000     -0.0079
##     220      0.3337      nan      0.1000     -0.0052
##     240      0.3055      nan      0.1000     -0.0048
##     250      0.2926      nan      0.1000     -0.0020
##
## Iter   TrainDeviance   ValidDeviance   StepSize   Improve
##      1           7.1126           nan      0.1000    0.7800
##      2           6.4349           nan      0.1000    0.5736
##      3           5.8711           nan      0.1000    0.5557
##      4           5.4355           nan      0.1000    0.3727
##      5           5.0058           nan      0.1000    0.2490
##      6           4.5813           nan      0.1000    0.2892
##      7           4.2200           nan      0.1000    0.3013
##      8           3.8990           nan      0.1000    0.2893
##      9           3.6677           nan      0.1000    0.1964
##     10           3.4290           nan      0.1000    0.1762
##     20           2.1135           nan      0.1000    0.0418
##     40           1.0761           nan      0.1000   -0.0015
##     60           0.7465           nan      0.1000   -0.0119
##     80           0.6002           nan      0.1000   -0.0082
##    100           0.4915           nan      0.1000   -0.0042
##    120           0.4217           nan      0.1000   -0.0055
##    140           0.3558           nan      0.1000   -0.0027
##    160           0.3103           nan      0.1000   -0.0039
##    180           0.2649           nan      0.1000   -0.0052
##    200           0.2323           nan      0.1000   -0.0017
##    220           0.2076           nan      0.1000   -0.0018
##    240           0.1858           nan      0.1000   -0.0025
##    250           0.1740           nan      0.1000   -0.0033
##
## Iter   TrainDeviance   ValidDeviance   StepSize   Improve
##      1           7.3456           nan      0.1000    0.2909
##      2           7.1114           nan      0.1000    0.1263
##      3           6.7570           nan      0.1000    0.2751
##      4           6.5482           nan      0.1000    0.1337
##      5           6.3257           nan      0.1000    0.2345

```

##	6	6.1206	nan	0.1000	0.1383
##	7	5.9301	nan	0.1000	0.1914
##	8	5.7797	nan	0.1000	0.1295
##	9	5.6019	nan	0.1000	0.1135
##	10	5.4685	nan	0.1000	0.0847
##	20	4.4722	nan	0.1000	0.0460
##	40	3.3338	nan	0.1000	0.0101
##	60	2.6477	nan	0.1000	0.0048
##	80	2.1648	nan	0.1000	0.0153
##	100	1.7916	nan	0.1000	-0.0011
##	120	1.5267	nan	0.1000	0.0018
##	140	1.3281	nan	0.1000	-0.0029
##	160	1.1995	nan	0.1000	-0.0011
##	180	1.1018	nan	0.1000	0.0001
##	200	1.0288	nan	0.1000	-0.0059
##	220	0.9667	nan	0.1000	-0.0033
##	240	0.9174	nan	0.1000	-0.0045
##	250	0.8974	nan	0.1000	-0.0025

##	Iter	TrainDeviance	ValidDeviance	StepSize	Improve
##	1	7.1892	nan	0.1000	0.5473
##	2	6.7512	nan	0.1000	0.4586
##	3	6.4426	nan	0.1000	0.2760
##	4	6.1358	nan	0.1000	0.2140
##	5	5.8079	nan	0.1000	0.2639
##	6	5.5512	nan	0.1000	0.2127
##	7	5.3613	nan	0.1000	0.1317
##	8	5.0590	nan	0.1000	0.2354
##	9	4.9117	nan	0.1000	0.1361
##	10	4.7130	nan	0.1000	0.1626
##	20	3.4020	nan	0.1000	0.0257
##	40	2.0751	nan	0.1000	0.0224
##	60	1.4101	nan	0.1000	-0.0005
##	80	1.1014	nan	0.1000	0.0065
##	100	0.9405	nan	0.1000	-0.0067
##	120	0.8391	nan	0.1000	-0.0066
##	140	0.7718	nan	0.1000	-0.0054
##	160	0.7291	nan	0.1000	-0.0095
##	180	0.6810	nan	0.1000	-0.0036
##	200	0.6457	nan	0.1000	-0.0071
##	220	0.6189	nan	0.1000	-0.0034
##	240	0.5895	nan	0.1000	-0.0023
##	250	0.5794	nan	0.1000	-0.0084

##	Iter	TrainDeviance	ValidDeviance	StepSize	Improve
##	1	7.1359	nan	0.1000	0.6309

```

##      2      6.5967      nan    0.1000    0.4591
##      3      6.1698      nan    0.1000    0.4132
##      4      5.7599      nan    0.1000    0.3928
##      5      5.4519      nan    0.1000    0.3237
##      6      5.1401      nan    0.1000    0.2872
##      7      4.9050      nan    0.1000    0.1392
##      8      4.6196      nan    0.1000    0.2846
##      9      4.3738      nan    0.1000    0.1828
##     10      4.1835      nan    0.1000    0.1700
##     20      2.8099      nan    0.1000    0.0580
##     40      1.6151      nan    0.1000   -0.0065
##     60      1.1301      nan    0.1000    0.0002
##     80      0.8944      nan    0.1000   -0.0077
##    100      0.7517      nan    0.1000   -0.0139
##    120      0.6730      nan    0.1000   -0.0085
##    140      0.6057      nan    0.1000   -0.0038
##    160      0.5547      nan    0.1000   -0.0097
##    180      0.5094      nan    0.1000   -0.0103
##    200      0.4766      nan    0.1000   -0.0066
##    220      0.4450      nan    0.1000   -0.0040
##    240      0.4151      nan    0.1000   -0.0027
##    250      0.4013      nan    0.1000   -0.0047
##
## Iter  TrainDeviance  ValidDeviance  StepSize  Improve
##      1      7.0294      nan    0.1000    0.7373
##      2      6.4316      nan    0.1000    0.6149
##      3      5.9113      nan    0.1000    0.3359
##      4      5.5546      nan    0.1000    0.3535
##      5      5.1917      nan    0.1000    0.2802
##      6      4.8371      nan    0.1000    0.3563
##      7      4.5468      nan    0.1000    0.2089
##      8      4.3248      nan    0.1000    0.1893
##      9      4.0651      nan    0.1000    0.1052
##     10      3.8138      nan    0.1000    0.1666
##     20      2.3320      nan    0.1000    0.0422
##     40      1.2332      nan    0.1000    0.0086
##     60      0.8464      nan    0.1000    0.0058
##     80      0.6737      nan    0.1000   -0.0128
##    100      0.5934      nan    0.1000   -0.0097
##    120      0.5242      nan    0.1000   -0.0050
##    140      0.4705      nan    0.1000   -0.0050
##    160      0.4233      nan    0.1000   -0.0078
##    180      0.3867      nan    0.1000   -0.0076
##    200      0.3531      nan    0.1000   -0.0040
##    220      0.3242      nan    0.1000   -0.0039
##    240      0.3013      nan    0.1000   -0.0036

```

```

##      250      0.2843      nan      0.1000     -0.0029
##
## Iter   TrainDeviance   ValidDeviance   StepSize   Improve
##      1      7.0098      nan      0.1000     0.5845
##      2      6.3343      nan      0.1000     0.4158
##      3      5.8514      nan      0.1000     0.3832
##      4      5.3437      nan      0.1000     0.4617
##      5      4.9793      nan      0.1000     0.3545
##      6      4.6111      nan      0.1000     0.3660
##      7      4.2957      nan      0.1000     0.2518
##      8      3.9570      nan      0.1000     0.2273
##      9      3.7320      nan      0.1000     0.1854
##     10      3.5140      nan      0.1000     0.1503
##     20      2.1057      nan      0.1000     0.0764
##     40      1.0592      nan      0.1000     0.0046
##     60      0.7323      nan      0.1000    -0.0010
##     80      0.5881      nan      0.1000    -0.0058
##    100      0.4956      nan      0.1000    -0.0051
##    120      0.4248      nan      0.1000    -0.0025
##    140      0.3718      nan      0.1000    -0.0073
##    160      0.3294      nan      0.1000    -0.0062
##    180      0.2819      nan      0.1000    -0.0026
##    200      0.2497      nan      0.1000    -0.0034
##    220      0.2233      nan      0.1000    -0.0058
##    240      0.1990      nan      0.1000    -0.0018
##    250      0.1871      nan      0.1000    -0.0021
##
## Iter   TrainDeviance   ValidDeviance   StepSize   Improve
##      1      7.5868      nan      0.1000     0.4115
##      2      7.2833      nan      0.1000     0.2331
##      3      7.0345      nan      0.1000     0.2570
##      4      6.8101      nan      0.1000     0.1071
##      5      6.5893      nan      0.1000     0.1563
##      6      6.3092      nan      0.1000     0.1664
##      7      6.1025      nan      0.1000     0.1476
##      8      5.9663      nan      0.1000     0.0749
##      9      5.7474      nan      0.1000     0.1445
##     10      5.5960      nan      0.1000     0.0876
##     20      4.4901      nan      0.1000     0.0532
##     40      3.2925      nan      0.1000     0.0364
##     60      2.6190      nan      0.1000     0.0090
##     80      2.1208      nan      0.1000     0.0132
##    100      1.7732      nan      0.1000    -0.0007
##    120      1.5132      nan      0.1000     0.0046
##    140      1.3283      nan      0.1000     0.0031
##    160      1.1925      nan      0.1000    -0.0002

```



```

##      180      1.0847      nan      0.1000     -0.0013
##      200      0.9981      nan      0.1000     -0.0003
##      220      0.9475      nan      0.1000     -0.0035
##      240      0.9021      nan      0.1000     -0.0047
##      250      0.8843      nan      0.1000     -0.0029
##
## Iter   TrainDeviance   ValidDeviance   StepSize   Improve
##      1           7.3282           nan      0.1000     0.5172
##      2           6.9512           nan      0.1000     0.3872
##      3           6.5318           nan      0.1000     0.3539
##      4           6.1978           nan      0.1000     0.2757
##      5           5.9066           nan      0.1000     0.2460
##      6           5.6366           nan      0.1000     0.2517
##      7           5.3731           nan      0.1000     0.1878
##      8           5.1982           nan      0.1000     0.1405
##      9           5.0115           nan      0.1000     0.1659
##     10           4.8495           nan      0.1000     0.1095
##     20           3.3891           nan      0.1000     0.1033
##     40           2.0749           nan      0.1000     0.0310
##     60           1.4400           nan      0.1000     0.0118
##     80           1.1141           nan      0.1000     0.0030
##    100           0.9360           nan      0.1000     0.0034
##    120           0.8234           nan      0.1000    -0.0051
##    140           0.7648           nan      0.1000    -0.0025
##    160           0.7108           nan      0.1000    -0.0075
##    180           0.6698           nan      0.1000    -0.0052
##    200           0.6366           nan      0.1000    -0.0069
##    220           0.6020           nan      0.1000    -0.0060
##    240           0.5770           nan      0.1000    -0.0049
##    250           0.5640           nan      0.1000    -0.0035
##
## Iter   TrainDeviance   ValidDeviance   StepSize   Improve
##      1           7.2616           nan      0.1000     0.6773
##      2           6.8094           nan      0.1000     0.2167
##      3           6.3219           nan      0.1000     0.4713
##      4           5.9148           nan      0.1000     0.3566
##      5           5.5959           nan      0.1000     0.2337
##      6           5.2724           nan      0.1000     0.3200
##      7           4.9638           nan      0.1000     0.1900
##      8           4.7283           nan      0.1000     0.1678
##      9           4.5198           nan      0.1000     0.1260
##     10           4.3024           nan      0.1000     0.1451
##     20           2.7404           nan      0.1000     0.0379
##     40           1.5392           nan      0.1000    -0.0055
##     60           1.0399           nan      0.1000     0.0123
##     80           0.8192           nan      0.1000     0.0004

```

##	100	0.7016	nan	0.1000	-0.0083
##	120	0.6325	nan	0.1000	-0.0080
##	140	0.5744	nan	0.1000	-0.0063
##	160	0.5259	nan	0.1000	-0.0054
##	180	0.4819	nan	0.1000	-0.0111
##	200	0.4502	nan	0.1000	-0.0055
##	220	0.4164	nan	0.1000	-0.0031
##	240	0.3874	nan	0.1000	-0.0056
##	250	0.3749	nan	0.1000	-0.0012

##	Iter	TrainDeviance	ValidDeviance	StepSize	Improve
##	1	7.4180	nan	0.1000	0.3827
##	2	6.7805	nan	0.1000	0.4994
##	3	6.2576	nan	0.1000	0.4550
##	4	5.8389	nan	0.1000	0.2288
##	5	5.4627	nan	0.1000	0.3687
##	6	5.1119	nan	0.1000	0.3303
##	7	4.7963	nan	0.1000	0.3248
##	8	4.5143	nan	0.1000	0.2401
##	9	4.2403	nan	0.1000	0.1986
##	10	4.0630	nan	0.1000	0.1325
##	20	2.4827	nan	0.1000	0.0652
##	40	1.2766	nan	0.1000	0.0153
##	60	0.8585	nan	0.1000	-0.0103
##	80	0.6931	nan	0.1000	-0.0009
##	100	0.5966	nan	0.1000	-0.0085
##	120	0.5283	nan	0.1000	-0.0111
##	140	0.4718	nan	0.1000	-0.0106
##	160	0.4263	nan	0.1000	-0.0058
##	180	0.3870	nan	0.1000	-0.0051
##	200	0.3533	nan	0.1000	-0.0027
##	220	0.3194	nan	0.1000	-0.0032
##	240	0.2922	nan	0.1000	-0.0019
##	250	0.2785	nan	0.1000	-0.0040

##	Iter	TrainDeviance	ValidDeviance	StepSize	Improve
##	1	7.1462	nan	0.1000	0.6854
##	2	6.5338	nan	0.1000	0.3961
##	3	6.0020	nan	0.1000	0.4159
##	4	5.5196	nan	0.1000	0.3879
##	5	5.0940	nan	0.1000	0.3601
##	6	4.7927	nan	0.1000	0.2963
##	7	4.4249	nan	0.1000	0.2827
##	8	4.1084	nan	0.1000	0.2133
##	9	3.8253	nan	0.1000	0.1881
##	10	3.6343	nan	0.1000	0.1233

```

##      20      2.1522      nan      0.1000      0.0354
##      40      1.0580      nan      0.1000      0.0161
##      60      0.7191      nan      0.1000      0.0018
##      80      0.5670      nan      0.1000     -0.0057
##     100      0.4832      nan      0.1000     -0.0050
##     120      0.4124      nan      0.1000     -0.0054
##     140      0.3602      nan      0.1000     -0.0078
##     160      0.3179      nan      0.1000     -0.0036
##     180      0.2778      nan      0.1000     -0.0047
##     200      0.2450      nan      0.1000     -0.0059
##     220      0.2183      nan      0.1000     -0.0040
##     240      0.1929      nan      0.1000     -0.0048
##     250      0.1839      nan      0.1000     -0.0040
##
## Iter  TrainDeviance  ValidDeviance  StepSize  Improve
##      1          7.4540          nan      0.1000    0.3641
##      2          7.1373          nan      0.1000    0.3291
##      3          6.8416          nan      0.1000    0.2730
##      4          6.5864          nan      0.1000    0.1830
##      5          6.3867          nan      0.1000    0.1603
##      6          6.1413          nan      0.1000    0.2019
##      7          5.9610          nan      0.1000    0.1395
##      8          5.8194          nan      0.1000    0.1027
##      9          5.6546          nan      0.1000    0.0827
##     10          5.4631          nan      0.1000    0.0971
##     20          4.3922          nan      0.1000    0.0188
##     40          3.2426          nan      0.1000    0.0345
##     60          2.5668          nan      0.1000    0.0223
##     80          2.0713          nan      0.1000    0.0088
##    100          1.7438          nan      0.1000   -0.0065
##    120          1.4921          nan      0.1000    0.0043
##    140          1.3219          nan      0.1000    0.0085
##    160          1.1951          nan      0.1000    0.0011
##    180          1.1011          nan      0.1000   -0.0045
##    200          1.0331          nan      0.1000   -0.0013
##    220          0.9765          nan      0.1000   -0.0020
##    240          0.9362          nan      0.1000   -0.0053
##    250          0.9130          nan      0.1000    0.0009
##
## Iter  TrainDeviance  ValidDeviance  StepSize  Improve
##      1          7.2752          nan      0.1000    0.5651
##      2          6.7895          nan      0.1000    0.4457
##      3          6.4533          nan      0.1000    0.2484
##      4          6.1292          nan      0.1000    0.3299
##      5          5.8643          nan      0.1000    0.2757
##      6          5.6456          nan      0.1000    0.1455

```

##	7	5.3780	nan	0.1000	0.2465
##	8	5.1363	nan	0.1000	0.1833
##	9	4.9400	nan	0.1000	0.1011
##	10	4.7344	nan	0.1000	0.1801
##	20	3.2834	nan	0.1000	0.0459
##	40	2.0454	nan	0.1000	0.0186
##	60	1.4428	nan	0.1000	0.0008
##	80	1.0931	nan	0.1000	0.0050
##	100	0.9245	nan	0.1000	-0.0066
##	120	0.8152	nan	0.1000	-0.0040
##	140	0.7456	nan	0.1000	-0.0028
##	160	0.6979	nan	0.1000	-0.0052
##	180	0.6573	nan	0.1000	-0.0035
##	200	0.6239	nan	0.1000	-0.0037
##	220	0.5949	nan	0.1000	-0.0042
##	240	0.5695	nan	0.1000	-0.0039
##	250	0.5578	nan	0.1000	-0.0039
##					
##	Iter	TrainDeviance	ValidDeviance	StepSize	Improve
##	1	7.2004	nan	0.1000	0.5575
##	2	6.6875	nan	0.1000	0.5080
##	3	6.1988	nan	0.1000	0.4364
##	4	5.7727	nan	0.1000	0.3191
##	5	5.4082	nan	0.1000	0.3497
##	6	5.1722	nan	0.1000	0.1589
##	7	4.8592	nan	0.1000	0.1847
##	8	4.6133	nan	0.1000	0.2159
##	9	4.3824	nan	0.1000	0.2002
##	10	4.1895	nan	0.1000	0.1182
##	20	2.7390	nan	0.1000	0.0733
##	40	1.5162	nan	0.1000	-0.0008
##	60	1.0207	nan	0.1000	0.0026
##	80	0.8292	nan	0.1000	-0.0016
##	100	0.7057	nan	0.1000	-0.0089
##	120	0.6346	nan	0.1000	-0.0085
##	140	0.5675	nan	0.1000	-0.0056
##	160	0.5176	nan	0.1000	-0.0074
##	180	0.4752	nan	0.1000	-0.0050
##	200	0.4333	nan	0.1000	-0.0054
##	220	0.4032	nan	0.1000	-0.0070
##	240	0.3743	nan	0.1000	-0.0022
##	250	0.3642	nan	0.1000	-0.0062
##					
##	Iter	TrainDeviance	ValidDeviance	StepSize	Improve
##	1	7.1706	nan	0.1000	0.6334
##	2	6.5472	nan	0.1000	0.5452

```

##      3      5.9610      nan      0.1000      0.4332
##      4      5.5223      nan      0.1000      0.3950
##      5      5.1475      nan      0.1000      0.2716
##      6      4.8185      nan      0.1000      0.2069
##      7      4.5333      nan      0.1000      0.2125
##      8      4.2886      nan      0.1000      0.2125
##      9      4.0286      nan      0.1000      0.1507
##     10      3.8317      nan      0.1000      0.1457
##     20      2.3895      nan      0.1000      0.0665
##     40      1.2552      nan      0.1000      0.0223
##     60      0.8795      nan      0.1000      0.0039
##     80      0.7069      nan      0.1000     -0.0107
##    100      0.6120      nan      0.1000     -0.0067
##    120      0.5331      nan      0.1000     -0.0065
##    140      0.4731      nan      0.1000     -0.0091
##    160      0.4256      nan      0.1000     -0.0038
##    180      0.3789      nan      0.1000     -0.0053
##    200      0.3447      nan      0.1000     -0.0045
##    220      0.3123      nan      0.1000     -0.0067
##    240      0.2807      nan      0.1000     -0.0018
##    250      0.2696      nan      0.1000     -0.0060
##
## Iter  TrainDeviance  ValidDeviance  StepSize  Improve
##      1      7.1181      nan      0.1000      0.6333
##      2      6.4637      nan      0.1000      0.5209
##      3      5.8982      nan      0.1000      0.5202
##      4      5.4003      nan      0.1000      0.3026
##      5      4.9882      nan      0.1000      0.2926
##      6      4.6393      nan      0.1000      0.2601
##      7      4.3294      nan      0.1000      0.2207
##      8      4.0488      nan      0.1000      0.2276
##      9      3.7542      nan      0.1000      0.2489
##     10      3.5442      nan      0.1000      0.1313
##     20      2.0640      nan      0.1000      0.0561
##     40      1.0503      nan      0.1000      0.0116
##     60      0.7338      nan      0.1000     -0.0011
##     80      0.5913      nan      0.1000     -0.0046
##    100      0.4891      nan      0.1000     -0.0047
##    120      0.4050      nan      0.1000     -0.0044
##    140      0.3428      nan      0.1000     -0.0083
##    160      0.2935      nan      0.1000     -0.0027
##    180      0.2561      nan      0.1000      0.0005
##    200      0.2280      nan      0.1000     -0.0051
##    220      0.1969      nan      0.1000     -0.0034
##    240      0.1764      nan      0.1000     -0.0047
##    250      0.1670      nan      0.1000     -0.0035

```

```

##
## Iter    TrainDeviance    ValidDeviance    StepSize    Improve
##      1          7.9913          nan      0.1000      0.3761
##      2          7.7852          nan      0.1000      0.1808
##      3          7.4353          nan      0.1000      0.3342
##      4          7.2002          nan      0.1000      0.1743
##      5          6.8764          nan      0.1000      0.3014
##      6          6.7249          nan      0.1000      0.1377
##      7          6.5608          nan      0.1000      0.0969
##      8          6.3857          nan      0.1000      0.1119
##      9          6.2081          nan      0.1000      0.2099
##     10          6.0676          nan      0.1000      0.0350
##     20          4.8481          nan      0.1000      0.0326
##     40          3.5311          nan      0.1000      0.0243
##     60          2.7637          nan      0.1000     -0.0042
##     80          2.2298          nan      0.1000      0.0035
##    100          1.8541          nan      0.1000      0.0040
##    120          1.5720          nan      0.1000     -0.0142
##    140          1.3723          nan      0.1000      0.0089
##    160          1.2296          nan      0.1000     -0.0073
##    180          1.1206          nan      0.1000     -0.0095
##    200          1.0353          nan      0.1000     -0.0009
##    220          0.9629          nan      0.1000     -0.0014
##    240          0.9290          nan      0.1000     -0.0011
##    250          0.9156          nan      0.1000     -0.0036
##
## Iter    TrainDeviance    ValidDeviance    StepSize    Improve
##      1          7.7709          nan      0.1000      0.4607
##      2          7.2930          nan      0.1000      0.4519
##      3          6.7154          nan      0.1000      0.2983
##      4          6.4052          nan      0.1000      0.2385
##      5          6.0937          nan      0.1000      0.2545
##      6          5.7996          nan      0.1000      0.2344
##      7          5.6159          nan      0.1000      0.1053
##      8          5.4356          nan      0.1000      0.1415
##      9          5.2146          nan      0.1000      0.1220
##     10          5.0671          nan      0.1000      0.1222
##     20          3.5111          nan      0.1000      0.0306
##     40          2.0860          nan      0.1000      0.0249
##     60          1.4474          nan      0.1000      0.0088
##     80          1.1308          nan      0.1000     -0.0062
##    100          0.9648          nan      0.1000      0.0020
##    120          0.8736          nan      0.1000     -0.0089
##    140          0.8043          nan      0.1000     -0.0064
##    160          0.7498          nan      0.1000     -0.0083
##    180          0.7018          nan      0.1000      0.0001

```

```

##      200      0.6598      nan      0.1000     -0.0037
##      220      0.6329      nan      0.1000     -0.0035
##      240      0.6049      nan      0.1000     -0.0068
##      250      0.5905      nan      0.1000     -0.0046
##
## Iter   TrainDeviance   ValidDeviance   StepSize   Improve
##      1           7.5947           nan      0.1000     0.5713
##      2           6.9919           nan      0.1000     0.5903
##      3           6.5417           nan      0.1000     0.4225
##      4           6.0968           nan      0.1000     0.3213
##      5           5.7124           nan      0.1000     0.3371
##      6           5.3694           nan      0.1000     0.2772
##      7           5.0464           nan      0.1000     0.2062
##      8           4.7692           nan      0.1000     0.2608
##      9           4.6153           nan      0.1000     0.0811
##     10           4.4195           nan      0.1000     0.0977
##     20           2.9502           nan      0.1000     0.1132
##     40           1.5511           nan      0.1000     0.0204
##     60           1.0570           nan      0.1000     0.0142
##     80           0.8503           nan      0.1000    -0.0047
##    100           0.7300           nan      0.1000    -0.0177
##    120           0.6452           nan      0.1000    -0.0077
##    140           0.5915           nan      0.1000    -0.0082
##    160           0.5400           nan      0.1000    -0.0058
##    180           0.5009           nan      0.1000    -0.0040
##    200           0.4669           nan      0.1000    -0.0059
##    220           0.4400           nan      0.1000    -0.0040
##    240           0.4131           nan      0.1000    -0.0079
##    250           0.4039           nan      0.1000    -0.0059
##
## Iter   TrainDeviance   ValidDeviance   StepSize   Improve
##      1           7.6580           nan      0.1000     0.6541
##      2           7.0177           nan      0.1000     0.5653
##      3           6.4657           nan      0.1000     0.5278
##      4           5.9869           nan      0.1000     0.4523
##      5           5.5715           nan      0.1000     0.3559
##      6           5.1874           nan      0.1000     0.2871
##      7           4.8558           nan      0.1000     0.2317
##      8           4.5402           nan      0.1000     0.2236
##      9           4.3011           nan      0.1000     0.1610
##     10           4.0022           nan      0.1000     0.2235
##     20           2.4352           nan      0.1000     0.0364
##     40           1.2432           nan      0.1000     0.0032
##     60           0.8629           nan      0.1000    -0.0036
##     80           0.7137           nan      0.1000    -0.0020
##    100           0.6171           nan      0.1000    -0.0102

```

##	120	0.5389	nan	0.1000	-0.0066
##	140	0.4772	nan	0.1000	-0.0067
##	160	0.4290	nan	0.1000	-0.0091
##	180	0.3873	nan	0.1000	-0.0116
##	200	0.3411	nan	0.1000	-0.0052
##	220	0.3122	nan	0.1000	-0.0030
##	240	0.2868	nan	0.1000	-0.0020
##	250	0.2726	nan	0.1000	-0.0027

##

##	Iter	TrainDeviance	ValidDeviance	StepSize	Improve
##	1	7.4799	nan	0.1000	0.8543
##	2	6.8255	nan	0.1000	0.5990
##	3	6.2454	nan	0.1000	0.4987
##	4	5.7558	nan	0.1000	0.3832
##	5	5.3790	nan	0.1000	0.3350
##	6	5.0578	nan	0.1000	0.2406
##	7	4.7324	nan	0.1000	0.1941
##	8	4.3901	nan	0.1000	0.2642
##	9	4.0528	nan	0.1000	0.2389
##	10	3.7874	nan	0.1000	0.2426
##	20	2.1952	nan	0.1000	0.0768
##	40	1.1025	nan	0.1000	-0.0081
##	60	0.7689	nan	0.1000	-0.0002
##	80	0.6122	nan	0.1000	-0.0061
##	100	0.5200	nan	0.1000	-0.0109
##	120	0.4400	nan	0.1000	-0.0069
##	140	0.3837	nan	0.1000	-0.0055
##	160	0.3328	nan	0.1000	-0.0034
##	180	0.2920	nan	0.1000	-0.0077
##	200	0.2544	nan	0.1000	-0.0035
##	220	0.2241	nan	0.1000	-0.0036
##	240	0.1988	nan	0.1000	-0.0047
##	250	0.1865	nan	0.1000	-0.0028

##

##	Iter	TrainDeviance	ValidDeviance	StepSize	Improve
##	1	7.2748	nan	0.1000	0.3849
##	2	6.9549	nan	0.1000	0.3086
##	3	6.6188	nan	0.1000	0.2131
##	4	6.3662	nan	0.1000	0.2414
##	5	6.1655	nan	0.1000	0.1744
##	6	5.9621	nan	0.1000	0.1800
##	7	5.7756	nan	0.1000	0.1729
##	8	5.6347	nan	0.1000	0.1348
##	9	5.5429	nan	0.1000	0.0634
##	10	5.3761	nan	0.1000	0.0928
##	20	4.3688	nan	0.1000	0.0591


```

##      40      3.2068      nan    0.1000    0.0402
##      60      2.5475      nan    0.1000    0.0114
##      80      2.0354      nan    0.1000    0.0023
##     100      1.7011      nan    0.1000   -0.0072
##     120      1.4362      nan    0.1000   -0.0015
##     140      1.2570      nan    0.1000   -0.0034
##     160      1.1369      nan    0.1000    0.0021
##     180      1.0223      nan    0.1000   -0.0062
##     200      0.9563      nan    0.1000   -0.0038
##     220      0.8991      nan    0.1000   -0.0058
##     240      0.8598      nan    0.1000   -0.0059
##     250      0.8460      nan    0.1000   -0.0086
##
## Iter   TrainDeviance   ValidDeviance   StepSize   Improve
##      1           7.1432           nan    0.1000    0.5851
##      2           6.7241           nan    0.1000    0.4090
##      3           6.2802           nan    0.1000    0.3640
##      4           5.9965           nan    0.1000    0.1559
##      5           5.7218           nan    0.1000    0.2864
##      6           5.4669           nan    0.1000    0.2326
##      7           5.1666           nan    0.1000    0.2190
##      8           4.9344           nan    0.1000    0.2265
##      9           4.7538           nan    0.1000    0.1466
##     10           4.5570           nan    0.1000    0.1311
##     20           3.2524           nan    0.1000    0.1132
##     40           1.9609           nan    0.1000    0.0491
##     60           1.3733           nan    0.1000   -0.0028
##     80           1.0764           nan    0.1000   -0.0083
##    100           0.8912           nan    0.1000   -0.0043
##    120           0.7917           nan    0.1000   -0.0084
##    140           0.7298           nan    0.1000   -0.0027
##    160           0.6894           nan    0.1000   -0.0061
##    180           0.6543           nan    0.1000   -0.0077
##    200           0.6175           nan    0.1000   -0.0009
##    220           0.5885           nan    0.1000   -0.0089
##    240           0.5540           nan    0.1000   -0.0036
##    250           0.5386           nan    0.1000   -0.0019
##
## Iter   TrainDeviance   ValidDeviance   StepSize   Improve
##      1           7.0286           nan    0.1000    0.5874
##      2           6.5422           nan    0.1000    0.4614
##      3           6.0807           nan    0.1000    0.4788
##      4           5.6220           nan    0.1000    0.3139
##      5           5.3339           nan    0.1000    0.1849
##      6           4.9990           nan    0.1000    0.2584
##      7           4.7023           nan    0.1000    0.2324

```

```

##      8      4.4521      nan      0.1000      0.2002
##      9      4.2254      nan      0.1000      0.1721
##     10      4.0158      nan      0.1000      0.1061
##     20      2.5603      nan      0.1000      0.0649
##     40      1.4633      nan      0.1000      0.0221
##     60      1.0126      nan      0.1000      0.0040
##     80      0.8089      nan      0.1000     -0.0011
##    100      0.6913      nan      0.1000     -0.0016
##    120      0.6212      nan      0.1000     -0.0039
##    140      0.5555      nan      0.1000      0.0013
##    160      0.5178      nan      0.1000     -0.0070
##    180      0.4779      nan      0.1000     -0.0062
##    200      0.4357      nan      0.1000     -0.0056
##    220      0.4076      nan      0.1000     -0.0034
##    240      0.3838      nan      0.1000     -0.0051
##    250      0.3724      nan      0.1000     -0.0045
##
## Iter  TrainDeviance  ValidDeviance  StepSize  Improve
##      1           7.0052           nan      0.1000      0.7357
##      2           6.3136           nan      0.1000      0.6391
##      3           5.8579           nan      0.1000      0.3395
##      4           5.4261           nan      0.1000      0.4118
##      5           5.0380           nan      0.1000      0.2541
##      6           4.7502           nan      0.1000      0.2364
##      7           4.5107           nan      0.1000      0.1393
##      8           4.2559           nan      0.1000      0.2325
##      9           4.0139           nan      0.1000      0.1431
##     10           3.7788           nan      0.1000      0.1690
##     20           2.2976           nan      0.1000      0.0670
##     40           1.1641           nan      0.1000      0.0258
##     60           0.7945           nan      0.1000     -0.0042
##     80           0.6385           nan      0.1000     -0.0010
##    100           0.5468           nan      0.1000     -0.0065
##    120           0.4711           nan      0.1000     -0.0070
##    140           0.4186           nan      0.1000     -0.0051
##    160           0.3792           nan      0.1000     -0.0062
##    180           0.3426           nan      0.1000     -0.0072
##    200           0.3149           nan      0.1000     -0.0066
##    220           0.2897           nan      0.1000     -0.0081
##    240           0.2642           nan      0.1000     -0.0020
##    250           0.2526           nan      0.1000     -0.0020
##
## Iter  TrainDeviance  ValidDeviance  StepSize  Improve
##      1           7.0074           nan      0.1000      0.6779
##      2           6.3049           nan      0.1000      0.6277
##      3           5.7428           nan      0.1000      0.5037

```

```

##      4      5.3843      nan    0.1000    0.2738
##      5      4.9412      nan    0.1000    0.4013
##      6      4.5370      nan    0.1000    0.2175
##      7      4.2296      nan    0.1000    0.2194
##      8      3.9306      nan    0.1000    0.2566
##      9      3.6976      nan    0.1000    0.1642
##     10      3.4742      nan    0.1000    0.1155
##     20      2.0717      nan    0.1000    0.0554
##     40      1.0961      nan    0.1000    0.0158
##     60      0.7387      nan    0.1000   -0.0071
##     80      0.5947      nan    0.1000   -0.0086
##    100      0.4878      nan    0.1000   -0.0035
##    120      0.4191      nan    0.1000   -0.0069
##    140      0.3558      nan    0.1000   -0.0047
##    160      0.3193      nan    0.1000   -0.0029
##    180      0.2757      nan    0.1000   -0.0060
##    200      0.2412      nan    0.1000   -0.0011
##    220      0.2154      nan    0.1000   -0.0022
##    240      0.1911      nan    0.1000   -0.0019
##    250      0.1798      nan    0.1000   -0.0025
##
## Iter  TrainDeviance  ValidDeviance  StepSize  Improve
##      1      7.5373      nan    0.1000    0.3255
##      2      7.2302      nan    0.1000    0.3081
##      3      7.0287      nan    0.1000    0.0792
##      4      6.8235      nan    0.1000    0.1648
##      5      6.5846      nan    0.1000    0.2528
##      6      6.3572      nan    0.1000    0.2142
##      7      6.2170      nan    0.1000    0.0731
##      8      6.0241      nan    0.1000    0.1693
##      9      5.8911      nan    0.1000    0.1379
##     10      5.7890      nan    0.1000    0.0396
##     20      4.6672      nan    0.1000    0.0610
##     40      3.3823      nan    0.1000    0.0241
##     60      2.6081      nan    0.1000    0.0076
##     80      2.1083      nan    0.1000    0.0027
##    100      1.7823      nan    0.1000   -0.0038
##    120      1.5329      nan    0.1000    0.0083
##    140      1.3371      nan    0.1000   -0.0004
##    160      1.1999      nan    0.1000   -0.0035
##    180      1.0969      nan    0.1000   -0.0001
##    200      1.0178      nan    0.1000   -0.0020
##    220      0.9601      nan    0.1000   -0.0065
##    240      0.9259      nan    0.1000   -0.0059
##    250      0.9093      nan    0.1000   -0.0037
##

```

##	Iter	TrainDeviance	ValidDeviance	StepSize	Improve
##	1	7.3367	nan	0.1000	0.5068
##	2	6.8403	nan	0.1000	0.4133
##	3	6.4575	nan	0.1000	0.3430
##	4	6.1040	nan	0.1000	0.2802
##	5	5.9071	nan	0.1000	0.1333
##	6	5.6534	nan	0.1000	0.2078
##	7	5.4050	nan	0.1000	0.2252
##	8	5.2113	nan	0.1000	0.1294
##	9	4.9828	nan	0.1000	0.1934
##	10	4.8171	nan	0.1000	0.1146
##	20	3.3732	nan	0.1000	0.0232
##	40	2.0529	nan	0.1000	0.0254
##	60	1.4630	nan	0.1000	0.0107
##	80	1.1301	nan	0.1000	-0.0184
##	100	0.9347	nan	0.1000	-0.0020
##	120	0.8346	nan	0.1000	-0.0053
##	140	0.7561	nan	0.1000	-0.0116
##	160	0.7038	nan	0.1000	-0.0083
##	180	0.6608	nan	0.1000	-0.0050
##	200	0.6231	nan	0.1000	-0.0062
##	220	0.5885	nan	0.1000	-0.0053
##	240	0.5666	nan	0.1000	-0.0073
##	250	0.5537	nan	0.1000	-0.0016
##					
##	Iter	TrainDeviance	ValidDeviance	StepSize	Improve
##	1	7.3914	nan	0.1000	0.5561
##	2	6.7379	nan	0.1000	0.5924
##	3	6.1931	nan	0.1000	0.3638
##	4	5.7353	nan	0.1000	0.2706
##	5	5.4272	nan	0.1000	0.2604
##	6	5.1269	nan	0.1000	0.1462
##	7	4.9406	nan	0.1000	0.0649
##	8	4.7122	nan	0.1000	0.1369
##	9	4.5181	nan	0.1000	0.1104
##	10	4.3030	nan	0.1000	0.1313
##	20	2.8485	nan	0.1000	0.0934
##	40	1.6130	nan	0.1000	-0.0081
##	60	1.1086	nan	0.1000	-0.0165
##	80	0.8744	nan	0.1000	-0.0020
##	100	0.7519	nan	0.1000	-0.0066
##	120	0.6638	nan	0.1000	-0.0039
##	140	0.5997	nan	0.1000	-0.0044
##	160	0.5553	nan	0.1000	-0.0082
##	180	0.5230	nan	0.1000	-0.0054
##	200	0.4907	nan	0.1000	-0.0056

```

##      220      0.4566      nan      0.1000     -0.0051
##      240      0.4234      nan      0.1000     -0.0047
##      250      0.4103      nan      0.1000     -0.0062
##
## Iter   TrainDeviance   ValidDeviance   StepSize   Improve
##      1           7.3959           nan      0.1000     0.5344
##      2           6.7840           nan      0.1000     0.6032
##      3           6.2042           nan      0.1000     0.5214
##      4           5.7561           nan      0.1000     0.3567
##      5           5.3542           nan      0.1000     0.3154
##      6           5.0748           nan      0.1000     0.2697
##      7           4.6856           nan      0.1000     0.3381
##      8           4.4209           nan      0.1000     0.2125
##      9           4.1834           nan      0.1000     0.1959
##     10           3.9558           nan      0.1000     0.1757
##     20           2.5107           nan      0.1000     0.0757
##     40           1.3385           nan      0.1000     0.0197
##     60           0.9193           nan      0.1000     0.0067
##     80           0.7247           nan      0.1000    -0.0105
##    100           0.6080           nan      0.1000    -0.0060
##    120           0.5370           nan      0.1000    -0.0086
##    140           0.4808           nan      0.1000    -0.0094
##    160           0.4331           nan      0.1000    -0.0018
##    180           0.3896           nan      0.1000    -0.0076
##    200           0.3487           nan      0.1000    -0.0060
##    220           0.3156           nan      0.1000    -0.0044
##    240           0.2857           nan      0.1000    -0.0023
##    250           0.2735           nan      0.1000    -0.0045
##
## Iter   TrainDeviance   ValidDeviance   StepSize   Improve
##      1           7.2002           nan      0.1000     0.7128
##      2           6.5508           nan      0.1000     0.5996
##      3           5.9803           nan      0.1000     0.5293
##      4           5.5269           nan      0.1000     0.4154
##      5           5.1301           nan      0.1000     0.3923
##      6           4.7897           nan      0.1000     0.2007
##      7           4.4224           nan      0.1000     0.2482
##      8           4.1528           nan      0.1000     0.1840
##      9           3.8740           nan      0.1000     0.1938
##     10           3.6367           nan      0.1000     0.1894
##     20           2.1725           nan      0.1000     0.0933
##     40           1.0934           nan      0.1000     0.0063
##     60           0.7734           nan      0.1000    -0.0003
##     80           0.6267           nan      0.1000    -0.0113
##    100           0.5248           nan      0.1000    -0.0115
##    120           0.4509           nan      0.1000    -0.0086

```


```
##      140      0.3909      nan      0.1000      -0.0042
##      160      0.3417      nan      0.1000      -0.0095
##      180      0.2979      nan      0.1000      -0.0010
##      200      0.2641      nan      0.1000      -0.0051
##      220      0.2324      nan      0.1000      -0.0048
##      240      0.2065      nan      0.1000      -0.0029
##      250      0.1944      nan      0.1000      -0.0033
##
## Iter   TrainDeviance   ValidDeviance   StepSize   Improve
##      1           7.5853           nan      0.1000     0.4622
##      2           7.2268           nan      0.1000     0.2276
##      3           6.9725           nan      0.1000     0.2317
##      4           6.7255           nan      0.1000     0.1945
##      5           6.5310           nan      0.1000     0.1967
##      6           6.3548           nan      0.1000     0.1452
##      7           6.1331           nan      0.1000     0.1877
##      8           6.0205           nan      0.1000     0.0816
##      9           5.8651           nan      0.1000     0.1151
##     10           5.7272           nan      0.1000     0.0983
##     20           4.5797           nan      0.1000     0.0457
##     40           3.3805           nan      0.1000     0.0281
##     60           2.6821           nan      0.1000    -0.0024
##     80           2.1887           nan      0.1000     0.0083
##    100           1.8328           nan      0.1000     0.0108
##    120           1.5572           nan      0.1000    -0.0039
##    140           1.3454           nan      0.1000     0.0011
##    160           1.2090           nan      0.1000    -0.0017
##    180           1.0998           nan      0.1000     0.0009
##    200           1.0155           nan      0.1000    -0.0061
##    220           0.9510           nan      0.1000    -0.0040
##    240           0.9159           nan      0.1000    -0.0058
##    250           0.8952           nan      0.1000    -0.0036
```

```
carseats.gbm
```

```
## Stochastic Gradient Boosting
##
## No pre-processing
## Resampling: Cross-Validated (10 fold)
## Summary of sample sizes: 289, 289, 289, 289, 289, 289, ...
## Resampling results across tuning parameters:
##
##  interaction.depth  n.trees  RMSE      Rsquared  MAE
##      1              50      1.853439  0.6550573  1.4985804
##      1             100      1.547356  0.7484921  1.2635775
```

```
##      1      150      1.359464  0.7941783  1.1116649
##      1      200      1.263360  0.8152576  1.0273078
##      1      250      1.216972  0.8238354  0.9842987
##      2       50      1.559349  0.7549692  1.2703898
##      2      100      1.295349  0.8121515  1.0550231
##      2      150      1.225522  0.8259530  0.9949898
##      2      200      1.219263  0.8282347  0.9847772
##      2      250      1.222610  0.8269006  0.9882843
##      3       50      1.434103  0.7828979  1.1618345
##      3      100      1.265869  0.8154077  1.0207212
##      3      150      1.242808  0.8187860  0.9946575
##      3      200      1.263512  0.8139170  1.0132197
##      3      250      1.266998  0.8123660  1.0169904
##      4       50      1.387541  0.7897254  1.1230060
##      4      100      1.279528  0.8095927  1.0299009
##      4      150      1.297827  0.8045587  1.0399179
##      4      200      1.316687  0.7979882  1.0556691
##      4      250      1.321038  0.7972775  1.0621347
##      5       50      1.348922  0.7980728  1.0974910
##      5      100      1.277990  0.8108521  1.0389748
##      5      150      1.286355  0.8066147  1.0393349
##      5      200      1.292184  0.8042104  1.0372671
##      5      250      1.314433  0.7973539  1.0572041
##
## Tuning parameter 'shrinkage' was held constant at a value of 0.1
##
## Tuning parameter 'n.minobsinnode' was held constant at a value of 10
## RMSE was used to select the optimal model using the smallest value.
## The final values used for the model were n.trees = 250,
## interaction.depth = 1, shrinkage = 0.1 and n.minobsinnode = 10.
```


```
plot(carseats.gbm)
```



```

carseats.pred <- predict(carseats.gbm, carseats_test, type = "raw")
plot(carseats_test$Sales, carseats.pred,
     main = "Gradient Boosting Regression: Predicted vs. Actual",
     xlab = "Actual",
     ylab = "Predicted")
abline(0,1)

```

```
(carseats.gbm.rmse <- RMSE(pred = carseats.pred,
  obs = carseats_test$Sales))
```

```
## [1] 1.402428
```

```
rm(carseats.pred)
```

```
#plot(varImp(carseats.gbm), main="Variable Importance with Gradient Boosting")
```

8.6 Summary

Okay, I'm going to tally up the results! For the classification division, the winner is the manual classification tree! Gradient boosting made a valiant run at it, but came up just a little short.

```
rbind(data.frame(model = "Manual Class", Acc = round(oj_model_1b_cm$overall["Accuracy"], 5)),
  data.frame(model = "Class w.tuneGrid", Acc = round(oj_model_3_cm$overall["Accuracy"], 5)),
  data.frame(model = "Bagging", Acc = round(oj.bag.acc, 5)),
  data.frame(model = "Random Forest", Acc = round(oj.frst.acc, 5)),
  data.frame(model = "Gradient Boosting", Acc = round(oj.gbm.acc, 5))
) %>% arrange(desc(Acc))
```

```
##           model      Acc
## 1      Manual Class 0.85915
## 2 Gradient Boosting 0.85446
## 3 Class w.tuneGrid 0.84507
## 4           Bagging 0.82629
## 5      Random Forest 0.82629
```

And now for the regression division, the winner is... gradient boosting!

```
rbind(data.frame(model = "Manual ANOVA", RMSE = round(carseats_model_1_pruned_rmse, 5)),
      data.frame(model = "ANOVA w.tuneGrid", RMSE = round(carseats_model_3_pruned_rmse, 5)),
      data.frame(model = "Bagging", RMSE = round(carseats.bag.rmse, 5)),
      data.frame(model = "Random Forest", RMSE = round(carseats.frst.rmse, 5)),
      data.frame(model = "Gradient Boosting", RMSE = round(carseats.gbm.rmse, 5))
) %>% arrange(RMSE)
```

```
##           model      RMSE
## 1 Gradient Boosting 1.40243
## 2      Random Forest 1.75811
## 3           Bagging 1.93279
## 4 ANOVA w.tuneGrid 2.29833
## 5      Manual ANOVA 2.38806
```


Here are plots of the ROC curves for all the models (one from each chapter) on the same graph. The ROCR package provides the `prediction()` and `performance()` functions which generate the data required for plotting the ROC curve, given a set of predictions and actual (true) values. The more “up and to the left” the ROC curve of a model is, the better the model. The AUC performance metric is literally the “Area Under the ROC Curve”, so the greater the area under this curve, the higher the AUC, and the better-performing the model is.

```
library(ROCR)
# List of predictions
oj.class.pred <- predict(oj_model_3, oj_test, type = "prob")[,2]
oj.bag.pred <- predict(oj.bag, oj_test, type = "prob")[,2]
oj.frst.pred <- predict(oj.frst, oj_test, type = "prob")[,2]
oj.gbm.pred <- predict(oj.gbm, oj_test, type = "prob")[,2]

preds_list <- list(oj.class.pred, oj.bag.pred, oj.frst.pred, oj.gbm.pred)
# preds_list <- list(oj.class.pred)

# List of actual values (same for all)
m <- length(preds_list)
actuals_list <- rep(list(oj_test$Purchase), m)
```

```
# Plot the ROC curves
pred <- prediction(preds_list, actuals_list)
#pred <- prediction(oj.class.pred[,2], oj_test$Purchase)
rocs <- performance(pred, "tpr", "fpr")
plot(rocs, col = as.list(1:m), main = "Test Set ROC Curves")
legend(x = "bottomright",
      legend = c("Decision Tree", "Bagged Trees", "Random Forest", "GBM"),
      fill = 1:m)
```



Chapter 9

Reference

Penn State University, STAT 508: Applied Data Mining and Statistical Learning, “Lesson 11: Tree-based Methods”. <https://newonlinecourses.science.psu.edu/stat508/lesson/11>.

Brownlee, Jason. “Classification And Regression Trees for Machine Learning”, Machine Learning Mastery. <https://machinelearningmastery.com/classification-and-regression-trees-for-machine-learning/>.

Brownlee, Jason. “A Gentle Introduction to the Gradient Boosting Algorithm for Machine Learning”, Machine Learning Mastery. <https://machinelearningmastery.com/gentle-introduction-gradient-boosting-algorithm-machine-learning/>.

DataCamp: Machine Learning with Tree-Based Models in R

An Introduction to Statistical Learning by Gareth James, et al.

SAS Documentation

StatMethods: Tree-Based Models

Machine Learning Plus

GBM (Boosted Models) Tuning Parameters from Listen Data

Harry Southworth on GitHub

Gradient Boosting Classification with GBM in R in DataTechNotes

Molnar, Christoph. “Interpretable machine learning. A Guide for Making Black Box Models Explainable”, 2019. <https://christophm.github.io/interpretable-ml-book/>.

Chapter 10

Support Vector Machines

These notes rely on (James et al., 2013), (Hastie et al., 2017), and (Kuhn and Johnson, 2016). I also reviewed the material in PSU’s Applied Data Mining and Statistical Learning (STAT 508), and the *e1071* Support Vector Machines vignette.

The Support Vector Machines (SVM) algorithm finds the optimal separating hyperplane between members of two classes using an appropriate nonlinear mapping to a sufficiently high dimension. The hyperplane is defined by the observations that lie within a margin optimized by a cost hyperparameter. These observations are called the *support vectors*.

SVM is an extension of the *support vector classifier* which in turn is a generalization of the simple and intuitive *maximal margin classifier*.

10.1 Maximal Margin Classifier

The maximal margin classifier is the optimal hyperplane defined in the (rare) case where two classes are *linearly separable*. Given an $n \times p$ data matrix X with binary response variable defined as $y \in [-1, 1]$ it *may* be possible to define a p -dimensional hyperplane $h(X) = \beta_0 + \beta_1 X_1 + \beta_2 X_2 \cdots + \beta_p X_p = x_i^T \beta + \beta_0 = 0$ such that all observations of each class fall on opposite sides of the hyperplane. This “separating hyperplane” has the property that if β is constrained to be a unit vector, $\|\beta\| = \sum \beta^2 = 1$, then the product of the hyperplane and response variables are positive perpendicular distances from the hyperplane, the smallest of which may be termed the hyperplane *margin*, M ,

$$y_i(x_i' \beta + \beta_0) \geq M.$$

The maximal margin classifier is the hyperplane with the maximum margin. That is, $\max\{M\}$ subject to $\|\beta\| = 1$. A separating hyperplane rarely exists. In fact, even if a separating hyperplane does exist, its resulting margin is probably undesirably narrow.

10.2 Support Vector Classifier

The maximal margin classifier can be generalized to non-separable cases using a so-called “soft margin”. The generalization is called the *support vector classifier*. The soft margin allows some misclassification in the interest of greater robustness to individual observations. The support vector classifier optimizes

$$y_i(x_i'\beta + \beta_0) \geq M(1 - \xi_i)$$

where the ξ_i are positive *slack variables* whose sum is bounded by some constant tuning parameter $\sum \xi_i \leq \text{constant}$. The slack variable values indicate where the observation lies: $\xi_i = 0$ observations lie on the correct side of the margin; $\xi_i > 0$ observations lie on the wrong side of the margin; $\xi_i > 1$ observations lie on the wrong side of the hyperplane. The constant sets the tolerance for margin violation. If $\text{constant} = 0$, then all observations must reside on the correct side of the margin, as in the maximal margin classifier. The *constant* controls the bias-variance trade-off. As the *constant* increases, the margin widens and allows more violations. The classifier bias increases but its variance decreases.

The support vector classifier is usually defined by dropping the $\|\beta\| = 1$ constraint, and defining $M = 1/\|\beta\|$. The optimization problem then becomes

$$\min \|\beta\| \quad s.t. \quad \begin{cases} y_i(x_i^T \beta + \beta_0) \geq 1 - \xi_i, & \forall i \\ \xi_i \geq 0, & \sum \xi_i \leq \text{constant}. \end{cases}$$

This is a quadratic equation with linear inequality constraints, so it is a convex optimization problem which can be solved using Lagrange multipliers. Re-express the optimization problem as

$$\min_{\beta_0, \beta} \frac{1}{2} \|\beta\|^2 = C \sum_{i=1}^N \xi_i \quad s.t. \quad \xi_i \geq 0, \quad y_i(x_i^T \beta + \beta_0) \geq 1 - \xi_i, \quad \forall i$$

where the “cost” parameter C replaces the constant and penalizes large residuals. This optimization problem is equivalent to *another* optimization problem, the familiar *loss + penalty* formulation:

$$\min_{\beta_0, \beta} \sum_{i=1}^N [1 - y_i f(x_i)]_+ + \frac{\lambda}{2} \|\beta\|^2$$

where $\lambda = 1/C$ and $[1 - y_i f(x_i)]_+$ is a “hinge” loss function with $f(x_i) = \text{sign}[Pr(Y = +1|x) - 1/2]$.

The parameter estimates can be written as functions of a set of unknown parameters (α_i) and data points. The solution to the optimization problem requires only the inner products of the observations, represented as $\langle x_i, x_j \rangle$,

$$f(x) = \beta_0 + \sum_{i=1}^n \alpha_i \langle x, x_i \rangle$$

The solution has the interesting property that only observations on or within the margin affect the hyperplane. These observations are known as support vectors. As the constant increases, the number of violating observations increase, and thus the number of support vectors increases. This property makes the algorithm robust to the extreme observations far away from the hyperplane.

The parameter estimators for α_i are nonzero only for the support vectors in the solution—that is, if a training observation is not a support vector, then its α_i equals zero.

The only shortcoming with the algorithm is that it presumes a linear decision boundary.

10.3 Support Vector Machines

Enlarging the feature space of the support vector classifier accommodates non-linear relationships. Support vector machines do this in a specific way, using *kernels*. The kernel is a generalization of the inner product with form $K(x_i, x'_i)$. So the linear kernel is simply

$$K(x_i, x'_i) = \langle x, x_i \rangle$$

and the solution is

$$f(x) = \beta_0 + \sum_{i=1}^n \alpha_i K(x_i, x'_i)$$

K can take another form instead, such as polynomial

$$K(x, x') = (\gamma \langle x, x' \rangle + c_0)^d$$

or radial

$$K(x, x') = \exp\{-\gamma \|x - x'\|^2\}.$$


10.4 Example

Here is a data set of two classes $y \in [-1, 1]$ described by two features $X1$ and $X2$.

```
library(tidyverse)
set.seed(1)
x <- matrix(rnorm(20*2), ncol=2)
y <- c(rep(-1, 10), rep(1, 10))
x[y==1, ] <- x[y==1, ] + 1
train_data <- data.frame(x, y)
train_data$y <- as.factor(y)
```


A scatter plot reveals whether the classes are linearly separable.

```
ggplot(train_data, aes(x = X1, y = X2, color = y)) +
  geom_point(size = 2) +
  labs(title = "Binary response with two features") +
  theme(legend.position = "top")
```



No, they are not linearly separable. Now fit a support vector machine. The **e1071** library implements the SVM algorithm. `svm(..., kernel="linear")` fits a support vector classifier. Change the kernel to `c("polynomial", "radial")` for SVM. Try a cost of 10.

```
library(e1071)
m <- svm(
  y ~ .,
  data = train_data,
  kernel = "linear",
  type = "C-classification", # (default) for classification
  cost = 10, # default is 1
  scale = FALSE # do not standardize features
)
plot(m, train_data)
```



The support vectors are plotted as “x’s”. There are seven of them.

```
m$index
```

```
## [1] 1 2 5 7 14 16 17
```

The summary shows adds additional information, including the distribution of the support vector classes.


```
summary(m)
```

```
##
```

```
## Call:
## svm(formula = y ~ ., data = train_data, kernel = "linear", type = "C-classification",
##      cost = 10, scale = FALSE)
##
##
## Parameters:
##   SVM-Type:  C-classification
## SVM-Kernel:  linear
##       cost:  10
##
## Number of Support Vectors:  7
##
##   ( 4 3 )
##
##
## Number of Classes:  2
##
## Levels:
##   -1 1
```

The seven support vectors are comprised of four in one class, three in the other. What if we lower the cost of margin violations? This will increase bias and lower variance.

```
m <- svm(
  y ~ .,
  data = train_data,
  kernel = "linear",
  type = "C-classification",
  cost = 0.1,
  scale = FALSE
)
plot(m, train_data)
```



There are many more support vectors now. *(In case you hoped to see the linear decision boundary formulation, or at least a graphical representation of the margins, keep hoping. The model is generalized beyond two features, so it evidently does not worry too much about supporting sanitized two-feature demos.)*

Which cost level yields the *best* predictive performance on holdout data? Use cross validation to find out. SVM defaults to 10-fold CV. I'll try seven candidate values for `cost`.

```
set.seed(1)
m_tune <- tune(
  svm,
  y ~ .,
  data = train_data,
  kernel = "linear",
  ranges = list(cost = c(0.001, 0.01, 0.1, 1, 5, 10, 100))
)
summary(m_tune)
```

```
##
## Parameter tuning of 'svm':
##
## - sampling method: 10-fold cross validation
##
## - best parameters:
```

```
## cost
## 0.1
##
## - best performance: 0.05
##
## - Detailed performance results:
## cost error dispersion
## 1 1e-03 0.55 0.4377975
## 2 1e-02 0.55 0.4377975
## 3 1e-01 0.05 0.1581139
## 4 1e+00 0.15 0.2415229
## 5 5e+00 0.15 0.2415229
## 6 1e+01 0.15 0.2415229
## 7 1e+02 0.15 0.2415229
```


The lowest cross-validation error rate is 0.10 with `cost = 0.1`. `tune()` saves the best tuning parameter value.

```
m_best <- m_tune$best.model
summary(m_best)
```

```
##
## Call:
## best.tune(method = svm, train.x = y ~ ., data = train_data, ranges = list(cost = c(
## 0.01, 0.1, 1, 5, 10, 100)), kernel = "linear")
##
##
## Parameters:
## SVM-Type: C-classification
## SVM-Kernel: linear
## cost: 0.1
##
## Number of Support Vectors: 16
##
## ( 8 8 )
##
##
## Number of Classes: 2
##
## Levels:
## -1 1
```


There are 16 support vectors, 8 in each class. This is a pretty wide margin.

```
plot(m_best, train_data)
```




What if the classes had been linearly separable? Then we could create a maximal margin classifier.

```
train_data_2 <- train_data %>%
  mutate(
    X1 = X1 + ifelse(y==1, 1.0, 0),
    X2 = X2 + ifelse(y==1, 1.0, 0)
  )
ggplot(train_data_2, aes(x = X1, y = X2, color = y)) +
  geom_point(size = 2) +
  labs(title = "Binary response with two features, linearly separable")
```



Specify a huge cost = 1e5 so that no support vectors violate the margin.

```
m2 <- svm(  
  y ~ .,  
  data = train_data_2,  
  kernel = "linear",  
  cost = 1e5,  
  scale = FALSE # do not standardize features  
)  
plot(m2, train_data_2)
```

```
summary(m2)
```

```
##
## Call:
## svm(formula = y ~ ., data = train_data_2, kernel = "linear",
##      cost = 1e+05, scale = FALSE)
##
##
## Parameters:
##   SVM-Type:  C-classification
##   SVM-Kernel: linear
##       cost:  1e+05
##
## Number of Support Vectors:  3
##
##   ( 1 2 )
##
##
## Number of Classes:  2
##
## Levels:
##   -1 1
```

This model will have very low bias, but very high variance. To fit an SVM, use a different kernel. You can use `kernel = c("polynomial", "radial",`

"sigmoid"). For a polynomial model, also specify the polynomial degree. For a radial model, include the gamma value.

```
set.seed(1)
m3_tune <- tune(
  svm,
  y ~ .,
  data = train_data,
  kernel = "polynomial",
  ranges = list(
    cost = c(0.001, 0.01, 0.1, 1, 5, 10, 100),
    degree = c(1, 2, 3)
  )
)
summary(m3_tune)
```

```
##
## Parameter tuning of 'svm':
##
## - sampling method: 10-fold cross validation
##
## - best parameters:
##   cost degree
##     1      1
##
## - best performance: 0.1
##
## - Detailed performance results:
##   cost degree error dispersion
## 1 1e-03      1 0.55 0.4377975
## 2 1e-02      1 0.55 0.4377975
## 3 1e-01      1 0.30 0.2581989
## 4 1e+00      1 0.10 0.2108185
## 5 5e+00      1 0.10 0.2108185
## 6 1e+01      1 0.15 0.2415229
## 7 1e+02      1 0.15 0.2415229
## 8 1e-03      2 0.70 0.4216370
## 9 1e-02      2 0.70 0.4216370
## 10 1e-01     2 0.70 0.4216370
## 11 1e+00     2 0.65 0.2415229
## 12 5e+00     2 0.50 0.3333333
## 13 1e+01     2 0.50 0.3333333
## 14 1e+02     2 0.50 0.3333333
## 15 1e-03     3 0.65 0.3374743
## 16 1e-02     3 0.65 0.3374743
```

```
## 17 1e-01      3  0.50  0.3333333
## 18 1e+00      3  0.40  0.3162278
## 19 5e+00      3  0.35  0.3374743
## 20 1e+01      3  0.35  0.3374743
## 21 1e+02      3  0.35  0.3374743
```


The lowest cross-validation error rate is 0.10 with cost = 1, polynomial degree 1.

```
m3_best <- m3_tune$best.model
summary(m3_best)
```

```
##
## Call:
## best.tune(method = svm, train.x = y ~ ., data = train_data, ranges = list(cost = c(0.001,
##      0.01, 0.1, 1, 5, 10, 100), degree = c(1, 2, 3)), kernel = "polynomial")
##
##
## Parameters:
##   SVM-Type:  C-classification
## SVM-Kernel:  polynomial
##      cost:   1
##    degree:   1
##   coef.0:   0
##
## Number of Support Vectors:  12
##
##   ( 6 6 )
##
##
## Number of Classes:  2
##
## Levels:
##   -1 1
```

There are 12 support vectors, 6 in each class. This is a pretty wide margin.

```
plot(m3_best, train_data)
```



10.5 Using Caret

The model can also be fit using **caret**. I'll used LOOCV since the data set is so small. Normalize the variables to make their scale comparable.

```
library(caret)
library(kernlab)


train_data_3 <- train_data %>%
  mutate(y = factor(y, labels = c("A", "B")))

m4 <- train(
  y ~ .,
  data = train_data_3,
  method = "svmPoly",
  preProcess = c("center", "scale"),
  trControl = trainControl(
    method = "cv",
    number = 5,
    summaryFunction = twoClassSummary, # Use AUC to pick the best model
    classProbs=TRUE
  )
)
```

```
m4$bestTune
```

```
## degree scale C
## 8      1  0.1 0.5
```

```
plot(m4)
```



Chapter 11

Principal Components Analysis

Chapter 12

Clustering

Chapter 13

Text Mining

Appendix

Here are miscellaneous skills, knowledge, and technologies I should know.

Publishing to BookDown

The **bookdown** package, written by Yihui Xie, is built on top of R Markdown and the **knitr** package. Use it to publish a book or long manuscript where each chapter is a separate file. There are instructions for how to author a book in his bookdown book (Xie, 2019). The main advantage of **bookdown** over R Markdown is that you can produce multi-page HTML output with numbered headers, equations, figures, etc., just like in a book. I'm using **bookdown** to create a compendium of all my data science notes.

The first step to using **bookdown** is installing the `**bookdown*` package with `install.packages("bookdown")`.

Next, create an account at bookdown.org, and connect the account to RStudio. Follow the instructions at <https://bookdown.org/home/about/>.

Finally, create a project in R Studio by creating a new project of type *Book Project using Bookdown*.

After creating all of your Markdown pages, knit the book or click the **Build Book** button in the Build panel.

Shiny Apps

Chapter 14

rstudio::conf

- 14.1 Open Source Software for Data Science
- 14.2 Data, visualization, and designing with AI
- 14.3 Deploying End-to-End Data Science with Shiny, Plumber, and Pins
- 14.4 Health Connected Siloed Data Sources and Streamlined Reporting Using R
- 14.5 Building a new data science pipeline for teh FT with RStudio Connect
- 14.6 How to win an AI Hackathon without using AI
- 14.7 Production-grade Shiny Apps with golem
- 14.8 Making the Shiny Contest
- 14.9 Styling Shiny apps with Sass and Bootstrap
4
- 14.10 R: Then and Now
- 14.11 Journalism with RStudio, R, and teh Tidyverse
- 14.12 Learning R with the RStudio::conf

Bibliography

Fawcett, T. (2005). *An introduction to ROC analysis*. ELSEVIER.

Hastie, T., Tibshirani, R., and Friedman, J. (2017). *The Elements of Statistical Learning*. Springer, New York, NY, 2nd edition. ISBN 978-0387848570-0387848570.

James, G., Witten, D., Hastie, T., and Tibshirani, R. (2013). *An Introduction to Statistical Learning: with Applications in R*. Springer, New York, NY, 1st edition. ISBN 978-1461471370.

Kuhn, M. and Johnson, K. (2016). *Applied Predictive Modeling*. Springer, New York, NY, 1st edition. ISBN 978-1-4614-6848-6.

Therneau, T. and Atkinson, E. (2019). *An Introduction to Recursive Partitioning Using the RPART Routines*. Chapman and Hall/CRC, Boca Raton, Florida.

Xie, Y. (2019). *bookdown: Authoring Books and Technical Documents with R Markdown*. Chapman and Hall/CRC, Boca Raton, Florida, 1st edition. ISBN 978113870010.